

Chapter 1: Neural Network Foundations with TensorFlow 2.0

tensorflow/tensorflow ● C++ ★ 123k

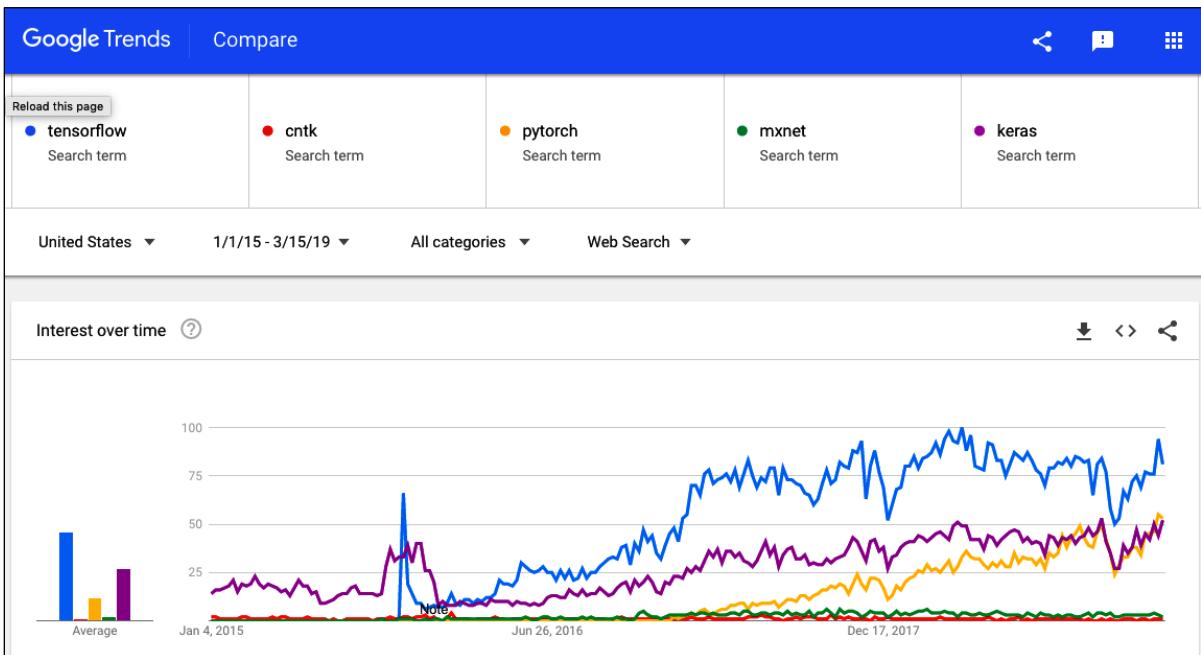
An Open Source Machine Learning Framework for Everyone

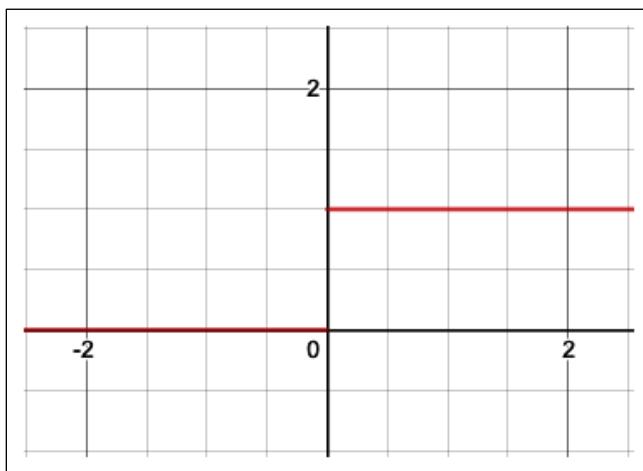
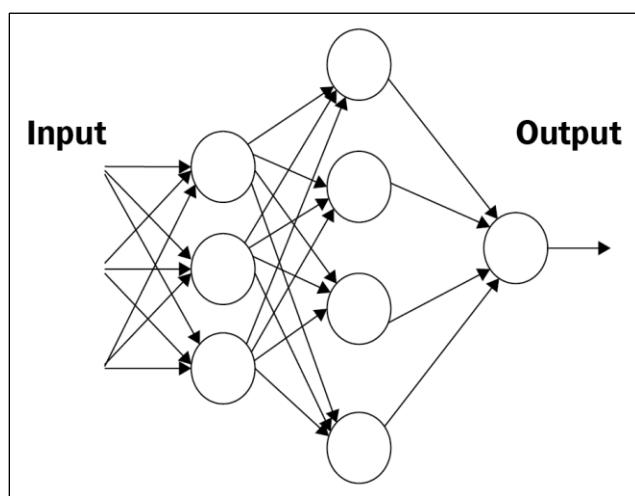
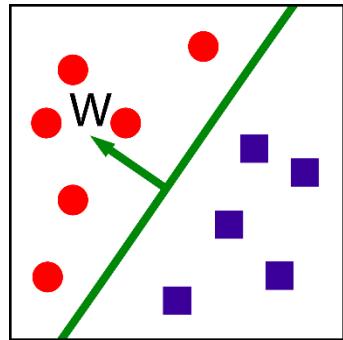
keras-team/keras ● Python ★ 39.1k

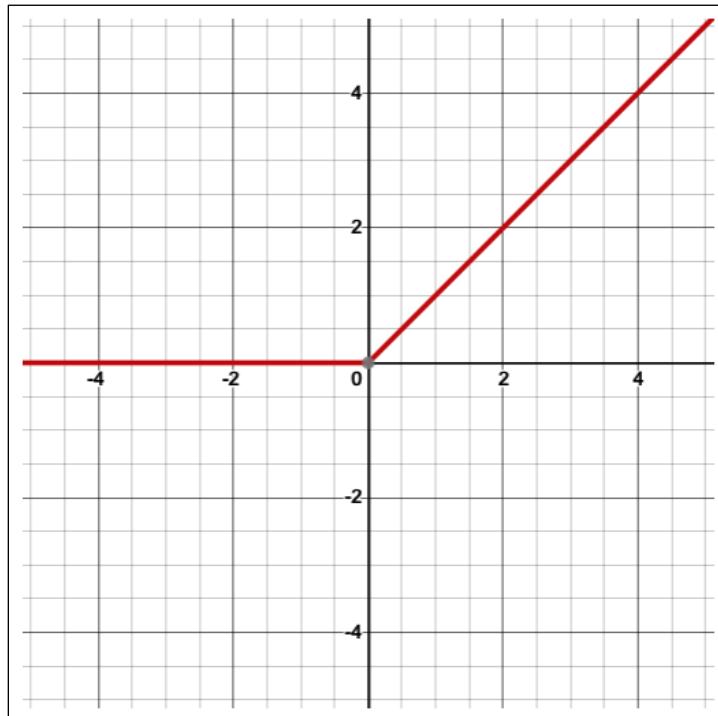
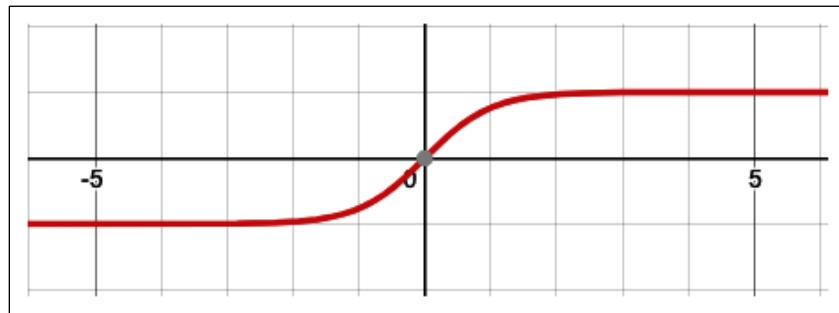
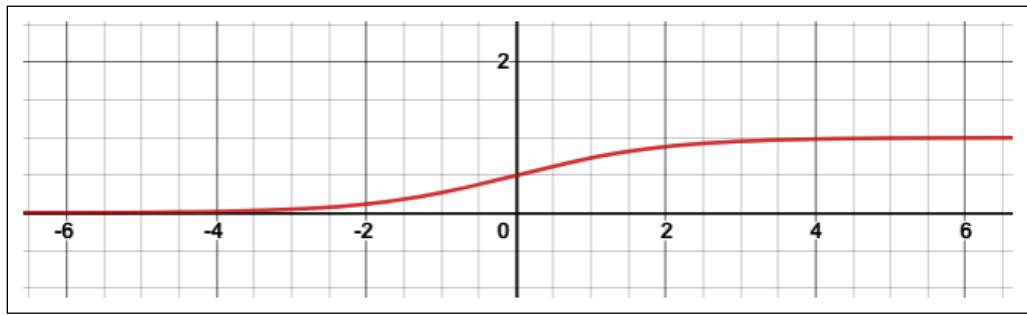
Deep Learning for humans

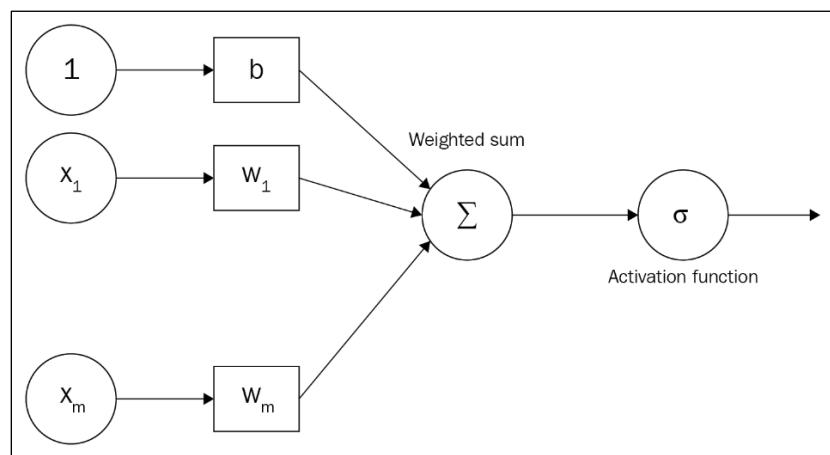
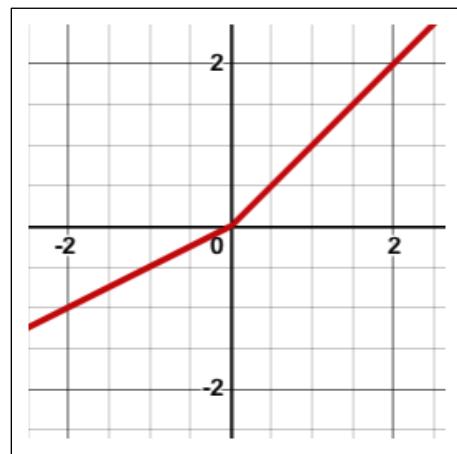
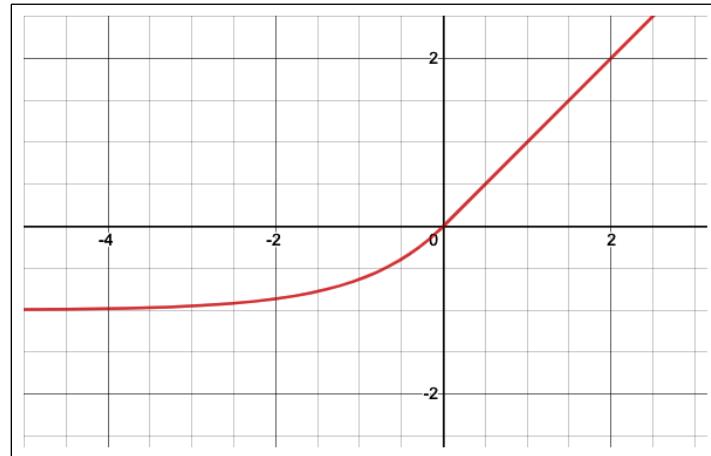
pytorch / pytorch ★ 25.8k

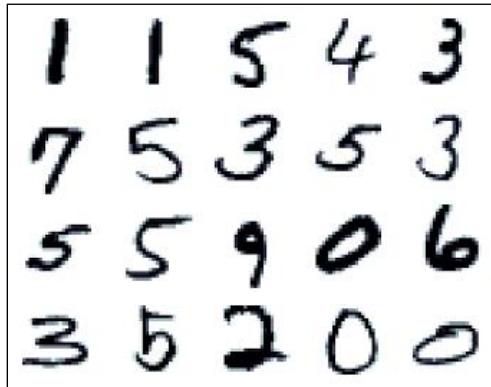
Tensors and Dynamic neural networks in Python with strong GPU acceleration











Model: "sequential"

Layer (type)	Output Shape	Param #
dense_layer (Dense)	(None, 10)	7850
Total params:	7,850	
Trainable params:	7,850	
Non-trainable params:	0	

Train on 48000 samples, validate on 12000 samples
Epoch 1/200
48000/48000 [=====] - 1s 31us/sample - loss: 2.1276 - accuracy: 0.2322 - val_loss: 1.9508 - val_accuracy: 0.3908
Epoch 2/200
48000/48000 [=====] - 1s 23us/sample - loss: 1.8251 - accuracy: 0.5141 - val_loss: 1.6848 - val_accuracy: 0.6277
Epoch 3/200
48000/48000 [=====] - 1s 25us/sample - loss: 1.5992 - accuracy: 0.6531 - val_loss: 1.4838 - val_accuracy: 0.7150
Epoch 4/200
48000/48000 [=====] - 1s 27us/sample - loss: 1.4281 - accuracy: 0.7115 - val_loss: 1.3304 - val_accuracy: 0.7551
Epoch 5/200

Epoch 199/200
48000/48000 [=====] - 1s 22us/sample - loss: 0.3684 - accuracy: 0.8995 - val_loss: 0.3464 - val_accuracy: 0.9071
Epoch 200/200
48000/48000 [=====] - 1s 23us/sample - loss: 0.3680 - accuracy: 0.8996 - val_loss: 0.3461 - val_accuracy: 0.9070
10000/10000 [=====] - 1s 54us/sample - loss: 0.3465 - accuracy: 0.9071
Test accuracy: 0.9071

Layer (type)	Output Shape	Param #
dense_layer (Dense)	(None, 128)	100480
dense_layer_2 (Dense)	(None, 128)	16512
dense_layer_3 (Dense)	(None, 10)	1290
Total params:	118,282	
Trainable params:	118,282	
Non-trainable params:	0	

Train on 48000 samples, validate on 12000 samples		
Epoch 1/200		
48000/48000 [=====] - 3s 63us/sample - loss: 2.2507 - accuracy: 0.2086 - val_loss: 2.1592 - val_accuracy: 0.3266		

```

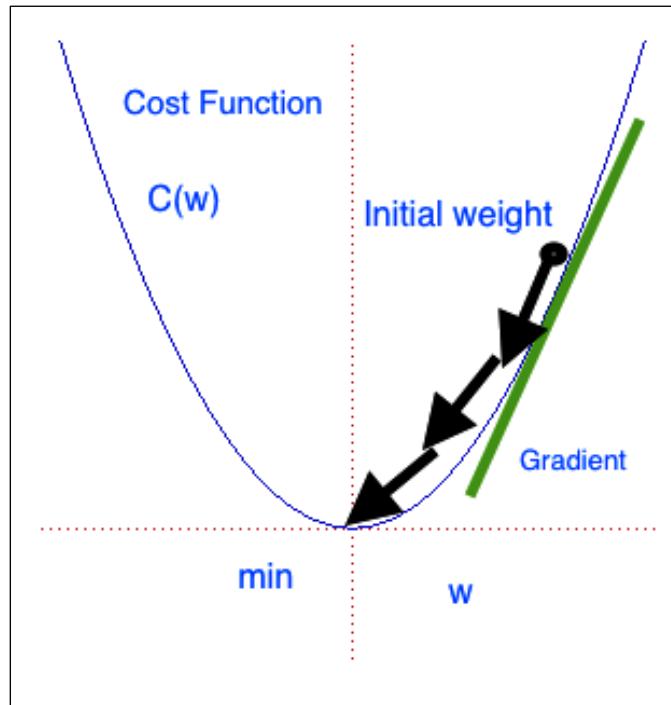
Epoch 49/50
48000/48000 [=====] - 1s 30us/sample - loss: 0.3347 - accuracy: 0.9075 - val_loss: 0.3126 - val_accuracy: 0.9136
Epoch 50/50
48000/48000 [=====] - 1s 28us/sample - loss: 0.3326 - accuracy: 0.9081 - val_loss: 0.3107 - val_accuracy: 0.9140
10000/10000 [=====] - 0s 40us/sample - loss: 0.3164 - accuracy: 0.9118
Test accuracy: 0.9118

```

```

Epoch 199/200
48000/48000 [=====] - 2s 45us/sample - loss: 0.2850 - accuracy: 0.9177 - val_loss: 0.1922 - val_accuracy: 0.9442
Epoch 200/200
48000/48000 [=====] - 2s 42us/sample - loss: 0.2845 - accuracy: 0.9170 - val_loss: 0.1917 - val_accuracy: 0.9442
10000/10000 [=====] - 1s 61us/sample - loss: 0.1927 - accuracy: 0.9415
Test accuracy: 0.9415

```



Layer (type)	Output Shape	Param #
dense_layer (Dense)	(None, 128)	100480
dropout (Dropout)	(None, 128)	0
dense_layer_2 (Dense)	(None, 128)	16512
dropout_1 (Dropout)	(None, 128)	0
dense_layer_3 (Dense)	(None, 10)	1290

Total params: 118,282
Trainable params: 118,282
Non-trainable params: 0

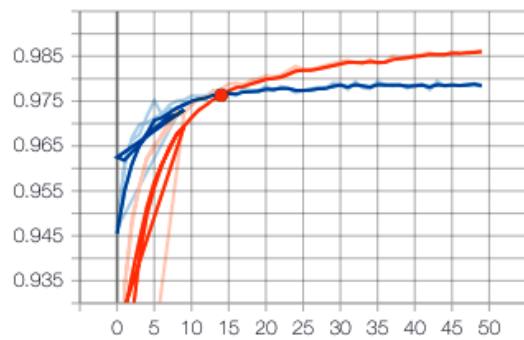
Train on 48000 samples, validate on 12000 samples

Epoch 1/10
48000/48000 [=====] - 2s 48us/sample - loss: 0.4715 -
accuracy: 0.8575 - val_loss: 0.1820 - val_accuracy: 0.9471
Epoch 2/10
48000/48000 [=====] - 2s 36us/sample - loss: 0.2215 -
accuracy: 0.9341 - val_loss: 0.1268 - val_accuracy: 0.9631
Epoch 3/10
48000/48000 [=====] - 2s 39us/sample - loss: 0.1684 -
accuracy: 0.9497 - val_loss: 0.1198 - val_accuracy: 0.9651
Epoch 4/10
48000/48000 [=====] - 2s 43us/sample - loss: 0.1459 -
accuracy: 0.9569 - val_loss: 0.1059 - val_accuracy: 0.9710
Epoch 5/10
48000/48000 [=====] - 2s 39us/sample - loss: 0.1273 -
accuracy: 0.9623 - val_loss: 0.1059 - val_accuracy: 0.9696
Epoch 6/10
48000/48000 [=====] - 2s 36us/sample - loss: 0.1177 -
accuracy: 0.9659 - val_loss: 0.0941 - val_accuracy: 0.9731
Epoch 7/10
48000/48000 [=====] - 2s 35us/sample - loss: 0.1083 -
accuracy: 0.9671 - val_loss: 0.1009 - val_accuracy: 0.9715
Epoch 8/10
48000/48000 [=====] - 2s 35us/sample - loss: 0.0971 -
accuracy: 0.9706 - val_loss: 0.0950 - val_accuracy: 0.9758
Epoch 9/10
48000/48000 [=====] - 2s 35us/sample - loss: 0.0969 -
accuracy: 0.9718 - val_loss: 0.0985 - val_accuracy: 0.9745
Epoch 10/10
48000/48000 [=====] - 2s 35us/sample - loss: 0.0873 -
accuracy: 0.9743 - val_loss: 0.0966 - val_accuracy: 0.9762
10000/10000 [=====] - 0s 37us/sample - loss: 0.0922 -
accuracy: 0.9764
Test accuracy: 0.9764

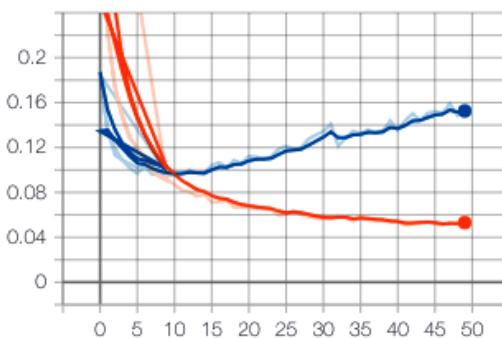
```
Epoch 248/250
48000/48000 [=====] - 2s 40us/sample - loss: 0.0506 -
accuracy: 0.9904 - val_loss: 0.3465 - val_accuracy: 0.9762
Epoch 249/250
48000/48000 [=====] - 2s 40us/sample - loss: 0.0490 -
accuracy: 0.9905 - val_loss: 0.3645 - val_accuracy: 0.9765
Epoch 250/250
48000/48000 [=====] - 2s 39us/sample - loss: 0.0547 -
accuracy: 0.9899 - val_loss: 0.3353 - val_accuracy: 0.9766
10000/10000 [=====] - 1s 58us/sample - loss: 0.3184 -
accuracy: 0.9779
Test accuracy: 0.9779
```

epoch_accuracy

epoch_accuracy



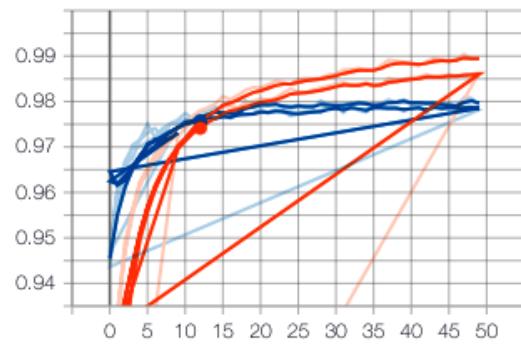
Name	Smoothed	Value	Step	Time	Relative
train	0.9763	0.9776	14	Sun Mar 24, 10:03:38	3m 43s
validation	0.9763	0.9762	14	Sun Mar 24, 10:03:38	3m 43s



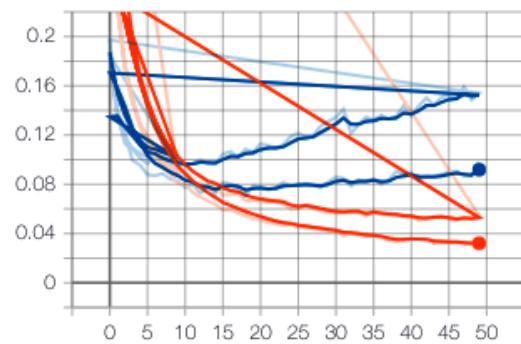
```
Epoch 49/50
48000/48000 [=====] - 3s 55us/sample - loss: 0.0313 -
accuracy: 0.9894 - val_loss: 0.0868 - val_accuracy: 0.9808
Epoch 50/50
48000/48000 [=====] - 2s 51us/sample - loss: 0.0321 -
accuracy: 0.9894 - val_loss: 0.0983 - val_accuracy: 0.9789
10000/10000 [=====] - 1s 66us/sample - loss: 0.0964 -
accuracy: 0.9782
Test accuracy: 0.9782
```

epoch_accuracy

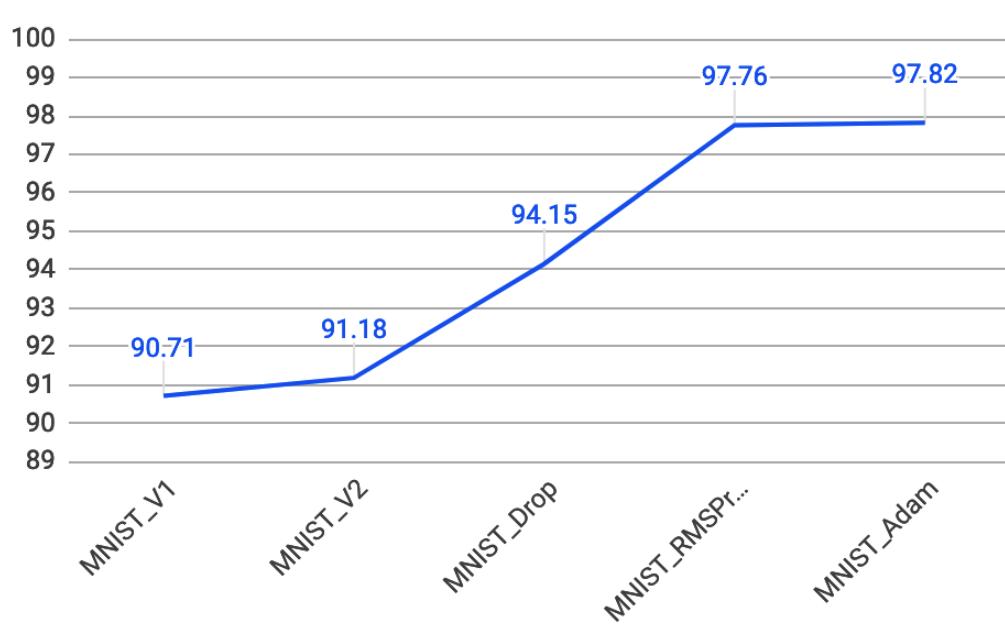
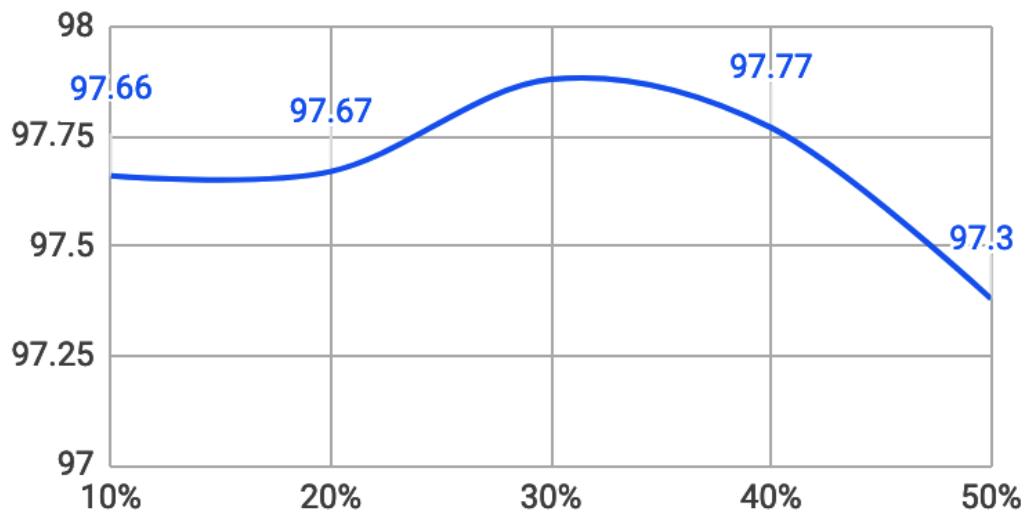
epoch_accuracy

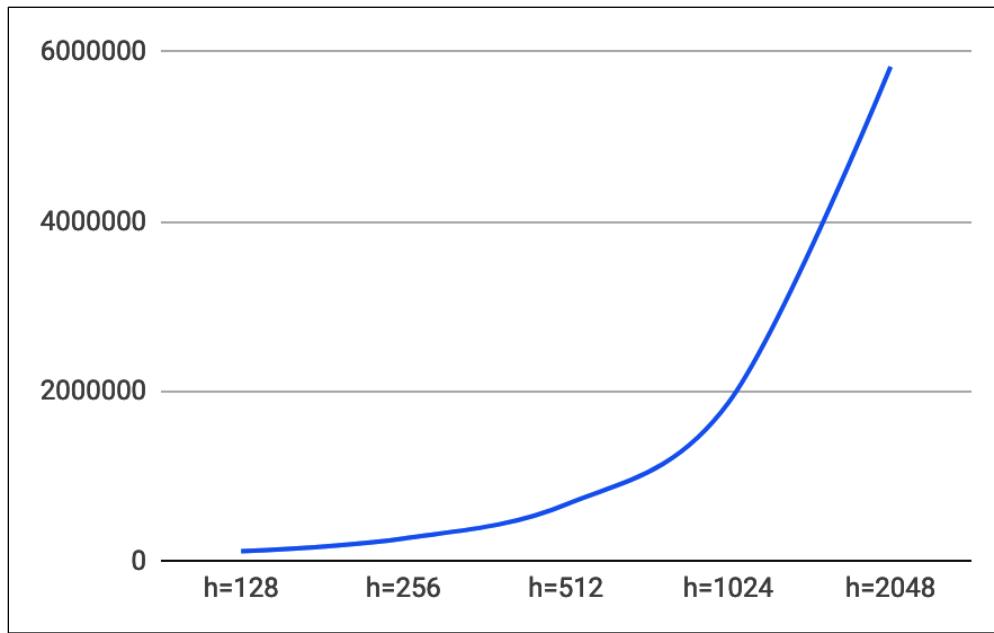
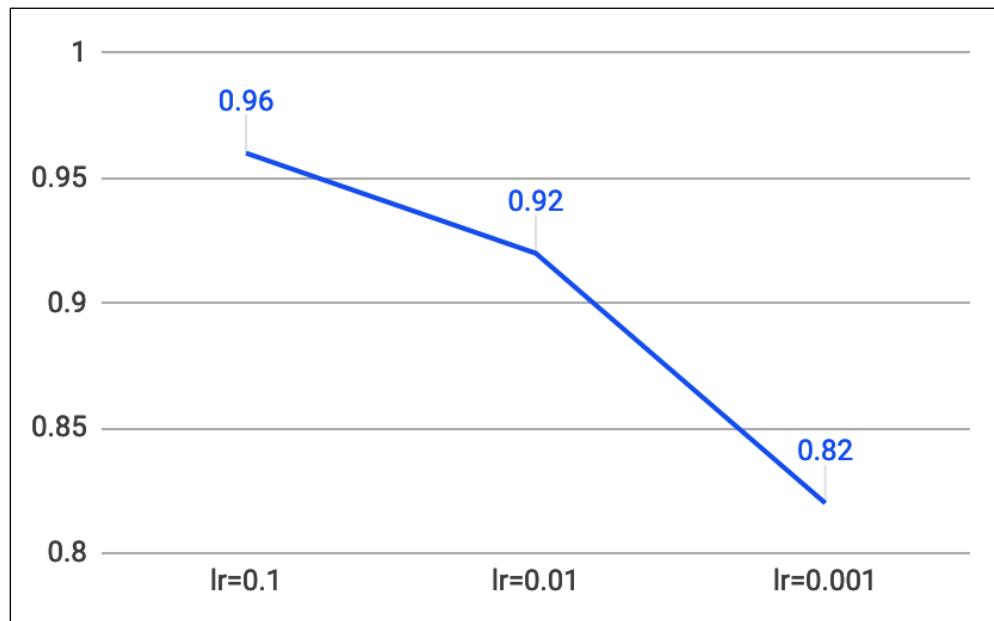


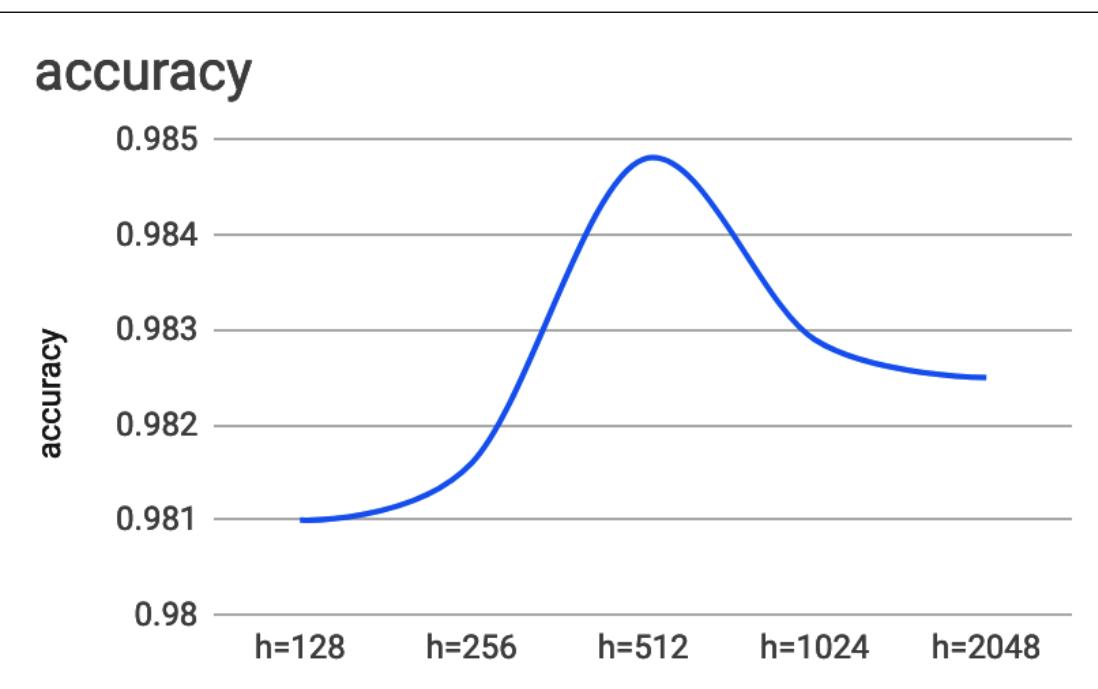
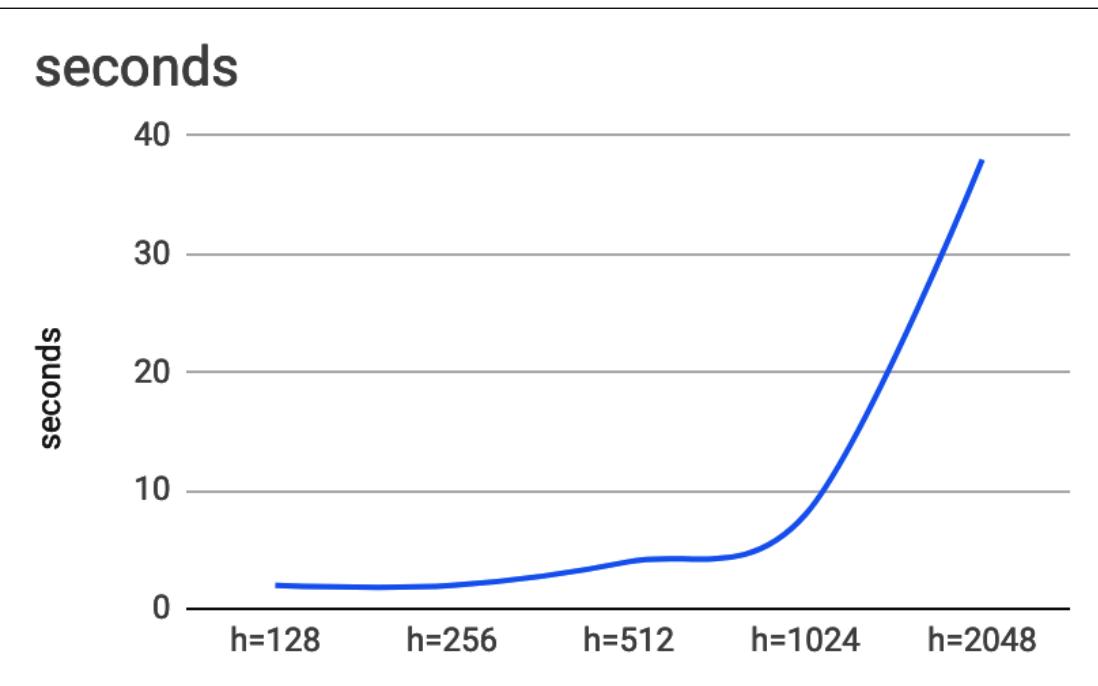
Name	Smoothed	Value	Step	Time	Relative
train	0.9742	0.9759	12	Sun Mar 24, 10:03:34	3m 40s
validation	0.9757	0.9762	12	Sun Mar 24, 10:03:34	3m 40s

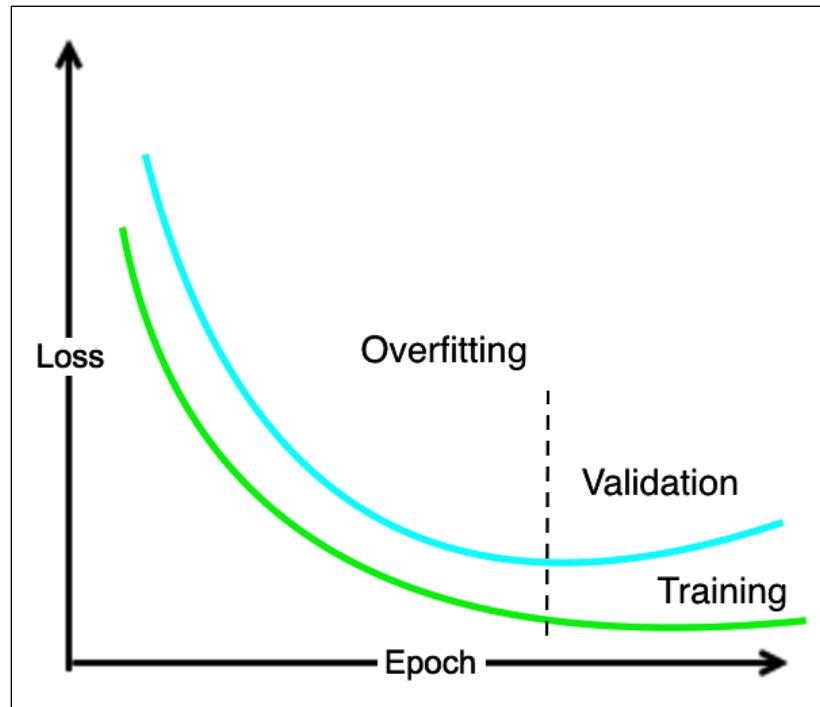
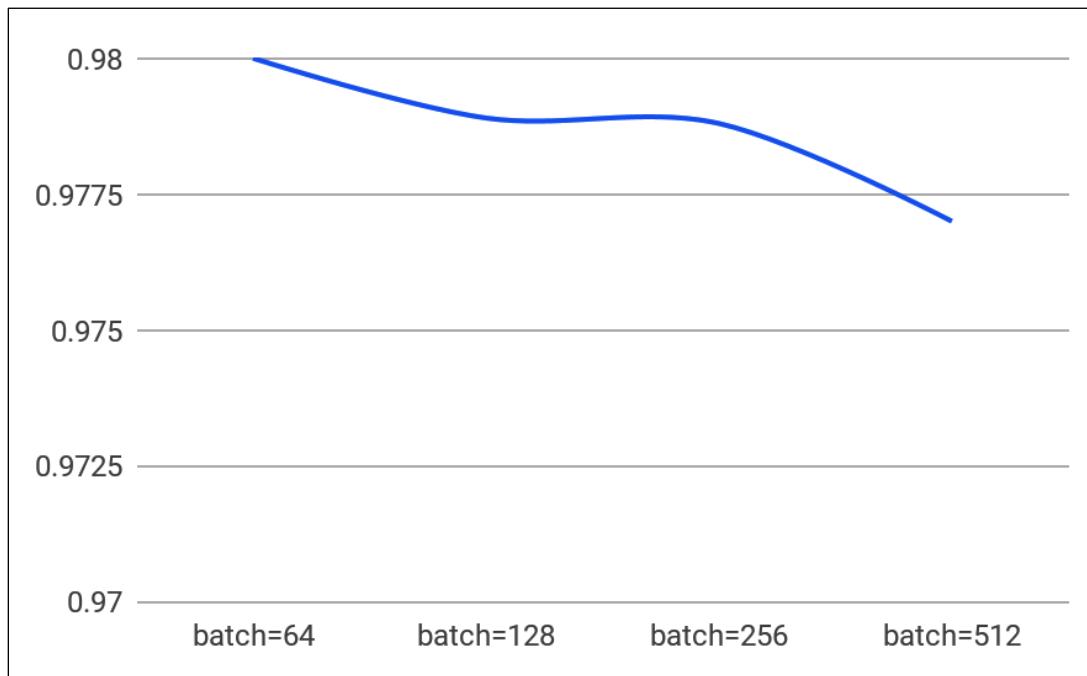


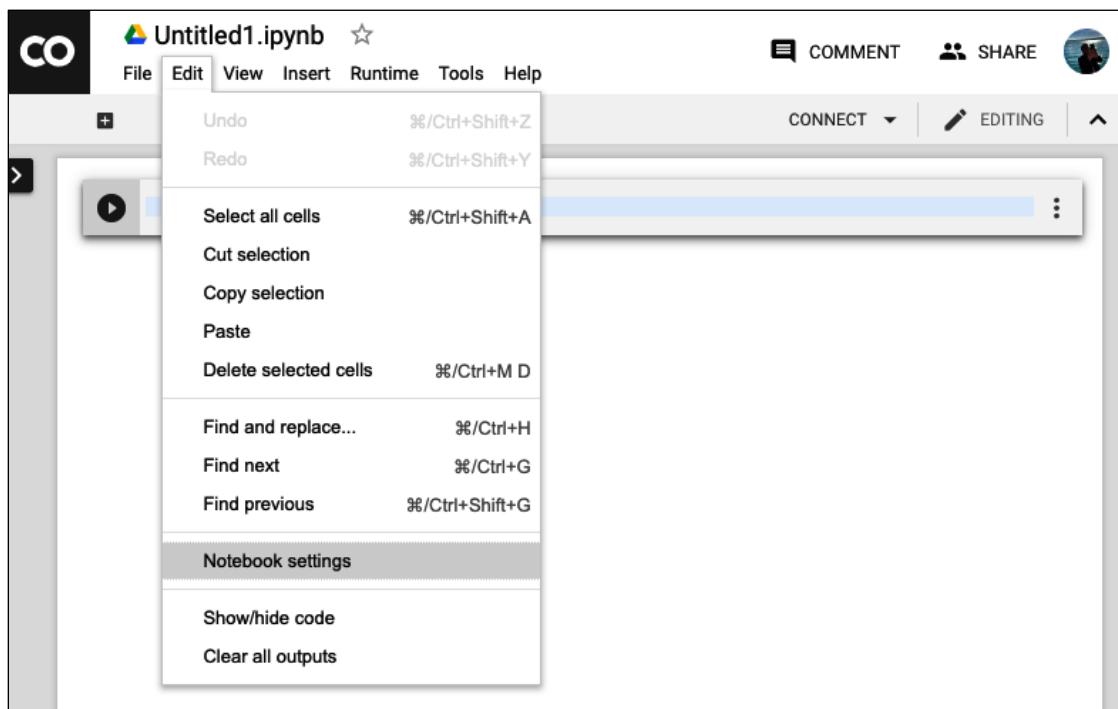
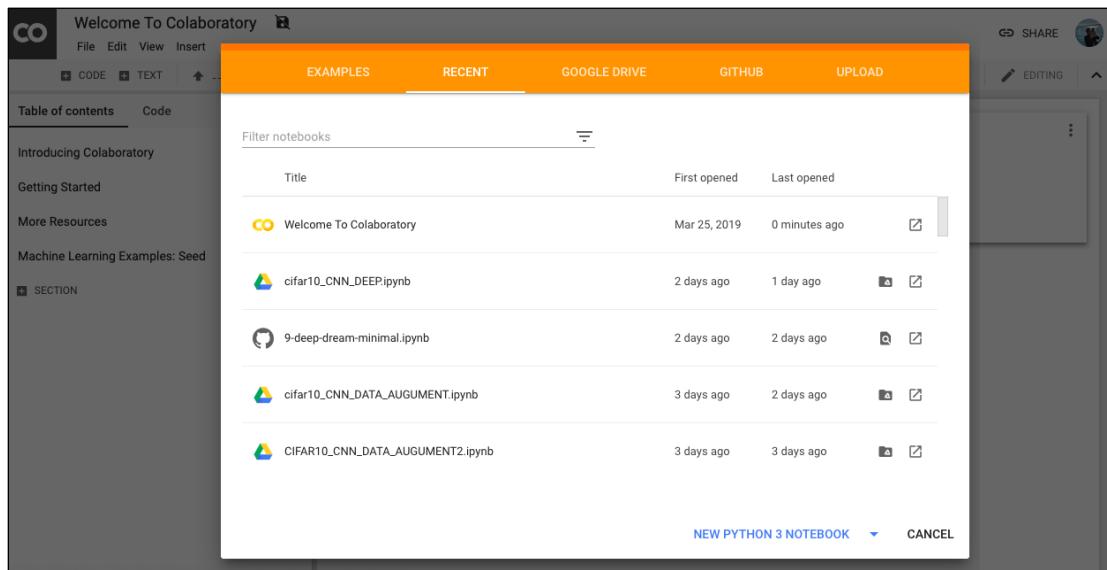
Dropout











Notebook settings

Runtime type

Python 3

Hardware accelerator

None



Omit code cell output when saving this notebook

CANCEL

SAVE

CODE TEXT

CELL CELL

```
return (X_train, y_train), (X_test, y_test)

def build_model():
    model = models.Sequential()
    #Input - Embedding Layer
    # the model will take as input an integer matrix of size (batch, input_length)
    # the model will output dimension (input_length, dim_embedding)
    # the largest integer in the input should be no larger
    # than n_words (vocabulary size).
    model.add(layers.Embedding(n_words,
                               dim_embedding, input_length=max_len))

    model.add(layers.Dropout(0.3))

    #takes the maximum value of either feature vector from each of the n_words features
    model.add(layers.GlobalMaxPooling1D())
    model.add(layers.Dense(128, activation='relu'))
    model.add(layers.Dropout(0.5))
    model.add(layers.Dense(1, activation='sigmoid'))

    return model

(X_train, y_train), (X_test, y_test) = load_data()
model=build_model()
model.summary()

model.compile(optimizer = "adam", loss = "binary_crossentropy",
               metrics = ["accuracy"]
              )

score = model.fit(X_train, y_train,
                   epochs= EPOCHS,
                   batch_size = BATCH_SIZE,
                   validation_data = (X_test, y_test)
                  )

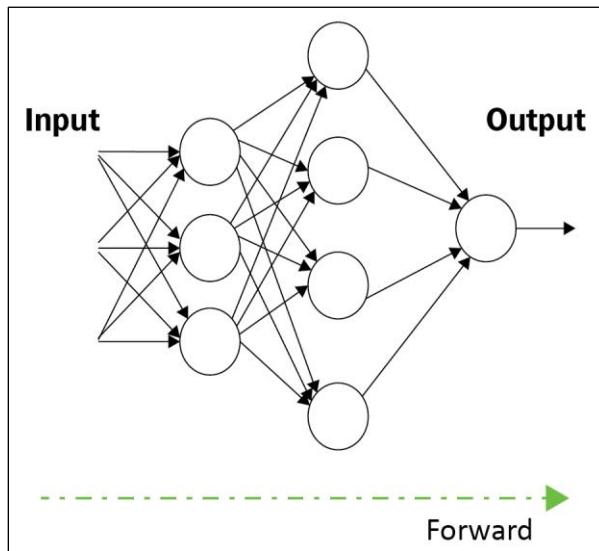
score = model.evaluate(X_test, y_test, batch_size=BATCH_SIZE)
print("\nTest score:", score[0])
print('Test accuracy:', score[1])
```

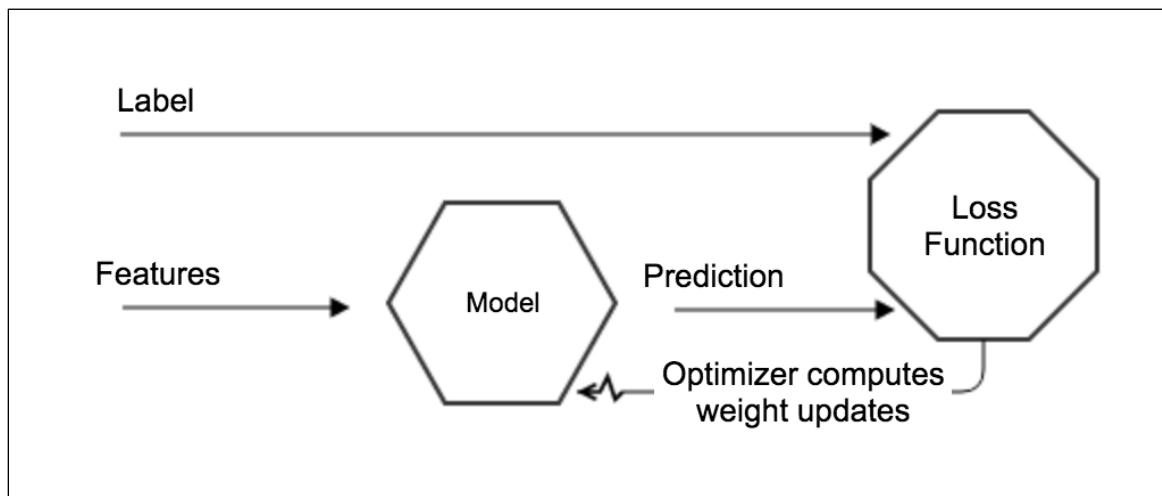
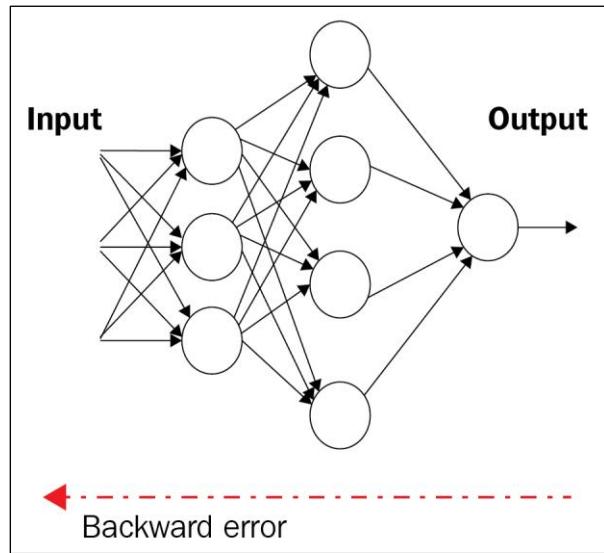
Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 200, 256)	2560000
dropout (Dropout)	(None, 200, 256)	0
global_max_pooling1d (Global)	(None, 256)	0
dense (Dense)	(None, 128)	32896
dropout_1 (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 1)	129

Total params: 2,593,025
Trainable params: 2,593,025
Non-trainable params: 0

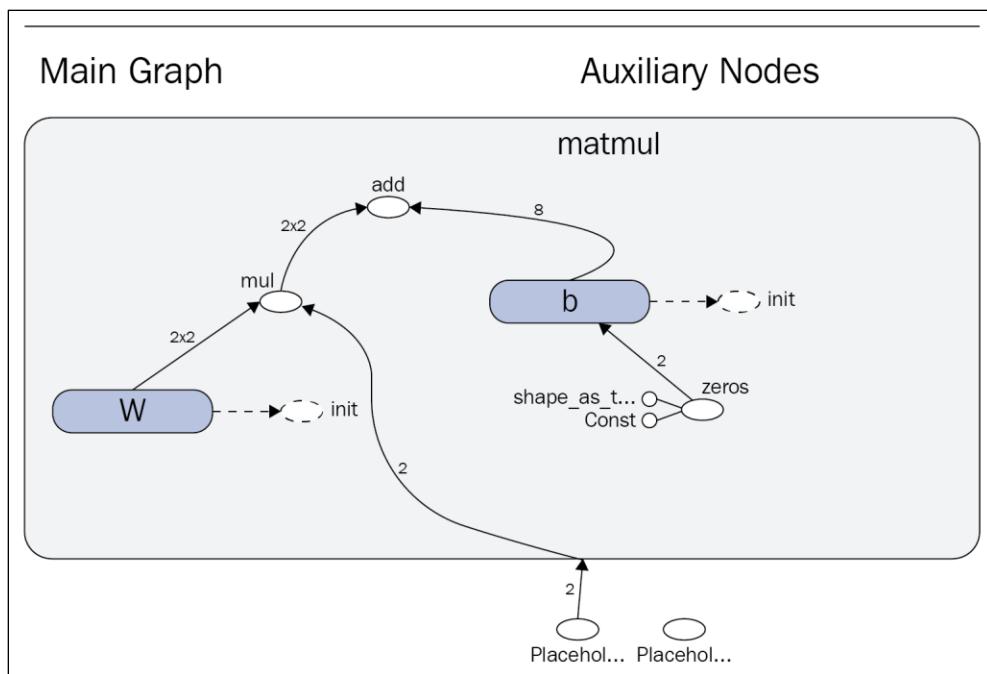
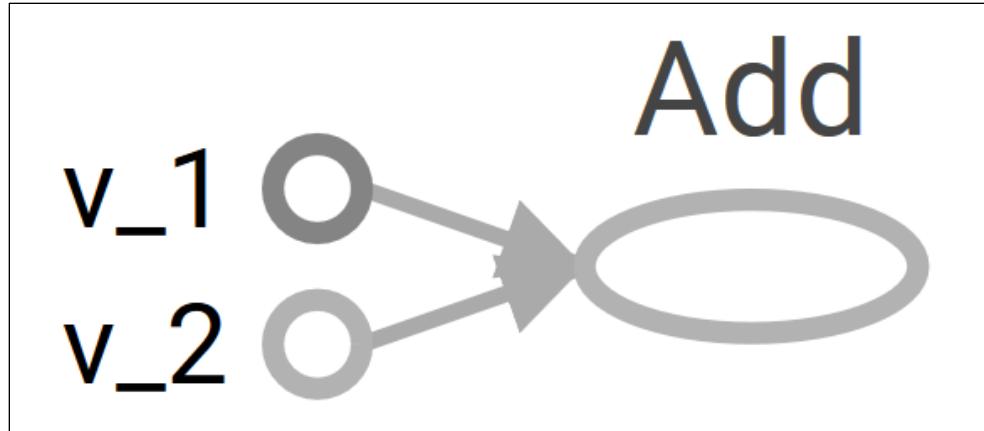
```
Epoch 20/20
25000/25000 [=====] - 23s 925us/sample - loss: 0.0053 - accuracy: 0.9991 - val_
loss: 0.4993 - val_accuracy: 0.8503
25000/25000 [=====] - 2s 74us/sample - loss: 0.4993 - accuracy: 0.8503
```

```
Test score: 0.4992710727453232
Test accuracy: 0.85028
```

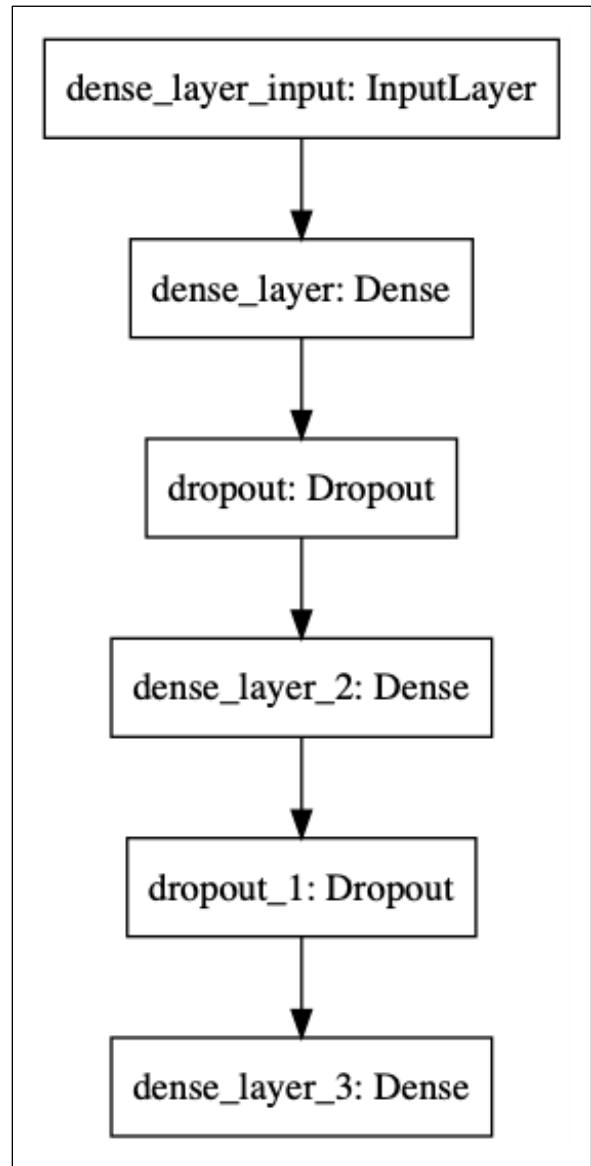


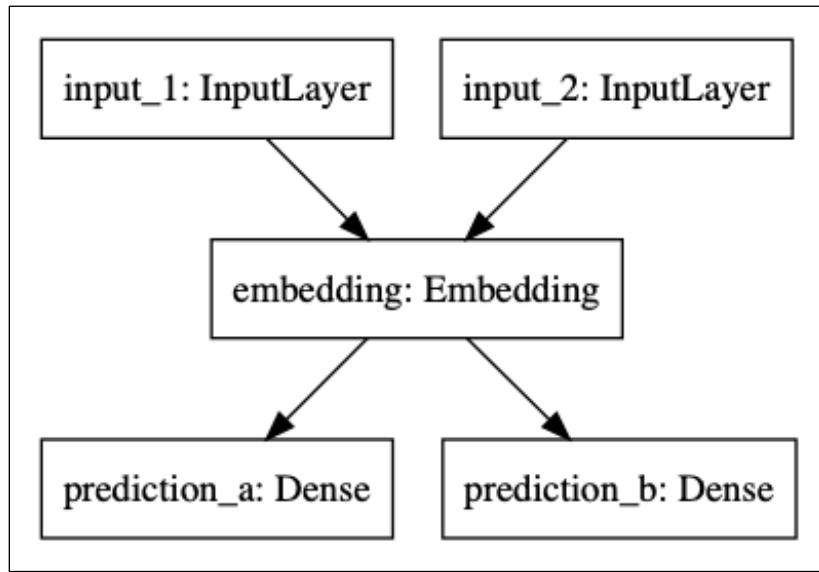


Chapter 2: TensorFlow 1.x and 2.x

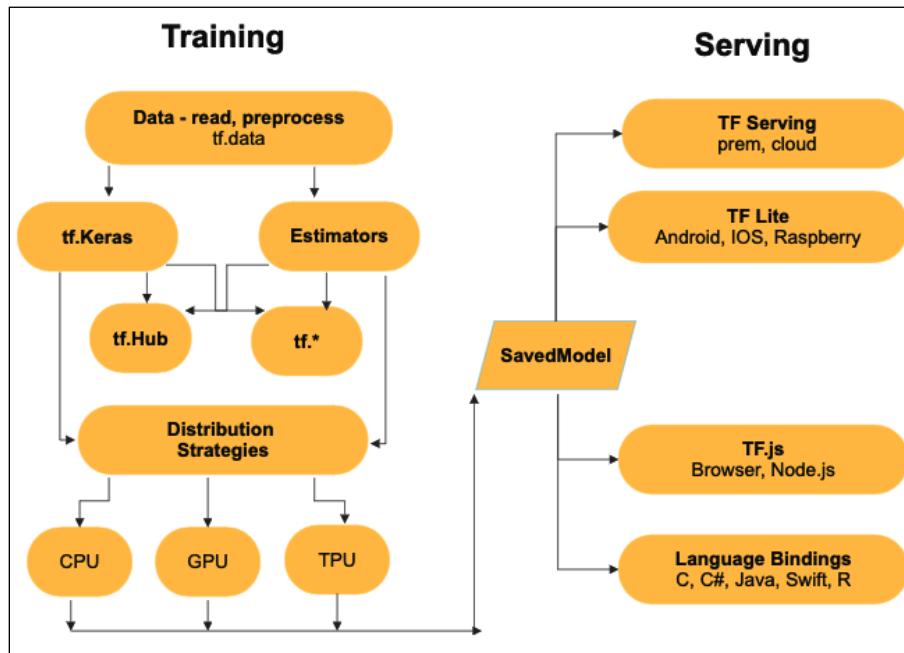
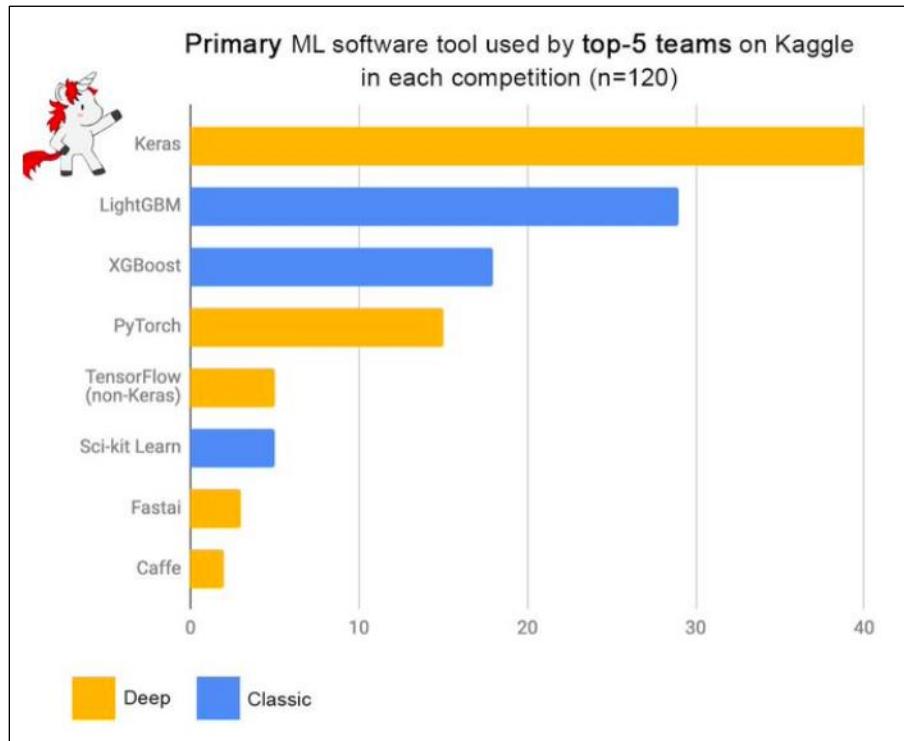


```
graph_time: 0.4504085020016646  
auto_graph_time: 0.07892408400221029
```

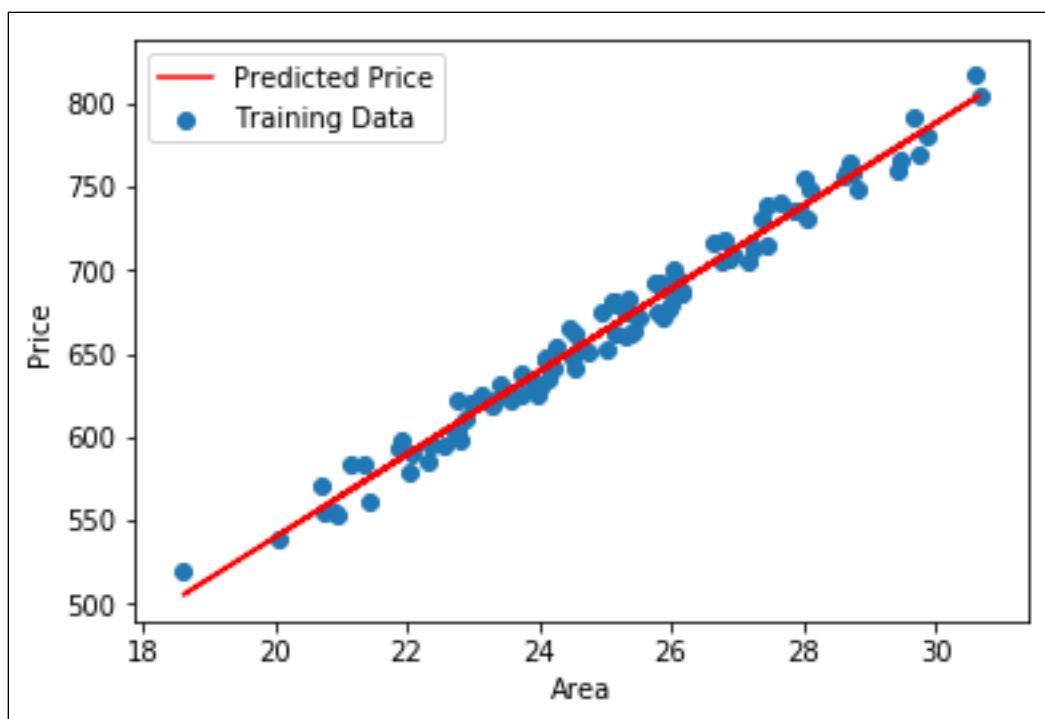
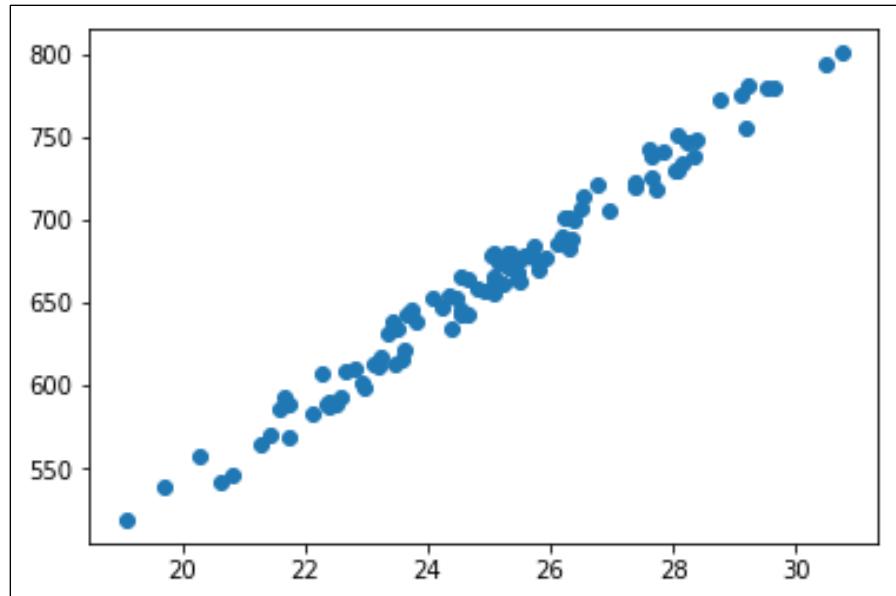




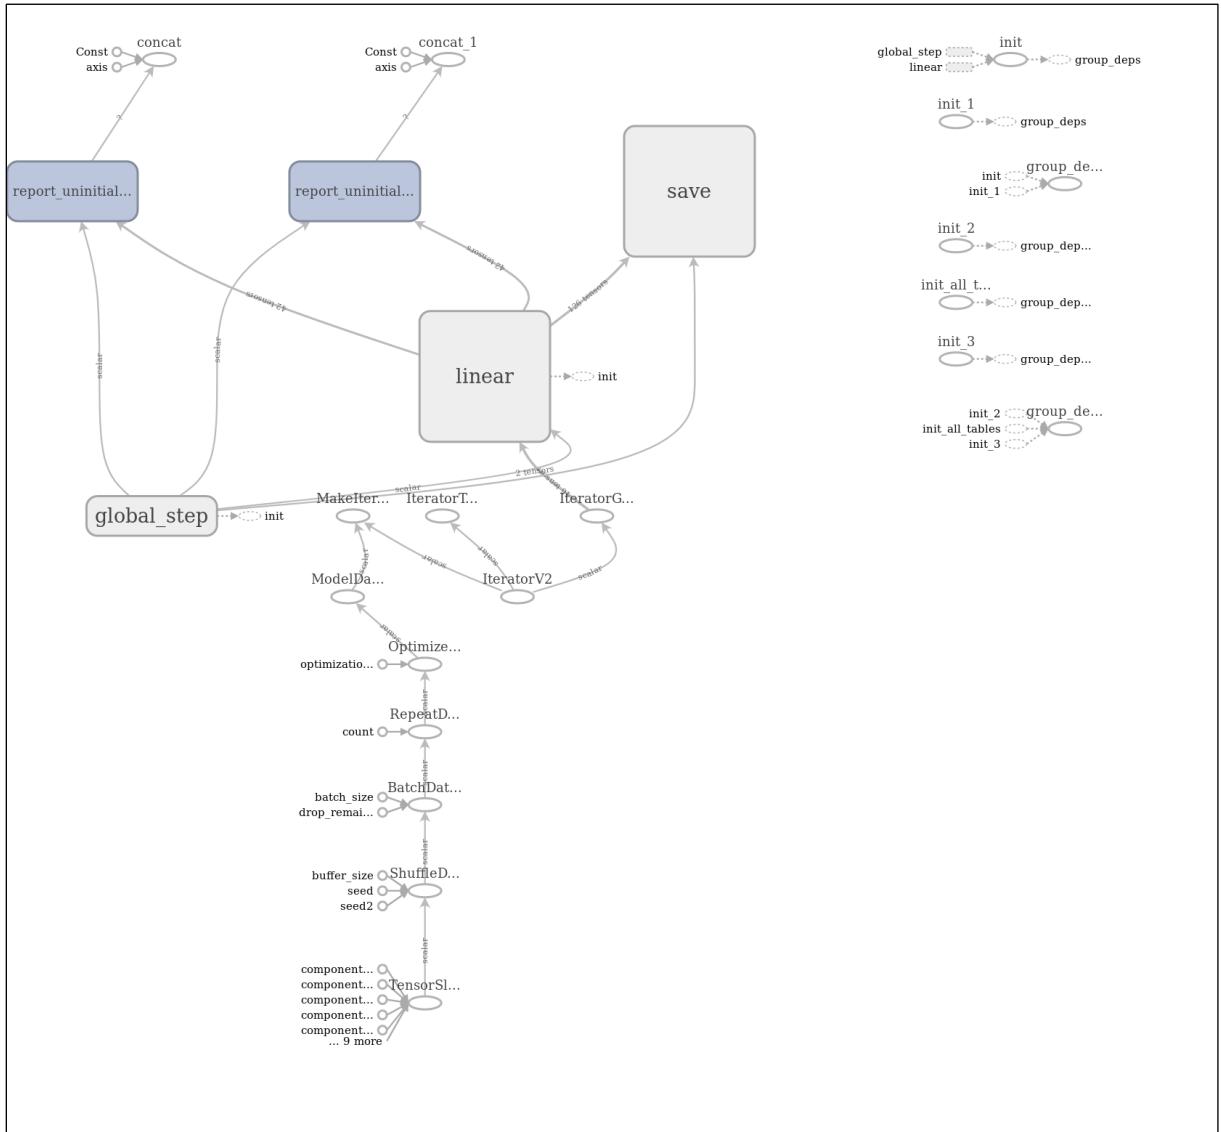
Training API	MirroredStrategy	TPUStrategy	MultiWorkerMirroredStrategy	CentralStorageStrategy	ParameterServerStrategy
Keras API	Supported	Experimental support	Experimental support	Experimental support	Supported planned post 2.0
Custom training loop	Experimental support	Experimental support	Support planned post 2.0	Support planned post 2.0	No support yet
Estimator API	Limited Support	Not supported	Limited Support	Limited Support	Limited Support

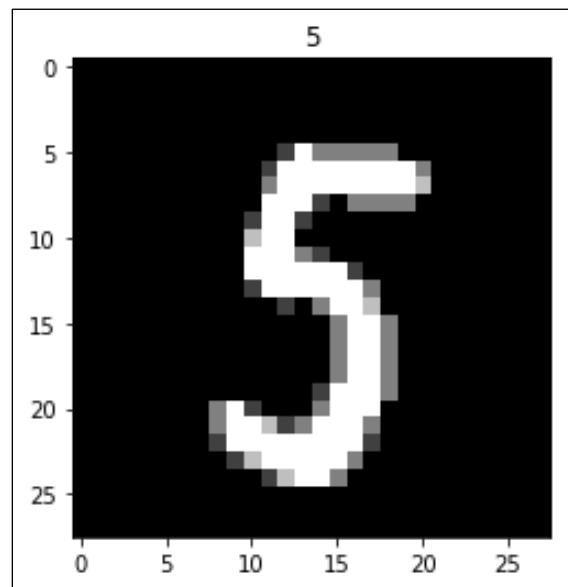
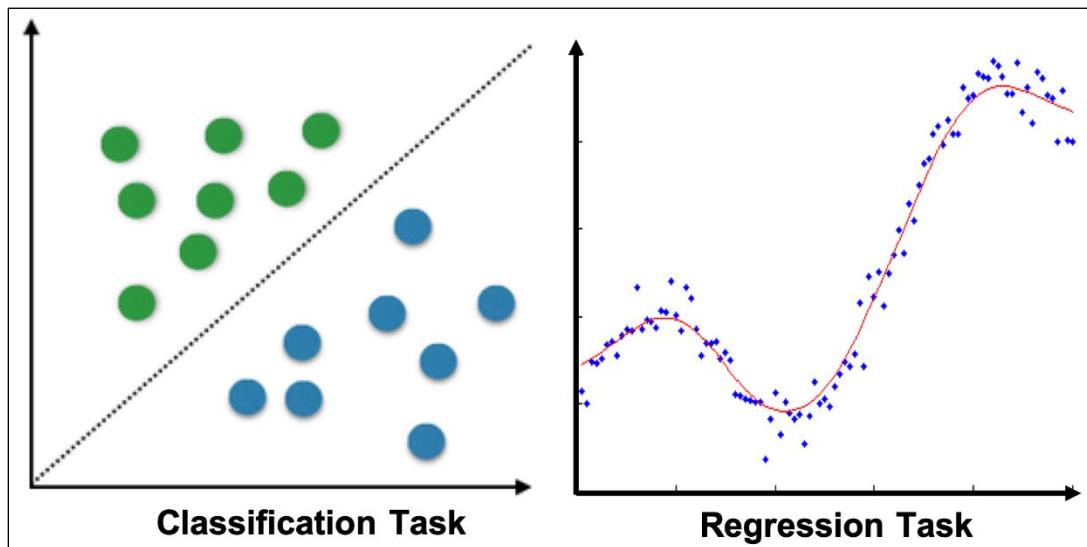


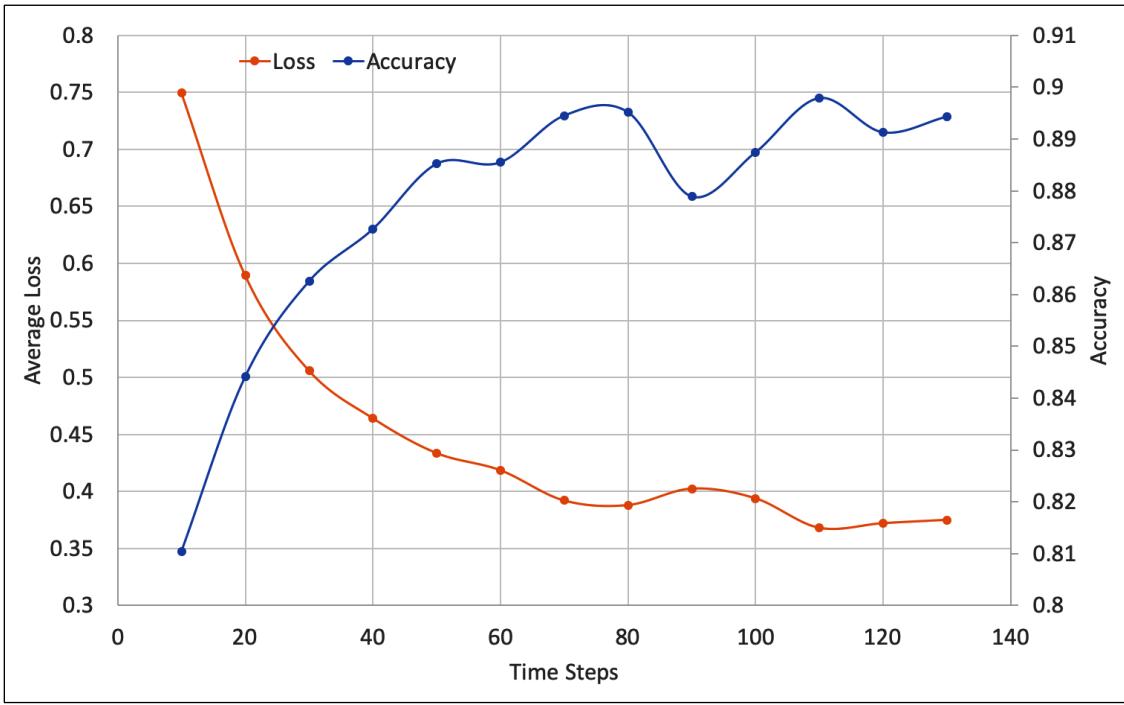
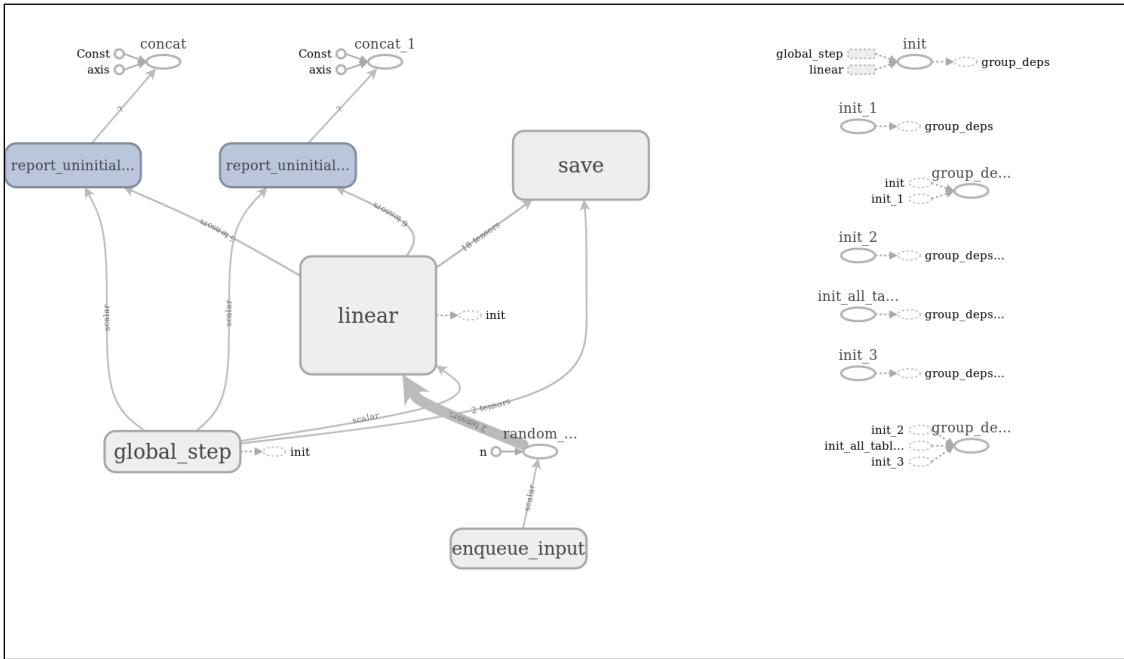
Chapter 3: Regression

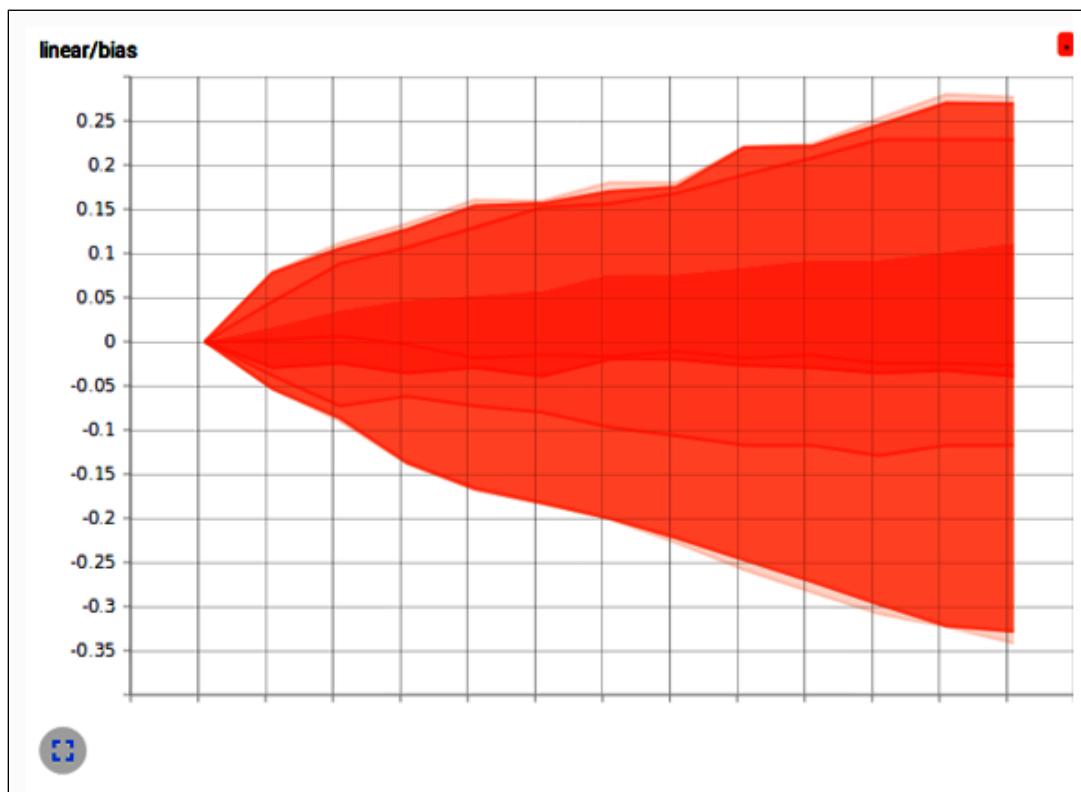


Predicted Value:	4.862152	Expected:	7.2
Predicted Value:	24.582247	Expected:	18.8
Predicted Value:	22.695276	Expected:	19.0
Predicted Value:	25.028057	Expected:	27.0
Predicted Value:	23.408998	Expected:	22.2
Predicted Value:	22.616102	Expected:	24.5
Predicted Value:	31.214731	Expected:	31.2
Predicted Value:	26.755243	Expected:	22.9
Predicted Value:	21.516464	Expected:	20.5
Predicted Value:	25.032785	Expected:	23.2
Predicted Value:	10.023388	Expected:	18.6
Predicted Value:	24.031082	Expected:	14.5
Predicted Value:	24.334019	Expected:	17.8
Predicted Value:	23.74925	Expected:	50.0
Predicted Value:	19.785368	Expected:	20.8
Predicted Value:	25.875463	Expected:	24.3
Predicted Value:	21.2129	Expected:	24.2
Predicted Value:	22.197586	Expected:	19.8
Predicted Value:	24.870373	Expected:	19.1
Predicted Value:	27.759129	Expected:	22.7
Predicted Value:	20.700903	Expected:	12.0
Predicted Value:	5.7440314	Expected:	10.2
Predicted Value:	22.404785	Expected:	20.0
Predicted Value:	25.772366	Expected:	18.5
Predicted Value:	33.465168	Expected:	20.9
Predicted Value:	25.10161	Expected:	23.0
Predicted Value:	26.143686	Expected:	27.5
Predicted Value:	35.51015	Expected:	30.1
Predicted Value:	8.041798	Expected:	9.5
Predicted Value:	24.381145	Expected:	22.0
Predicted Value:	24.351122	Expected:	21.2
Predicted Value:	9.700583	Expected:	14.1

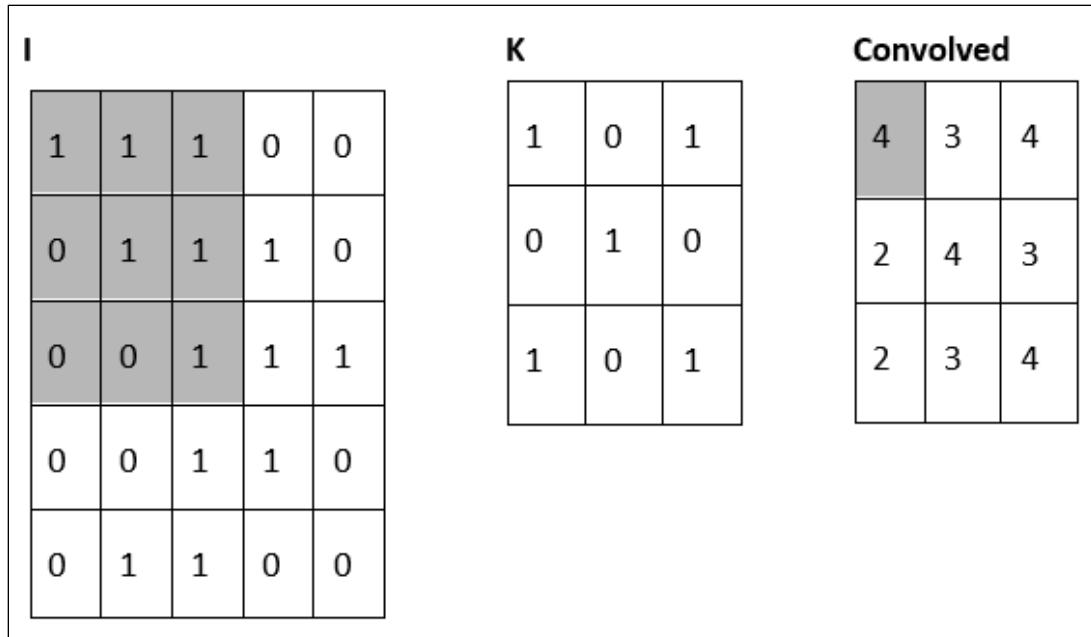
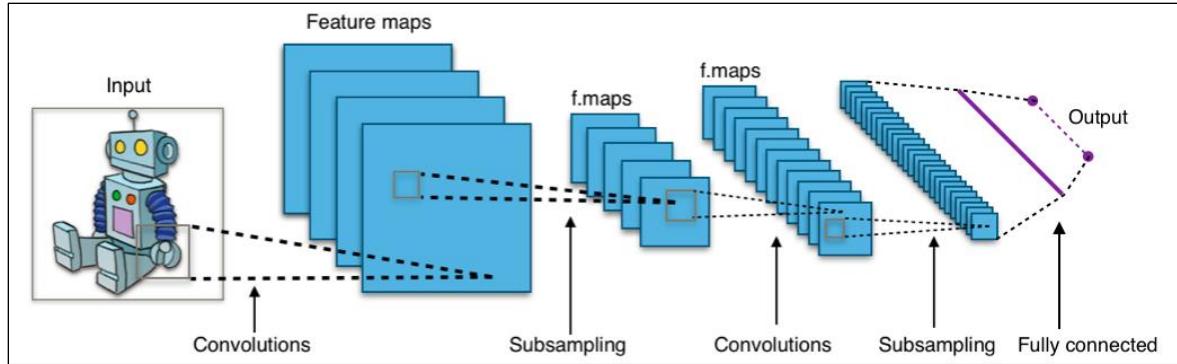


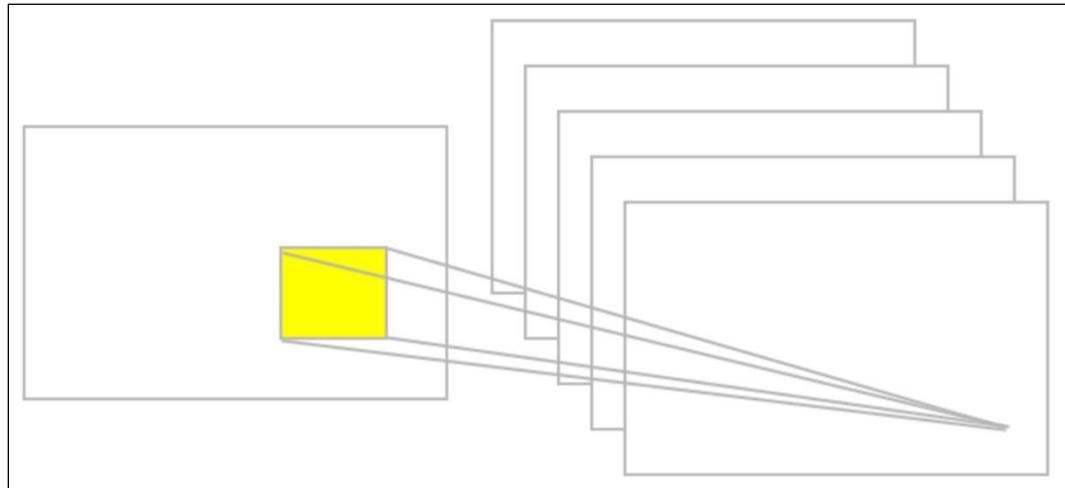






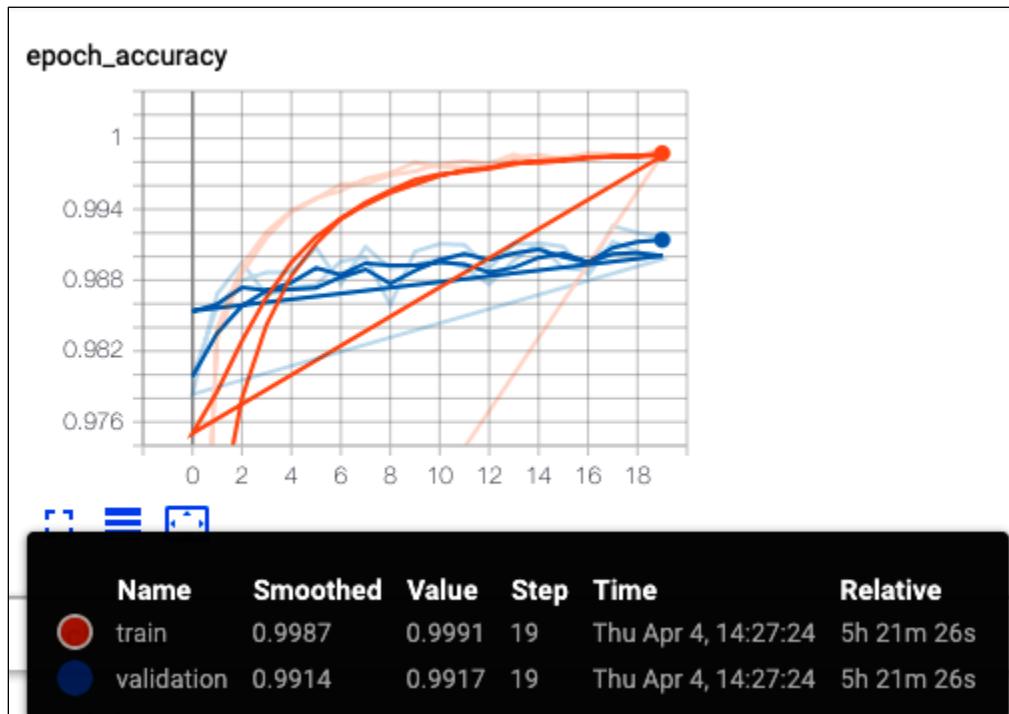
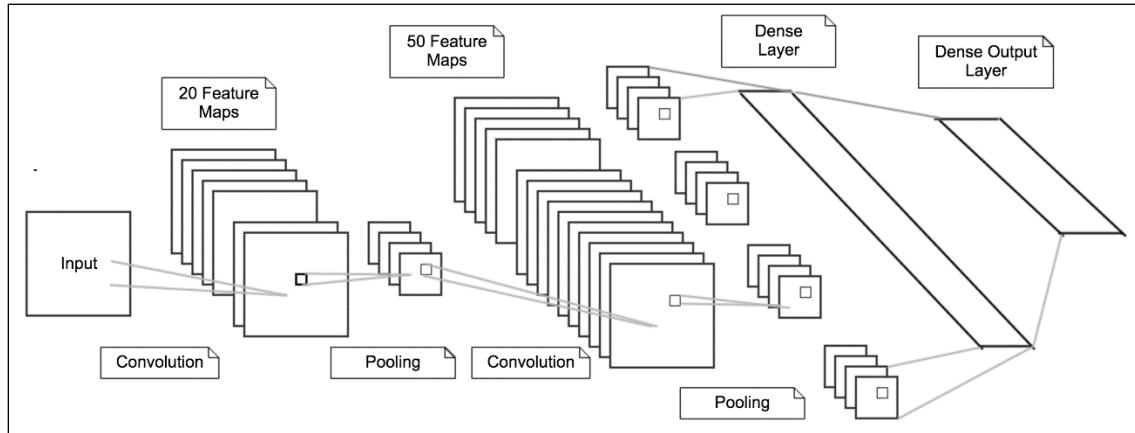
Chapter 4: Convolutional Neural Networks





1	0	3	6
2	4	5	2
2	6	2	0
3	4	1	7

A 4x4 input matrix is shown on the left, divided into four 2x2 blocks. The top-left block is yellow, the top-right is cyan, the bottom-left is green, and the bottom-right is orange. To its right is a smaller 2x2 output matrix. The top-left cell of the output matrix is yellow and contains the value 4. The top-right cell is cyan and contains 6. The bottom-left cell is green and contains 6. The bottom-right cell is orange and contains 7. This represents a stride of 2 convolution operation where the kernel size is 2x2.



```

Model: "sequential"
-----  

Layer (type)      Output Shape       Param #  

-----  

[conv2d (Conv2D)  (None, 24, 24, 20)    520  

max_pooling2d (MaxPooling2D) (None, 12, 12, 20)    0  

conv2d_1 (Conv2D)  (None, 8, 8, 50)     25050  

max_pooling2d_1 (MaxPooling2D) (None, 4, 4, 50)    0  

flatten (Flatten) (None, 800)          0  

dense (Dense)     (None, 500)          400500  

dense_1 (Dense)   (None, 10)           5010  

-----  

Total params: 431,080  

Trainable params: 431,080  

Non-trainable params: 0  

-----  

[Train on 48000 samples, validate on 12000 samples  

Epoch 1/20  

[2019-04-04 14:18:28.546158: I tensorflow/core/profiler/lib/profiler_session.cc:164] Profile Session started.  

48000/48000 [=====] - 28s 594us/sample - loss: 0.2035 - accuracy: 0.9398 - val_loss: 0.0739 - val_accuracy: 0.9783  

Epoch 2/20  

48000/48000 [=====] - 26s 534us/sample - loss: 0.0520 - accuracy: 0.9839 - val_loss: 0.0435 - val_accuracy: 0.9868  

Epoch 3/20  

48000/48000 [=====] - 27s 564us/sample - loss: 0.0343 - accuracy: 0.9893 - val_loss: 0.0365 - val_accuracy: 0.9895  

Epoch 4/20  

48000/48000 [=====] - 27s 562us/sample - loss: 0.0248 - accuracy: 0.9921 - val_loss: 0.0452 - val_accuracy: 0.9868  

Epoch 5/20  

48000/48000 [=====] - 27s 562us/sample - loss: 0.0195 - accuracy: 0.9939 - val_loss: 0.0428 - val_accuracy: 0.9873  

Epoch 6/20  

48000/48000 [=====] - 28s 588us/sample - loss: 0.0153 - accuracy: 0.9950 - val_loss: 0.0417 - val_accuracy: 0.9876  

Epoch 7/20  

48000/48000 [=====] - 26s 537us/sample - loss: 0.0134 - accuracy: 0.9955 - val_loss: 0.0388 - val_accuracy: 0.9896  

Epoch 8/20  

[48000/48000 [=====] - 29s 598us/sample - loss: 0.0097 - accuracy: 0.9966 - val_loss: 0.0347 - val_accuracy: 0.9899  

Epoch 9/20  

48000/48000 [=====] - 29s 607us/sample - loss: 0.0091 - accuracy: 0.9971 - val_loss: 0.0315 - val_accuracy: 0.9859  

Epoch 10/20  

48000/48000 [=====] - 27s 565us/sample - loss: 0.0062 - accuracy: 0.9980 - val_loss: 0.0376 - val_accuracy: 0.9904  

Epoch 11/20  

48000/48000 [=====] - 30s 627us/sample - loss: 0.0068 - accuracy: 0.9976 - val_loss: 0.0366 - val_accuracy: 0.9911  

Epoch 12/20  

48000/48000 [=====] - 24s 505us/sample - loss: 0.0079 - accuracy: 0.9975 - val_loss: 0.0389 - val_accuracy: 0.9918  

Epoch 13/20  

48000/48000 [=====] - 28s 584us/sample - loss: 0.0057 - accuracy: 0.9978 - val_loss: 0.0331 - val_accuracy: 0.9890  

Epoch 14/20  

48000/48000 [=====] - 28s 580us/sample - loss: 0.0045 - accuracy: 0.9984 - val_loss: 0.0409 - val_accuracy: 0.9911  

Epoch 15/20  

48000/48000 [=====] - 26s 537us/sample - loss: 0.0039 - accuracy: 0.9986 - val_loss: 0.0436 - val_accuracy: 0.9911  

Epoch 16/20  

48000/48000 [=====] - 25s 513us/sample - loss: 0.0059 - accuracy: 0.9983 - val_loss: 0.0480 - val_accuracy: 0.9890  

Epoch 17/20  

[48000/48000 [=====] - 24s 499us/sample - loss: 0.0042 - accuracy: 0.9988 - val_loss: 0.0535 - val_accuracy: 0.9888  

Epoch 18/20  

[48000/48000 [=====] - 24s 505us/sample - loss: 0.0042 - accuracy: 0.9986 - val_loss: 0.0349 - val_accuracy: 0.9926  

Epoch 19/20  

48000/48000 [=====] - 29s 599us/sample - loss: 0.0052 - accuracy: 0.9984 - val_loss: 0.0377 - val_accuracy: 0.9920  

Epoch 20/20  

48000/48000 [=====] - 25s 524us/sample - loss: 0.0028 - accuracy: 0.9991 - val_loss: 0.0477 - val_accuracy: 0.9917  

10000/10000 [=====] - 2s 240us/sample - loss: 0.0383 - accuracy: 0.9915  

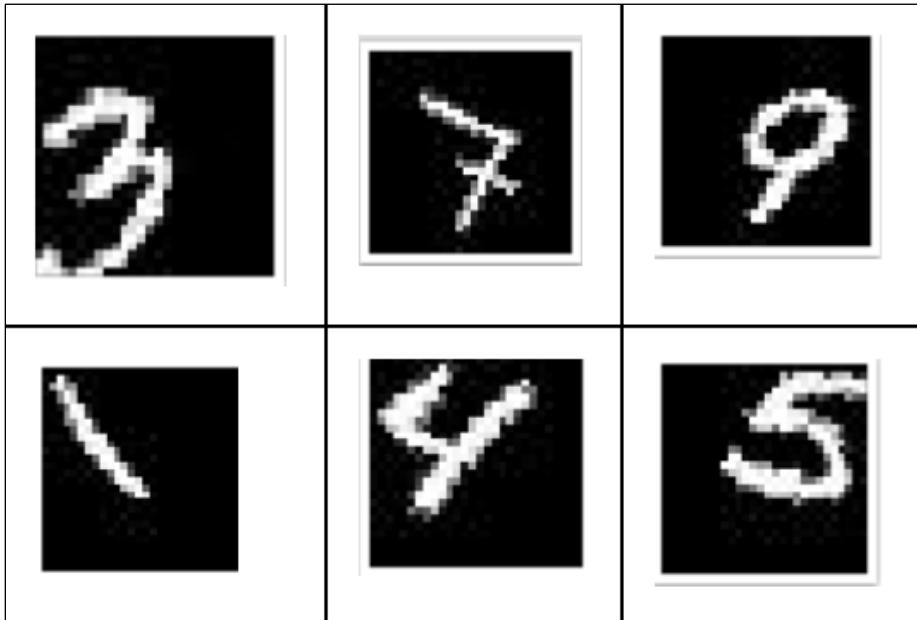
[  

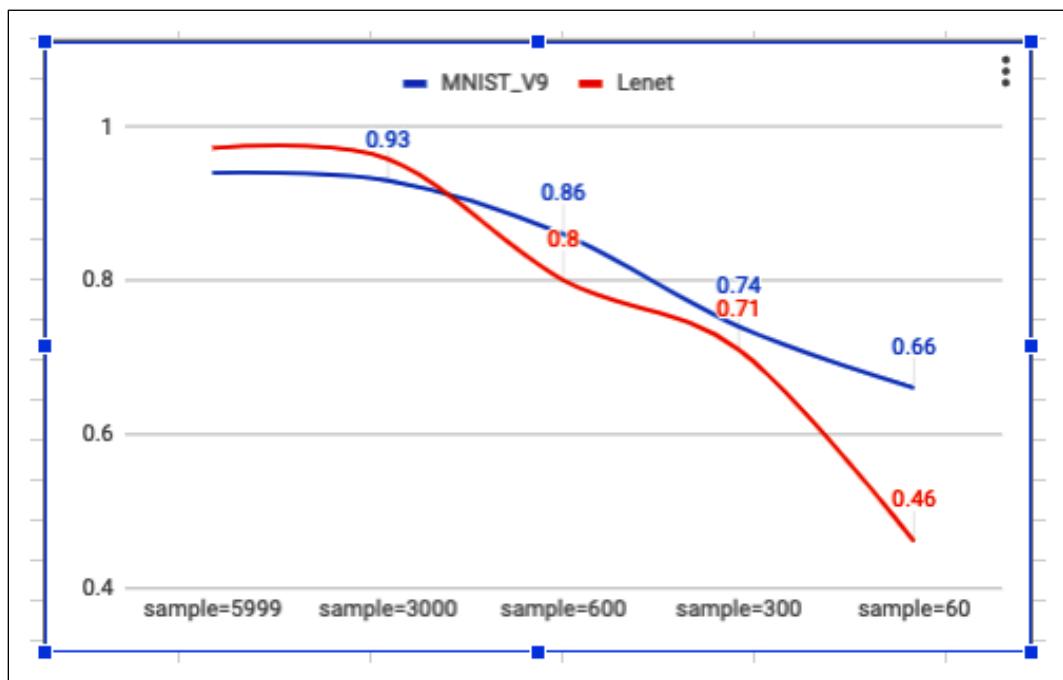
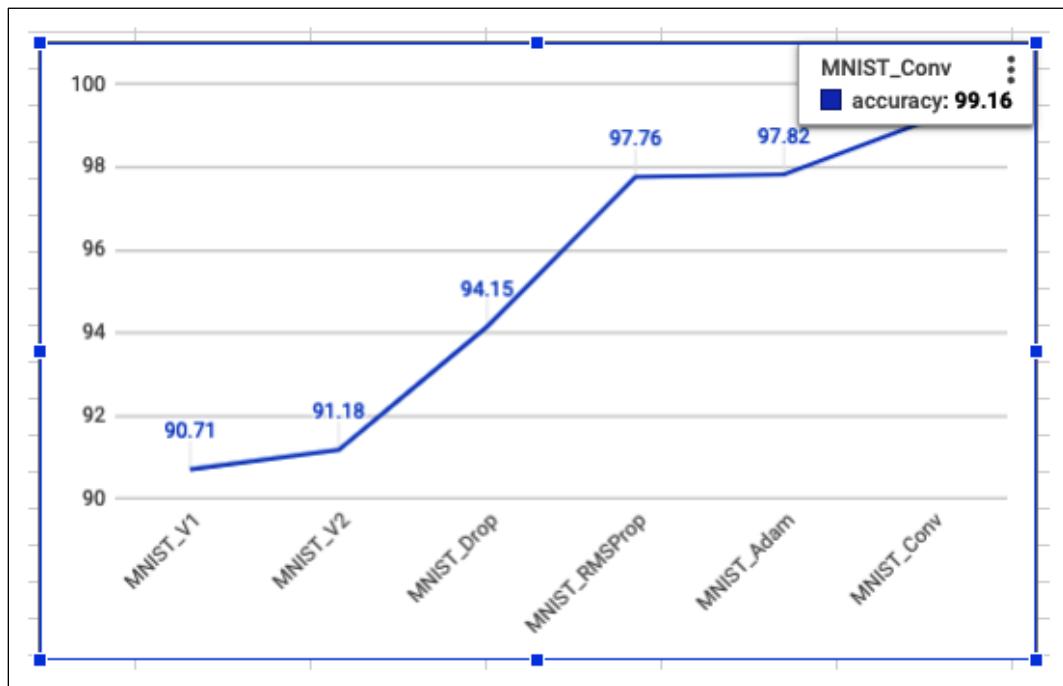
Test score: 0.03832608199457617  

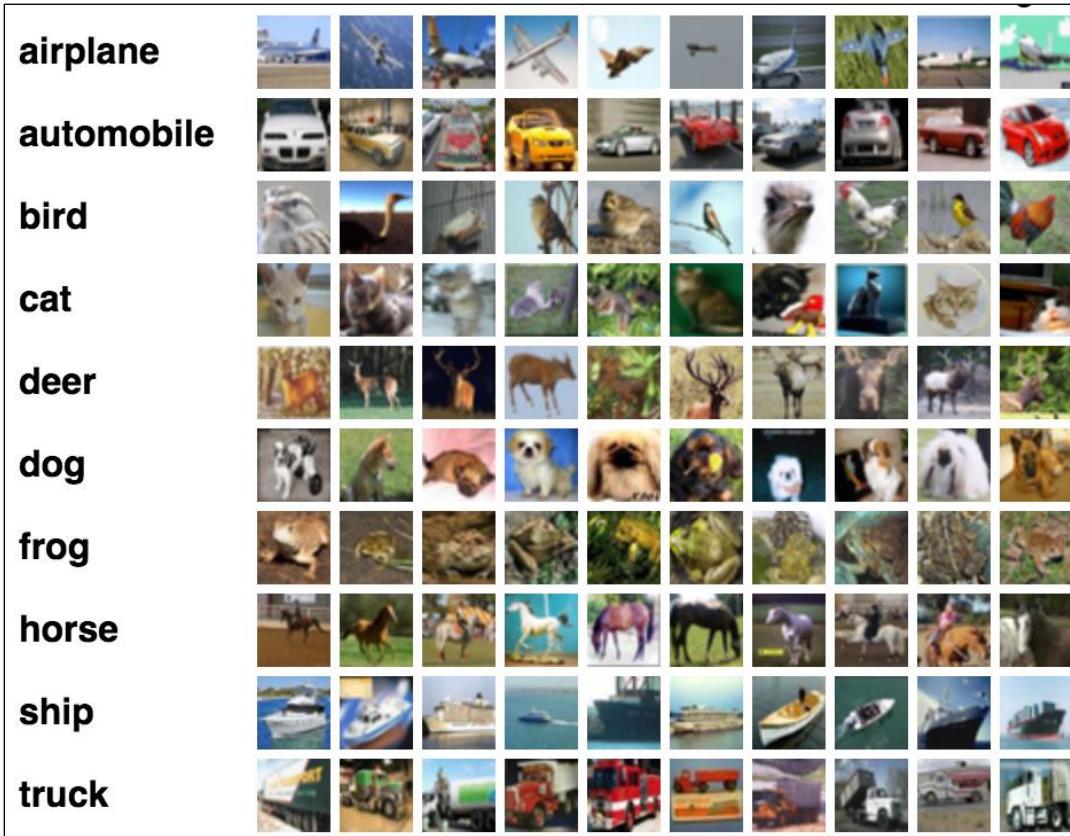
Test accuracy: 0.9915

```

```
-----  
Train on 48000 samples, validate on 12000 samples  
Epoch 1/10  
2019-04-04 15:57:17.848186: I tensorflow/core/profiler/lib/profiler_session.cc:164] Profile Session started.  
48000/48000 [=====] - 26s 544us/sample - loss: 0.2134 - accuracy: 0.9361 - val_loss: 0.0688 - val_accuracy: 0.9783  
Epoch 2/10  
48000/48000 [=====] - 30s 633us/sample - loss: 0.0550 - accuracy: 0.9831 - val_loss: 0.0533 - val_accuracy: 0.9843  
Epoch 3/10  
48000/48000 [=====] - 38s 621us/sample - loss: 0.0353 - accuracy: 0.9884 - val_loss: 0.0410 - val_accuracy: 0.9874  
Epoch 4/10  
48000/48000 [=====] - 37s 767us/sample - loss: 0.0276 - accuracy: 0.9910 - val_loss: 0.0381 - val_accuracy: 0.9887  
Epoch 5/10  
48000/48000 [=====] - 24s 509us/sample - loss: 0.0200 - accuracy: 0.9932 - val_loss: 0.0406 - val_accuracy: 0.9881  
Epoch 6/10  
48000/48000 [=====] - 31s 641us/sample - loss: 0.0161 - accuracy: 0.9950 - val_loss: 0.0423 - val_accuracy: 0.9881  
Epoch 7/10  
48000/48000 [=====] - 29s 613us/sample - loss: 0.0129 - accuracy: 0.9955 - val_loss: 0.0396 - val_accuracy: 0.9894  
Epoch 8/10  
48000/48000 [=====] - 27s 554us/sample - loss: 0.0107 - accuracy: 0.9965 - val_loss: 0.0454 - val_accuracy: 0.9871  
Epoch 9/10  
48000/48000 [=====] - 24s 510us/sample - loss: 0.0082 - accuracy: 0.9973 - val_loss: 0.0388 - val_accuracy: 0.9902  
Epoch 10/10  
48000/48000 [=====] - 26s 542us/sample - loss: 0.0083 - accuracy: 0.9970 - val_loss: 0.0440 - val_accuracy: 0.9892  
10000/10000 [=====] - 2s 196us/sample - loss: 0.0327 - accuracy: 0.9910  
  
Test score: 0.03265062951518773  
Test accuracy: 0.991
```





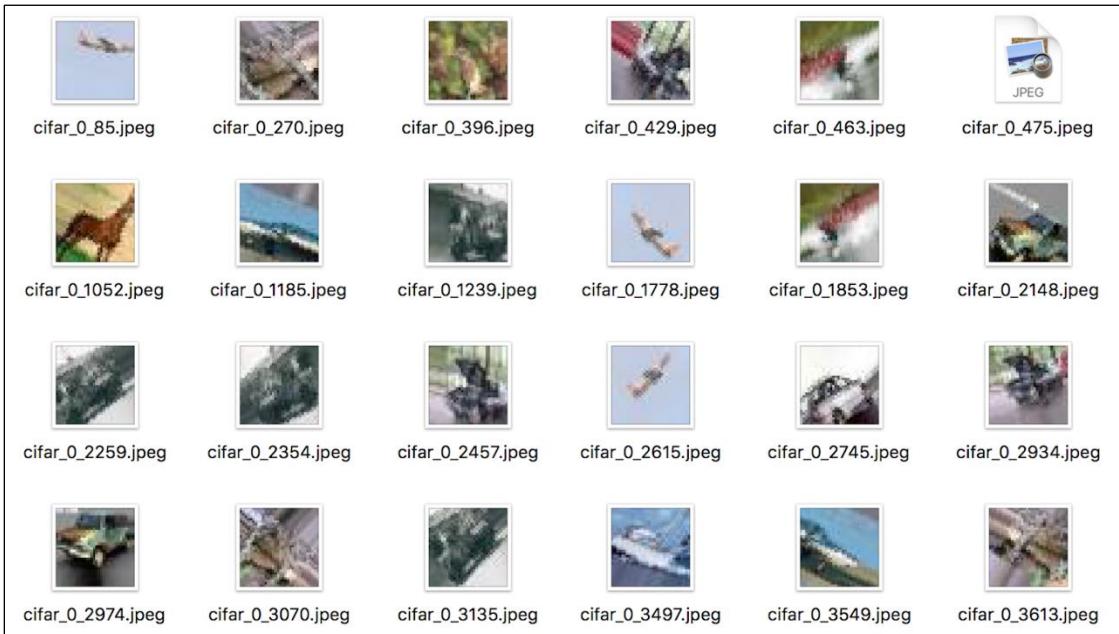
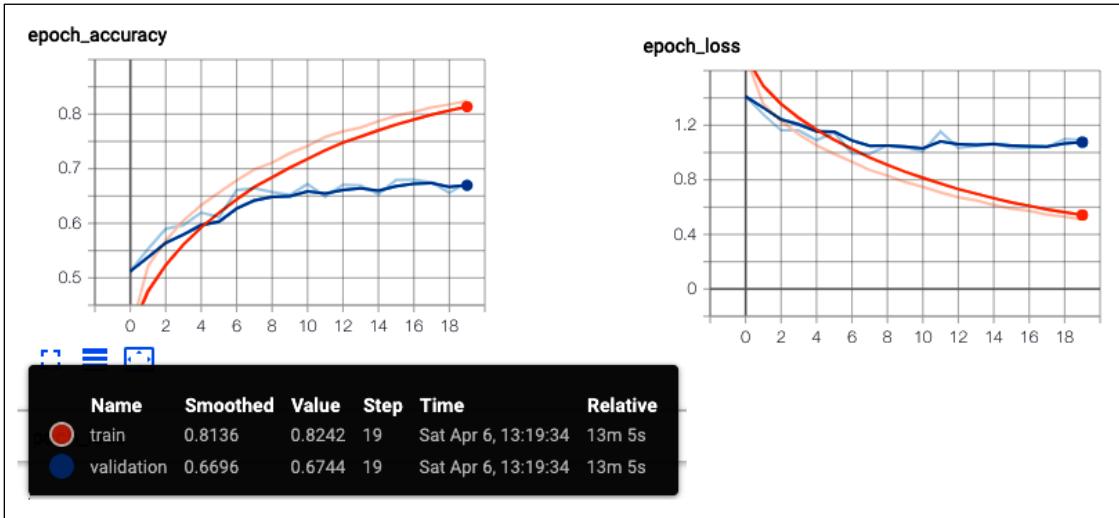


```

Epoch 17/20
40000/40000 [=====] - 112s 3ms/sample - loss: 0.6282 - accuracy: 0.7841 - val_loss: 1.0296 -
val_accuracy: 0.6734
Epoch 18/20
40000/40000 [=====] - 76s 2ms/sample - loss: 0.6140 - accuracy: 0.7879 - val_loss: 1.0789 -
val_accuracy: 0.6489
Epoch 19/20
40000/40000 [=====] - 74s 2ms/sample - loss: 0.5931 - accuracy: 0.7958 - val_loss: 1.0461 -
val_accuracy: 0.6811
Epoch 20/20
40000/40000 [=====] - 71s 2ms/sample - loss: 0.5724 - accuracy: 0.8042 - val_loss: 1.0527 -
val_accuracy: 0.6773
10000/10000 [=====] - 5s 472us/sample - loss: 1.0423 - accuracy: 0.6686

Test score: 1.0423416819572449
Test accuracy: 0.6686

```

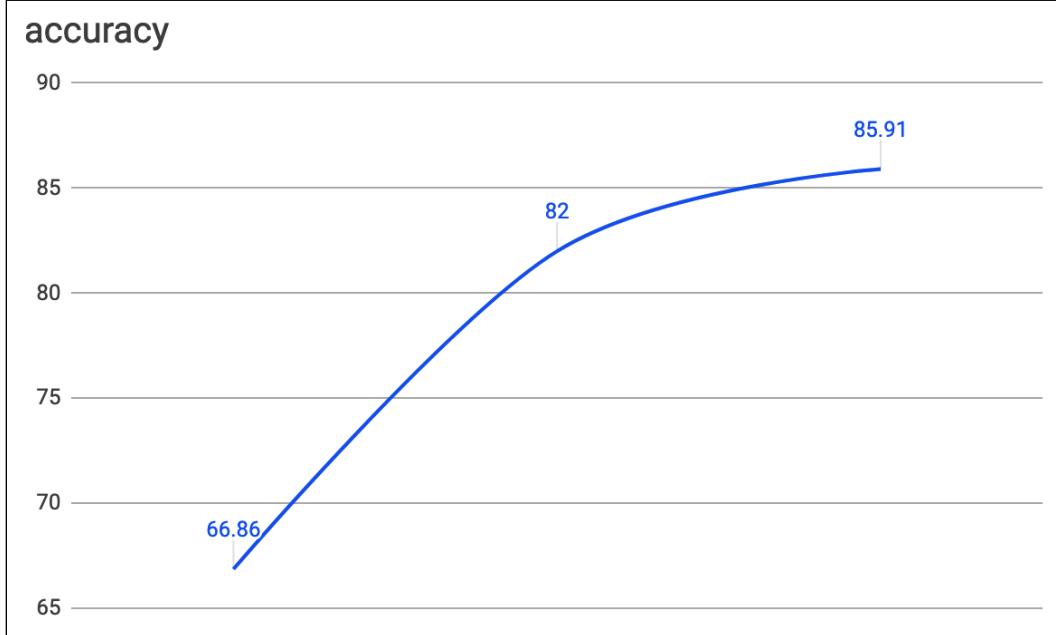


```

Epoch 46/50
50000/50000 [=====] - 36s 722us/sample - loss: 0.2440 - acc: 0.9183 - val_loss: 0.4918 - val_acc: 0.8546
Epoch 47/50
50000/50000 [=====] - 34s 685us/sample - loss: 0.2338 - acc: 0.9208 - val_loss: 0.4884 - val_acc: 0.8574
Epoch 48/50
50000/50000 [=====] - 32s 643us/sample - loss: 0.2383 - acc: 0.9189 - val_loss: 0.5106 - val_acc: 0.8556
Epoch 49/50
50000/50000 [=====] - 37s 734us/sample - loss: 0.2285 - acc: 0.9212 - val_loss: 0.5017 - val_acc: 0.8581
Epoch 50/50
50000/50000 [=====] - 36s 712us/sample - loss: 0.2263 - acc: 0.9228 - val_loss: 0.4911 - val_acc: 0.8591

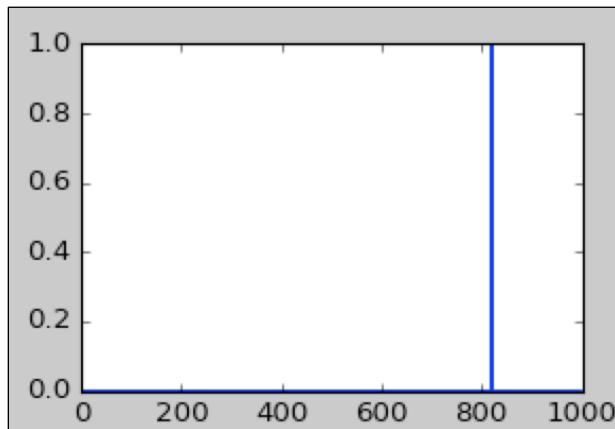
```

```
10000/10000 [=====] - 2s 160us/sample - loss: 0.4911 - acc: 0.8591
Test score: 0.4911323667049408
Test accuracy: 0.8591
```



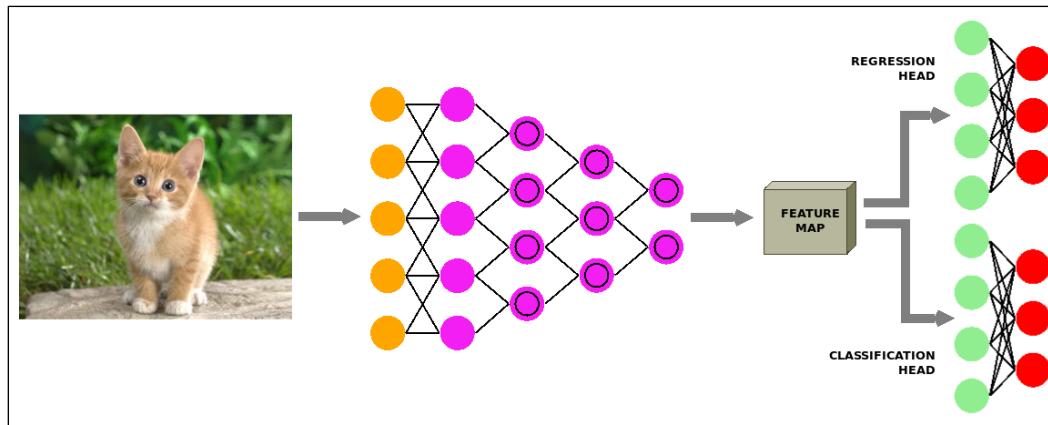
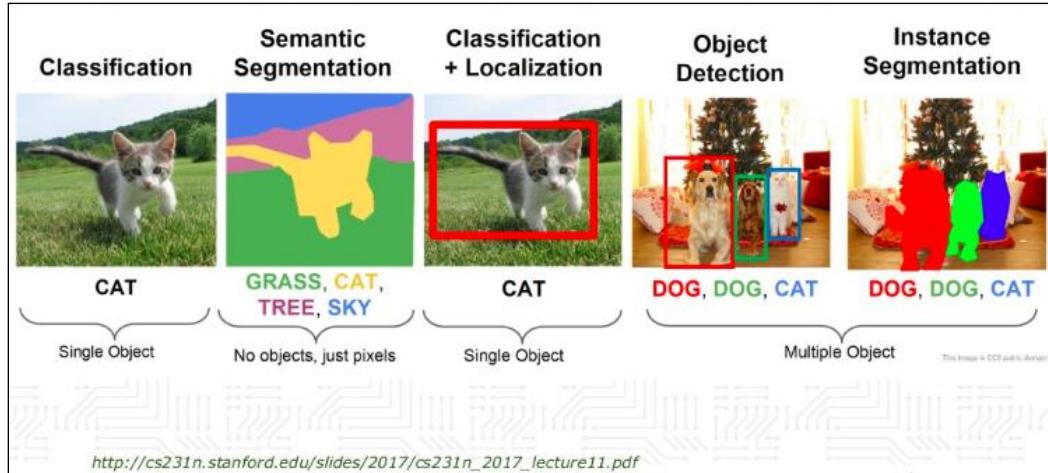
```
Total params: 138,357,544
Trainable params: 138,357,544
Non-trainable params: 0
```

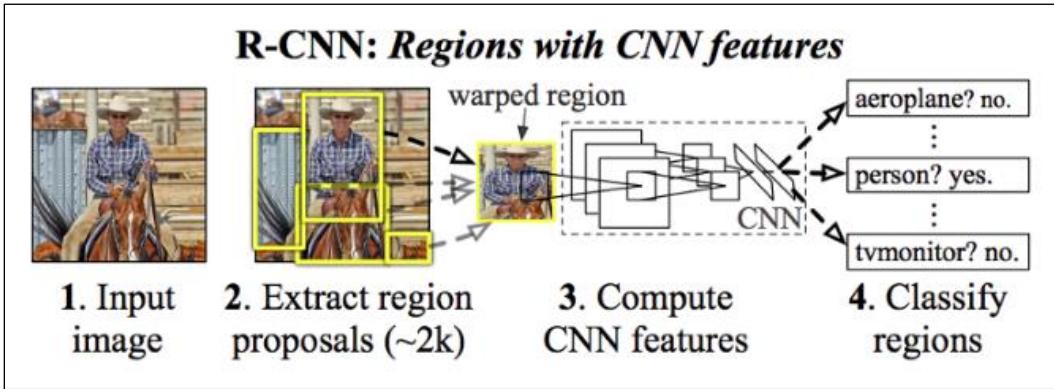
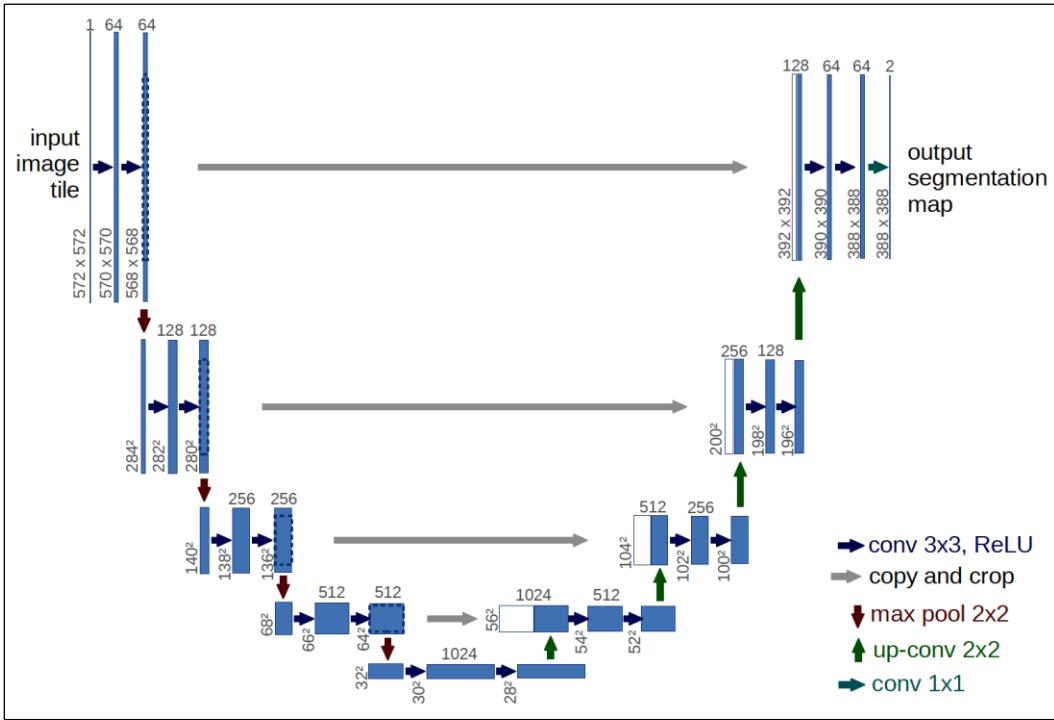
```
-----  
285
```

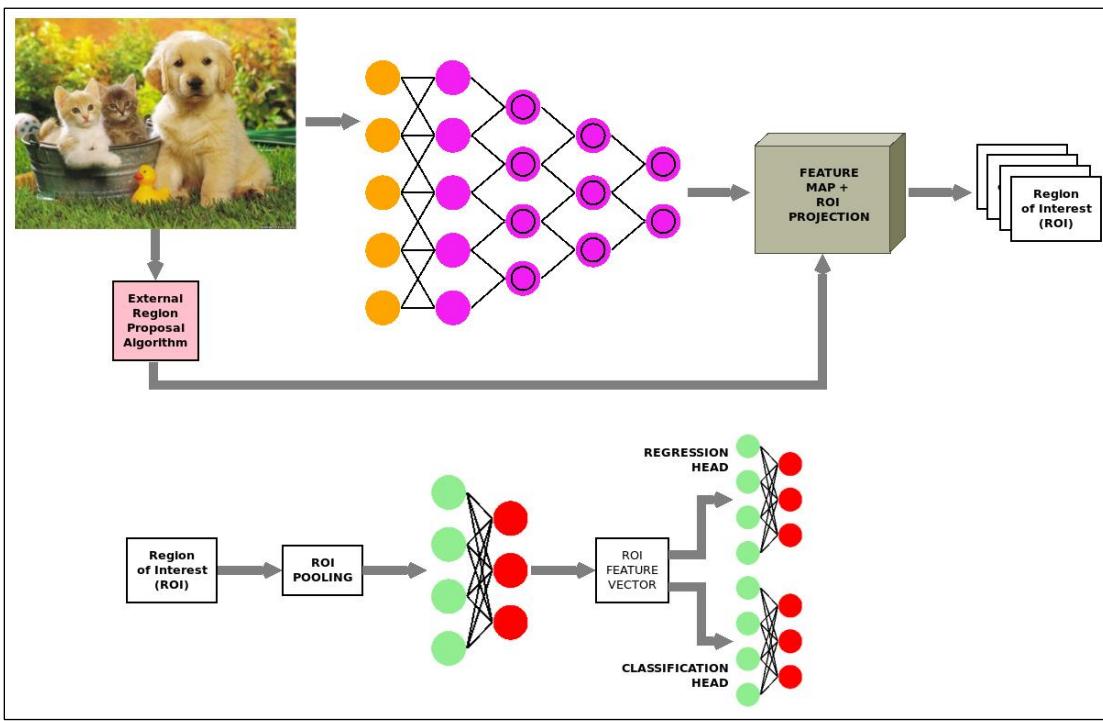
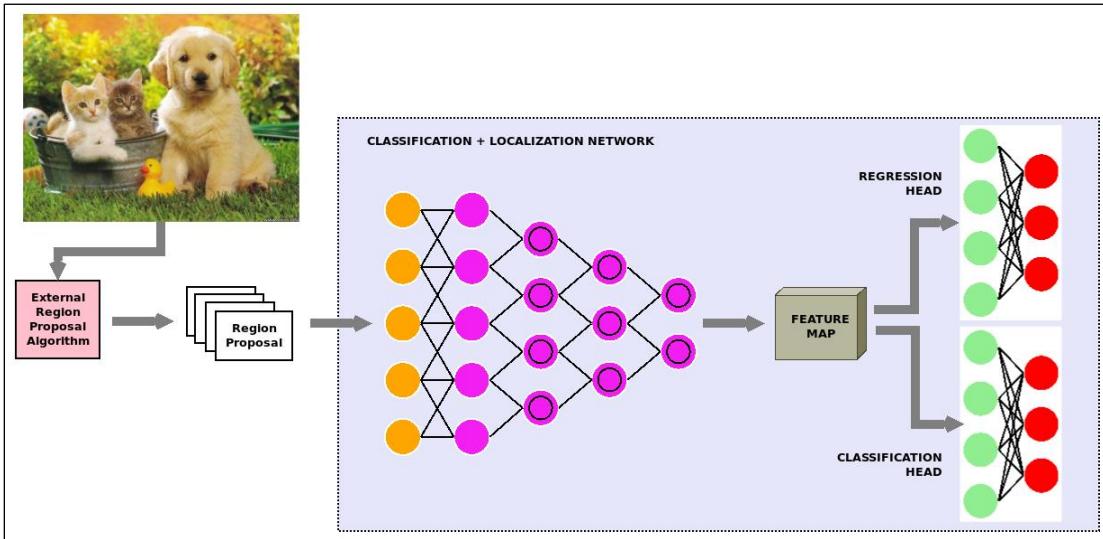


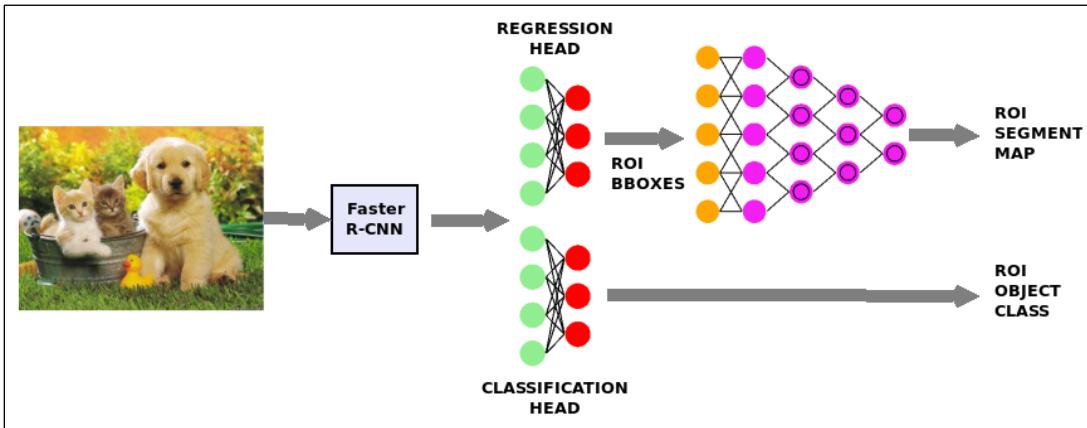
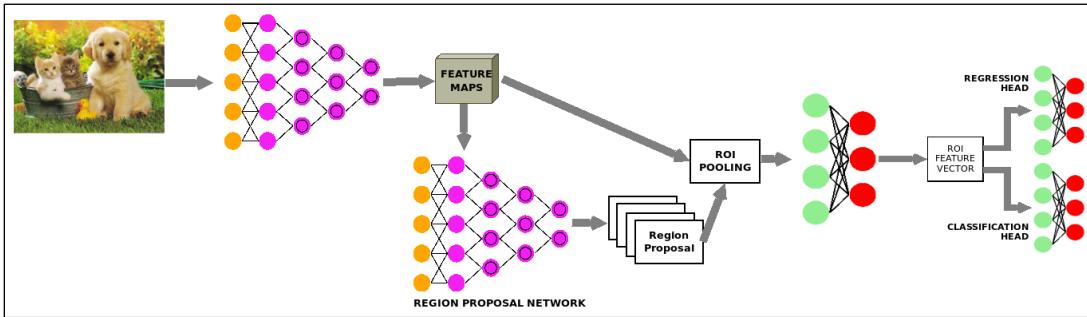


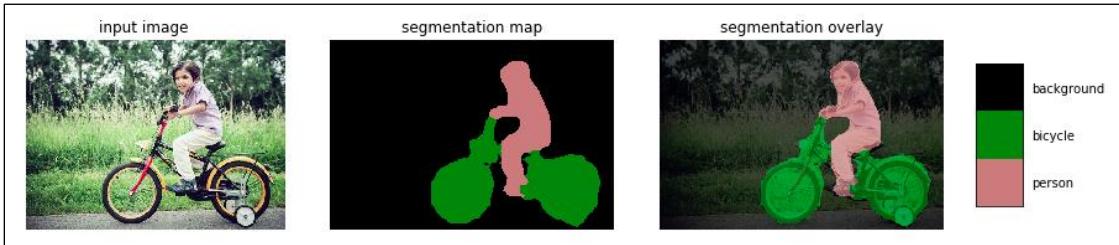
Chapter 5: Advanced Convolutional Neural Networks











```

    estimator.evaluate(lambda:input_fn(test_images,
                                         test_labels,
                                         epochs=1,
                                         batch_size=BATCH_SIZE))

    {'accuracy': 0.7162, 'global_step': 5860, 'loss': 0.77385104}

```

```

[8] #strategy = None
strategy = tf.distribute.MirroredStrategy()
config = tf.estimator.RunConfig(train_distribute=strategy)

```

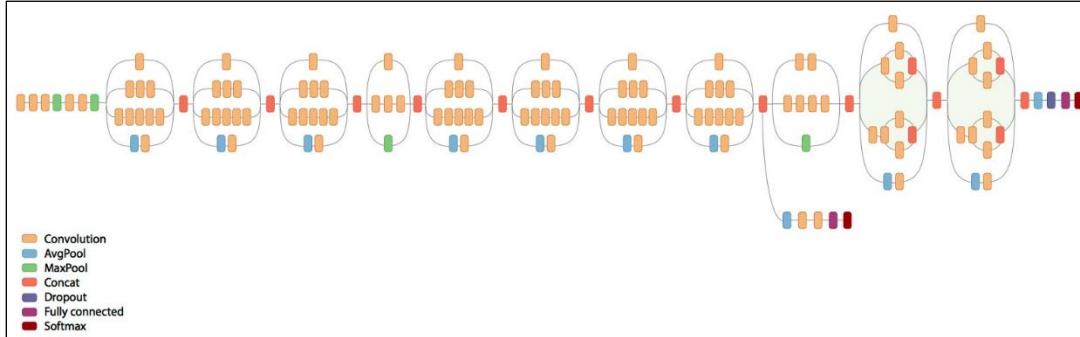
```
BATCH_SIZE = 512
EPOCHS = 50

#time_hist = TimeHistory()

estimator_train_result = estimator.train(input_fn=lambda:input_fn(train_images,
                                                               train_labels,
                                                               epochs=EPOCHS,
                                                               batch_size=BATCH_SIZE))
print(estimator_train_result)
```

```
[12] estimator.evaluate(lambda:input_fn(test_images,
                                         test_labels,
                                         epochs=1,
                                         batch_size=BATCH_SIZE))

[?] {'acc': 0.8215, 'global_step': 5860, 'loss': 0.48483768}
```



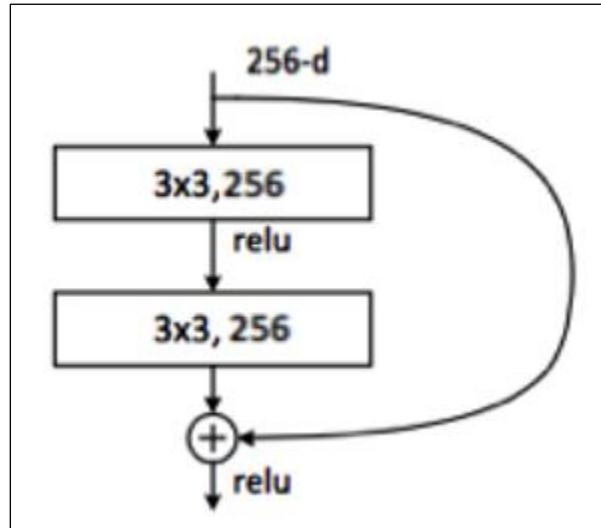
```
Downloading data from https://github.com/fchollet/deep-learning-models/releases/download/v0.5/inception_v3_weights_tf_dim_ordering_tf_kernels_notop.h5
87916544/87910968 [=====] - 26s 0us/step
```

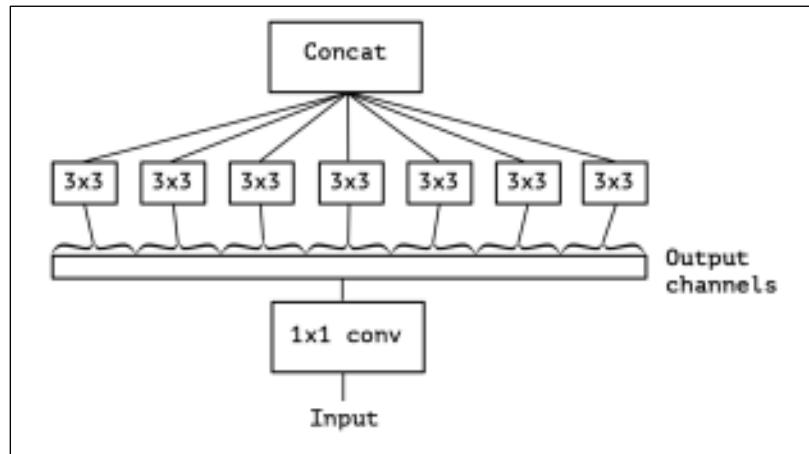
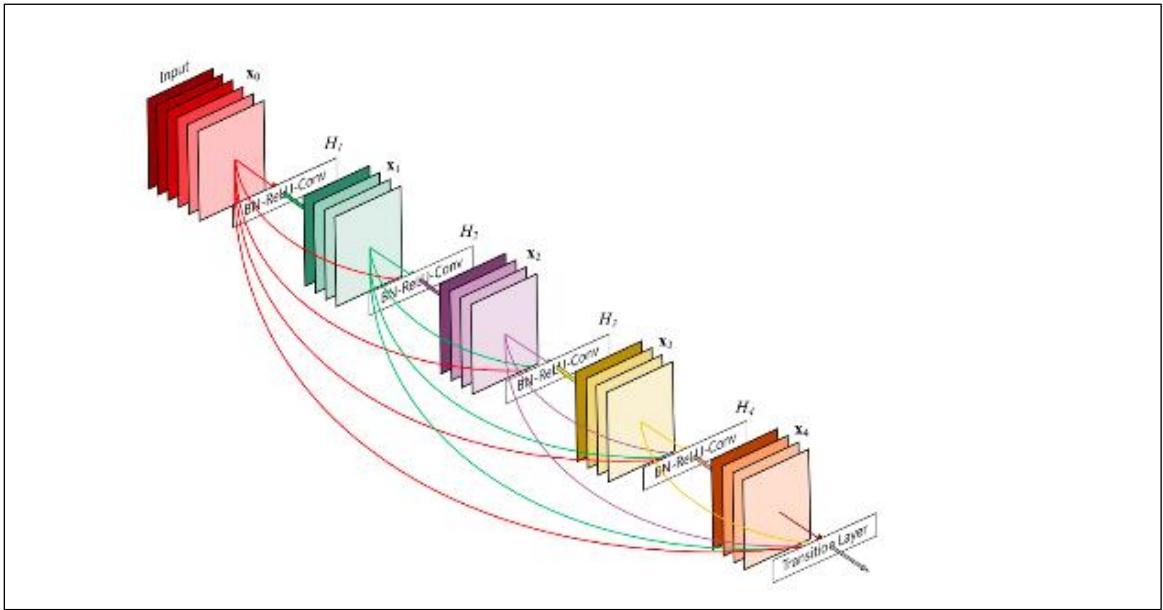
```
model.summary()
```

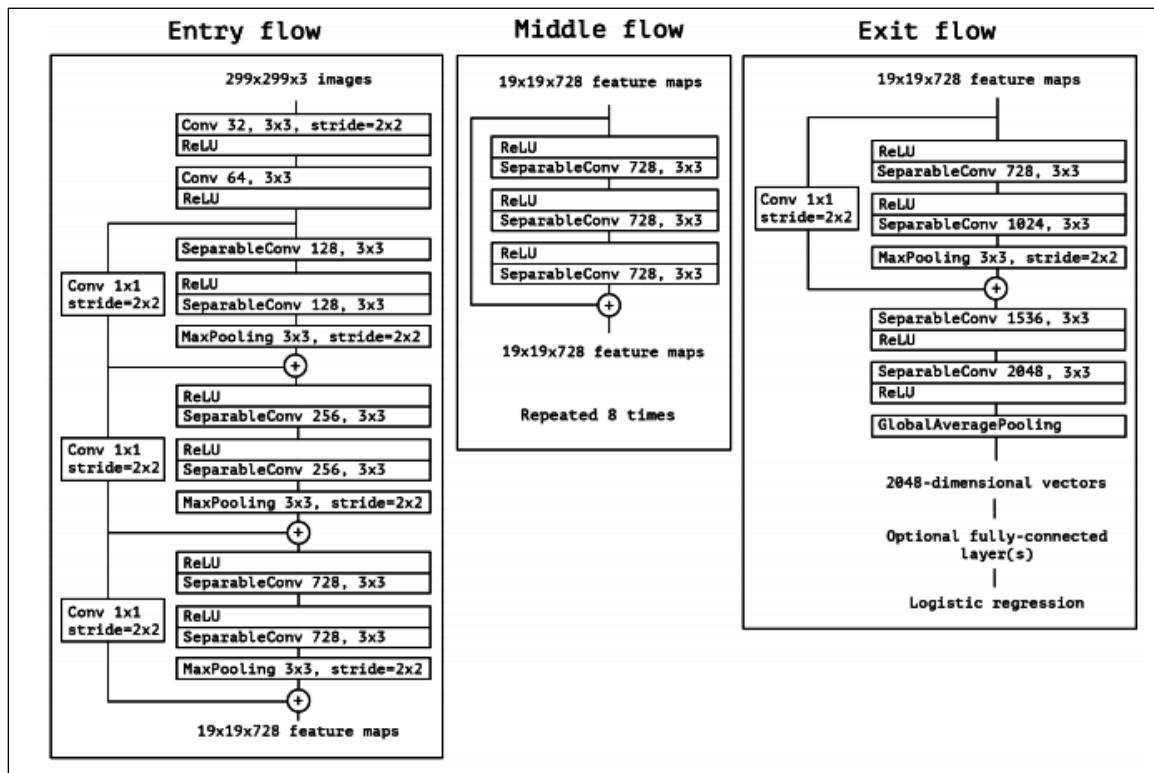
```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
mobilenetv2_1.00_160 (Model)	(None, 5, 5, 1280)	2257984
global_average_pooling2d (G1)	(None, 1280)	0
dense (Dense)	(None, 1)	1281
Total params:	2,259,265	
Trainable params:	1,281	
Non-trainable params:	2,257,984	

```
Epoch 18/20  
26/26 [=====] - 5s 198ms/step - loss: 0.1675 - accuracy: 0.9661 - val_loss: 0.0451 - val_accuracy: 0.9800  
Epoch 19/20  
26/26 [=====] - 6s 223ms/step - loss: 0.1222 - accuracy: 0.9722 - val_loss: 0.0381 - val_accuracy: 0.9800  
Epoch 20/20  
26/26 [=====] - 6s 225ms/step - loss: 0.1087 - accuracy: 0.9807 - val_loss: 0.0359 - val_accuracy: 0.9800
```







Model	Size	Top-1 Accuracy	Top-5 Accuracy	Parameters	Depth
Xception	88 MB	0.790	0.945	22,910,480	126
VGG16	528 MB	0.713	0.901	138,357,544	23
VGG19	549 MB	0.713	0.900	143,667,240	26
ResNet50	98 MB	0.749	0.921	25,636,712	-
ResNet101	171 MB	0.764	0.928	44,707,176	-
ResNet152	232 MB	0.766	0.931	60,419,944	-
ResNet50V2	98 MB	0.760	0.930	25,613,800	-
ResNet101V2	171 MB	0.772	0.938	44,675,560	-
ResNet152V2	232 MB	0.780	0.942	60,380,648	-
ResNeXt50	96 MB	0.777	0.938	25,097,128	-
ResNeXt101	170 MB	0.787	0.943	44,315,560	-
InceptionV3	92 MB	0.779	0.937	23,851,784	159
InceptionResNetV2	215 MB	0.803	0.953	55,873,736	572
MobileNet	16 MB	0.704	0.895	4,253,864	88
MobileNetV2	14 MB	0.713	0.901	3,538,984	88
DenseNet121	33 MB	0.750	0.923	8,062,504	121
DenseNet169	57 MB	0.762	0.932	14,307,880	169
DenseNet201	80 MB	0.773	0.936	20,242,984	201
NASNetMobile	23 MB	0.744	0.919	5,326,716	-
NASNetLarge	343 MB	0.825	0.960	88,949,818	-

The top-1 and top-5 accuracy refers to the model's performance on the ImageNet validation dataset.





how many umbrellas are here Submit

Predicted top-5 answers with confidence:

2	54.694%
3	26.443%
1	13.868%
4	3.647%
5	0.873%



is it day or night? Submit

Predicted top-5 answers with confidence:

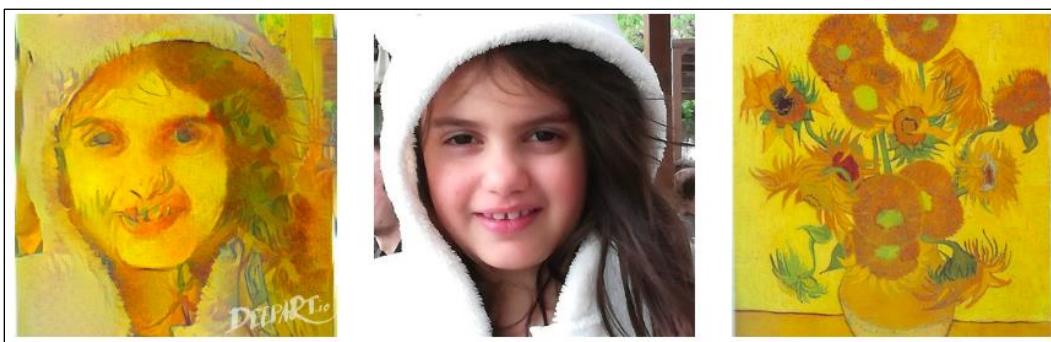
day	97.569%
night	1.403%
afternoon	0.011%
morning	0.002%
daytime	0.001%



what is this Submit

Predicted top-5 answers with confidence:

surfboard	57.696%
frisbee	11.463%
plane	1.398%
airplane	1.332%
boat	1.687%







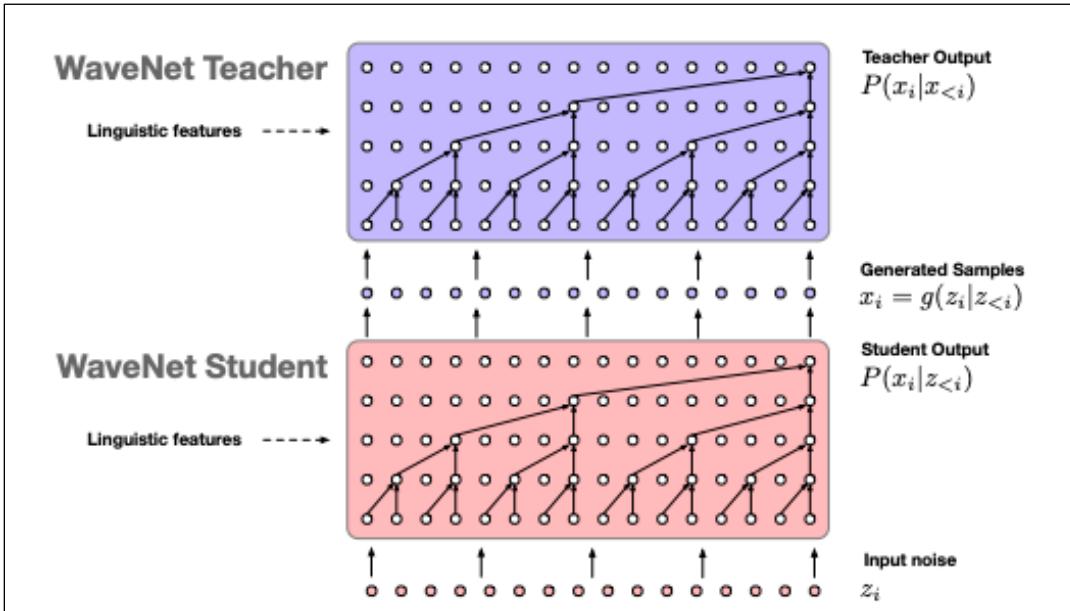
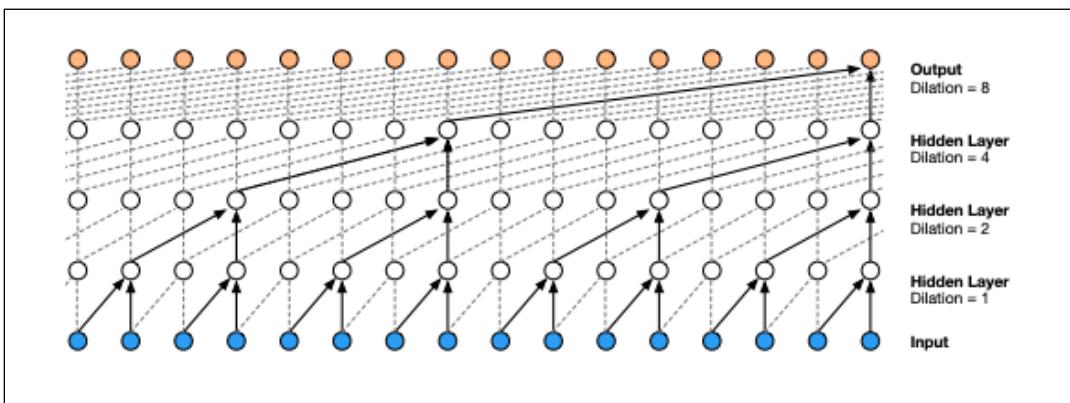
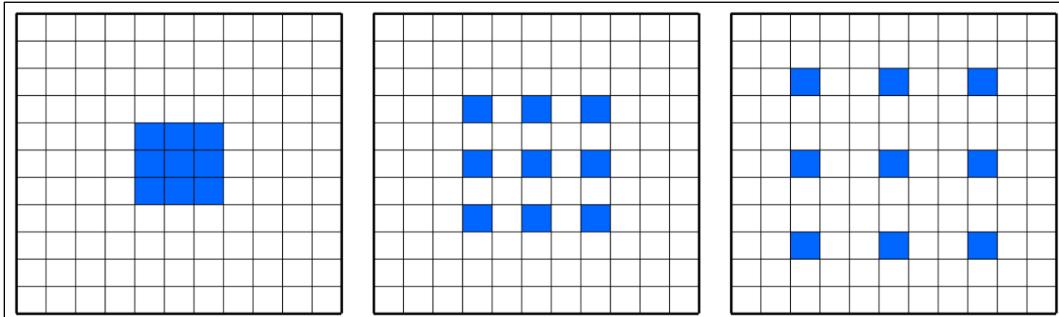
Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 200, 256)	2560000
dropout (Dropout)	(None, 200, 256)	0
conv1d (Conv1D)	(None, 198, 256)	196864
global_max_pooling1d (Global)	(None, 256)	0
dense (Dense)	(None, 128)	32896
dropout_1 (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 1)	129
<hr/>		
Total params: 2,789,889		
Trainable params: 2,789,889		
Non-trainable params: 0		

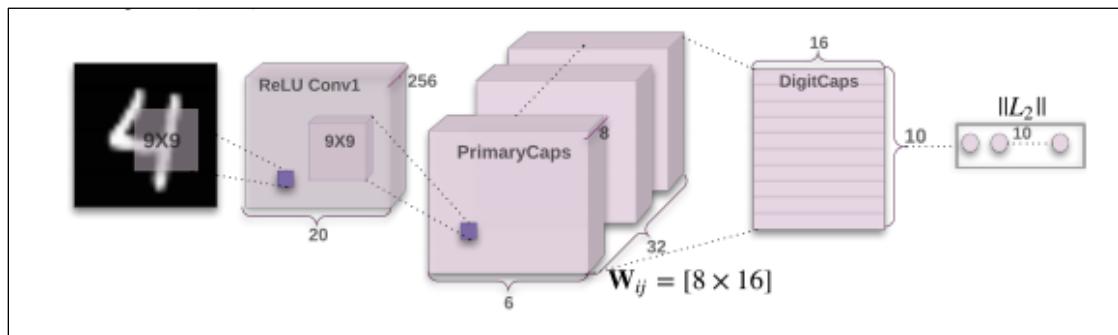
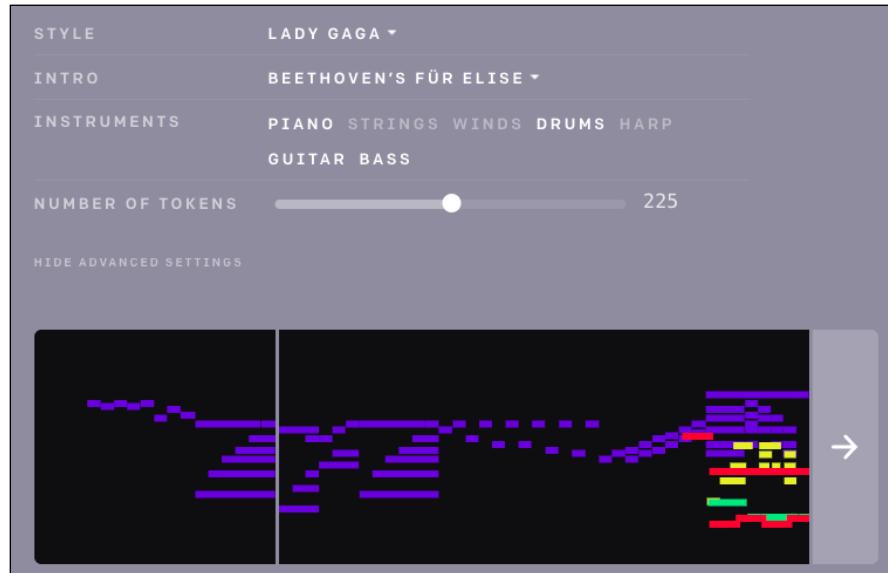
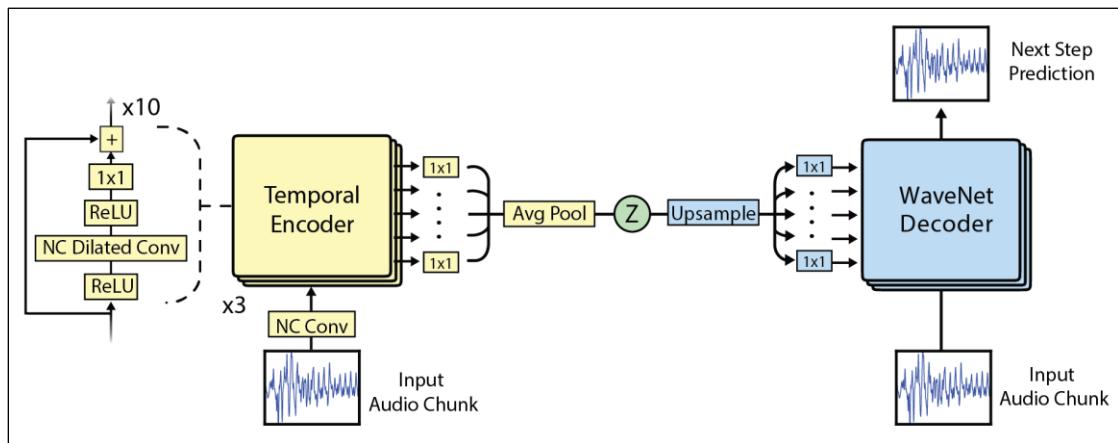
```

Epoch 19/20
25000/25000 [=====] - 135s 5ms/sample - loss: 7.5276e-04 - accuracy: 1.0000 - val_loss: 0.5753 - val_accuracy: 0.8818
Epoch 20/20
25000/25000 [=====] - 129s 5ms/sample - loss: 6.7755e-04 - accuracy: 0.9999 - val_loss: 0.5802 - val_accuracy: 0.8821
25000/25000 [=====] - 23s 916us/sample - loss: 0.5802 - accuracy: 0.8821

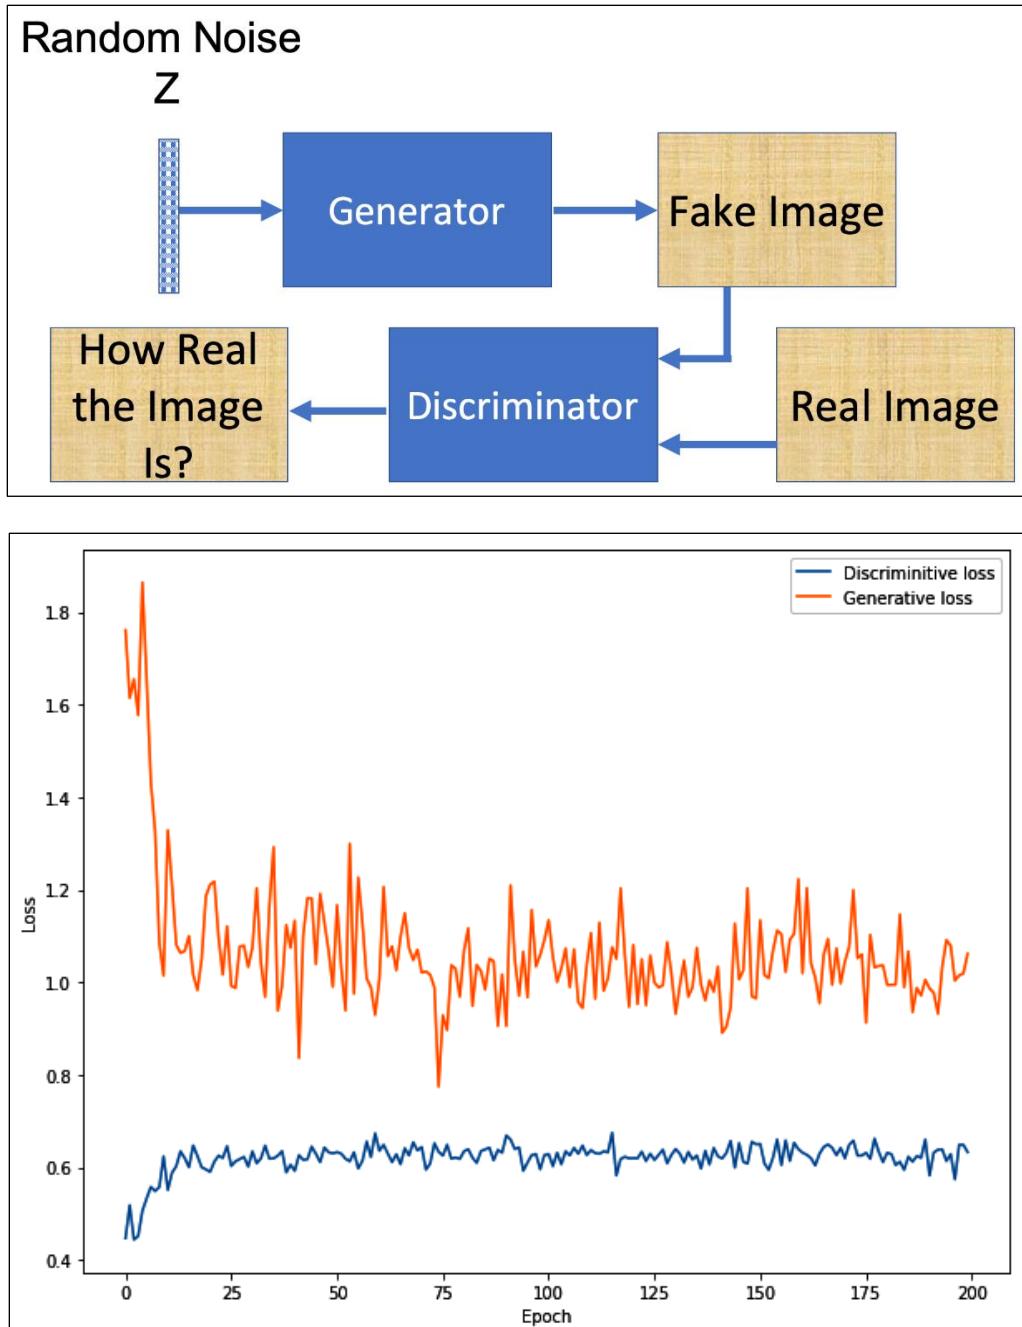
Test score: 0.5801781857013703
Test accuracy: 0.88212

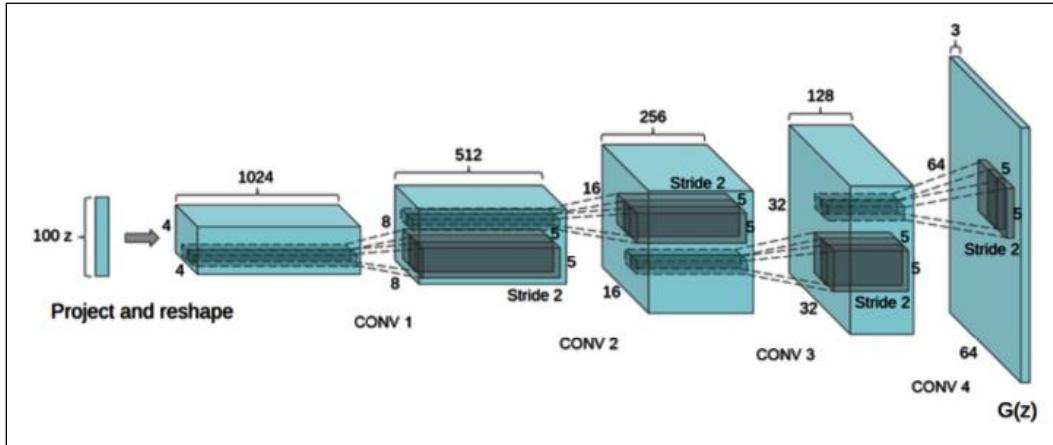
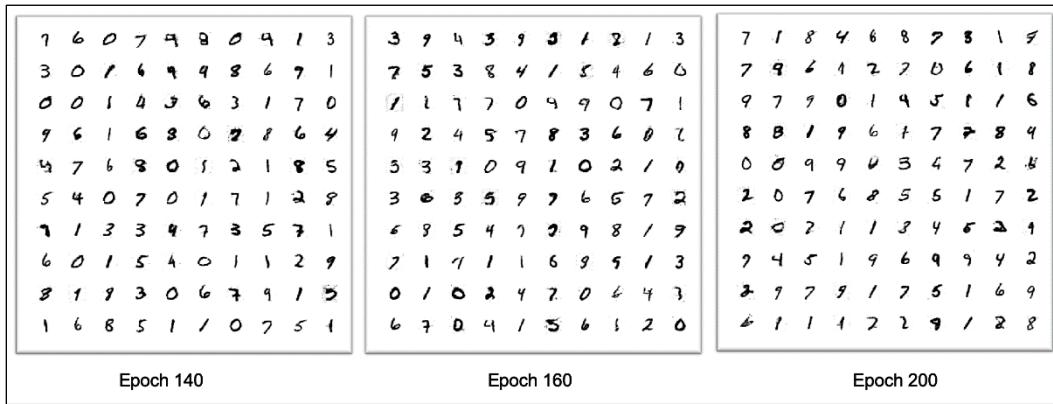
```

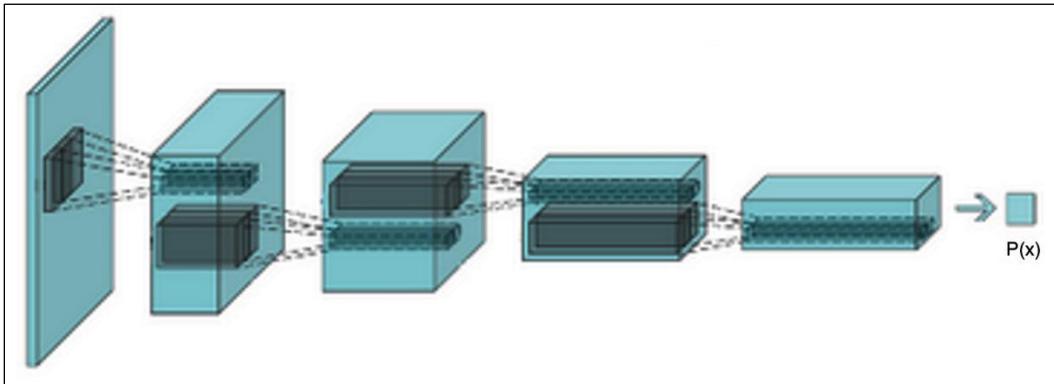




Chapter 6: Generative Adversarial Networks





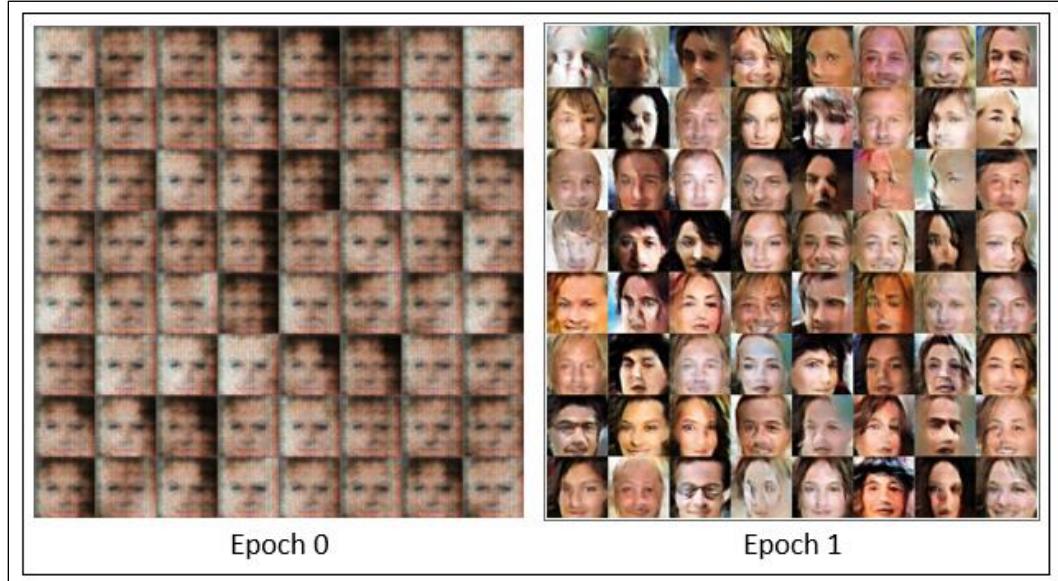
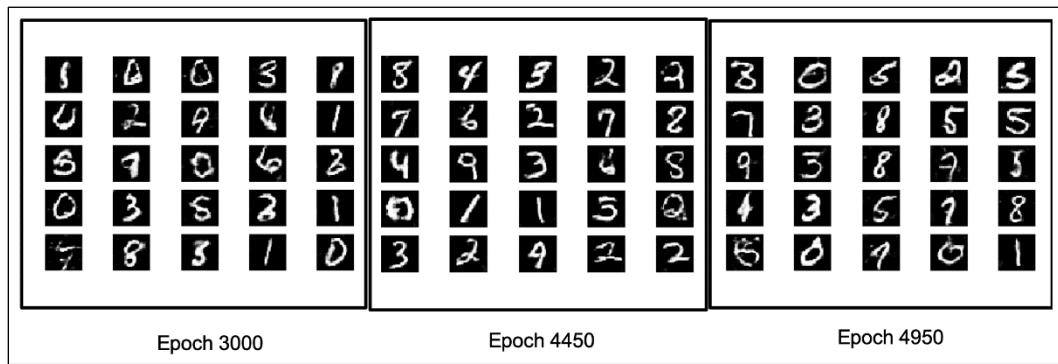
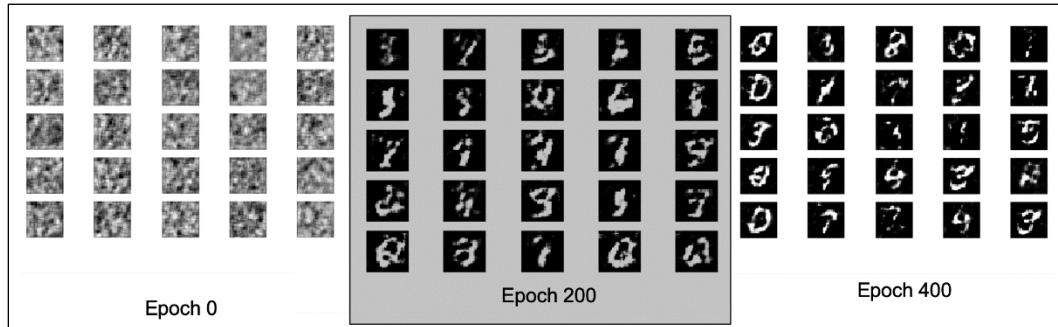


Model: "sequential_1"

Layer (type)	Output Shape	Param #
<hr/>		
dense_1 (Dense)	(None, 6272)	633472
reshape (Reshape)	(None, 7, 7, 128)	0
up_sampling2d (UpSampling2D)	(None, 14, 14, 128)	0
conv2d_4 (Conv2D)	(None, 14, 14, 128)	147584
batch_normalization_v2_3 (BatchNormalization)	(None, 14, 14, 128)	512
activation (Activation)	(None, 14, 14, 128)	0
up_sampling2d_1 (UpSampling2D)	(None, 28, 28, 128)	0
conv2d_5 (Conv2D)	(None, 28, 28, 64)	73792
batch_normalization_v2_4 (BatchNormalization)	(None, 28, 28, 64)	256
activation_1 (Activation)	(None, 28, 28, 64)	0
conv2d_6 (Conv2D)	(None, 28, 28, 1)	577
activation_2 (Activation)	(None, 28, 28, 1)	0
<hr/>		
Total params:	856,193	
Trainable params:	855,809	
Non-trainable params:	384	

Model: "sequential"		
Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 14, 14, 32)	320
leaky_re_lu (LeakyReLU)	(None, 14, 14, 32)	0
dropout (Dropout)	(None, 14, 14, 32)	0
conv2d_1 (Conv2D)	(None, 7, 7, 64)	18496
zero_padding2d (ZeroPadding2D)	(None, 8, 8, 64)	0
batch_normalization_v2 (BatchNormalization)	(None, 8, 8, 64)	256
leaky_re_lu_1 (LeakyReLU)	(None, 8, 8, 64)	0
dropout_1 (Dropout)	(None, 8, 8, 64)	0
conv2d_2 (Conv2D)	(None, 4, 4, 128)	73856
batch_normalization_v2_1 (BatchNormalization)	(None, 4, 4, 128)	512
leaky_re_lu_2 (LeakyReLU)	(None, 4, 4, 128)	0
dropout_2 (Dropout)	(None, 4, 4, 128)	0
conv2d_3 (Conv2D)	(None, 4, 4, 256)	295168
batch_normalization_v2_2 (BatchNormalization)	(None, 4, 4, 256)	1024
leaky_re_lu_3 (LeakyReLU)	(None, 4, 4, 256)	0
dropout_3 (Dropout)	(None, 4, 4, 256)	0
flatten (Flatten)	(None, 4096)	0
dense (Dense)	(None, 1)	4097

Total params: 393,729
Trainable params: 392,833
Non-trainable params: 896

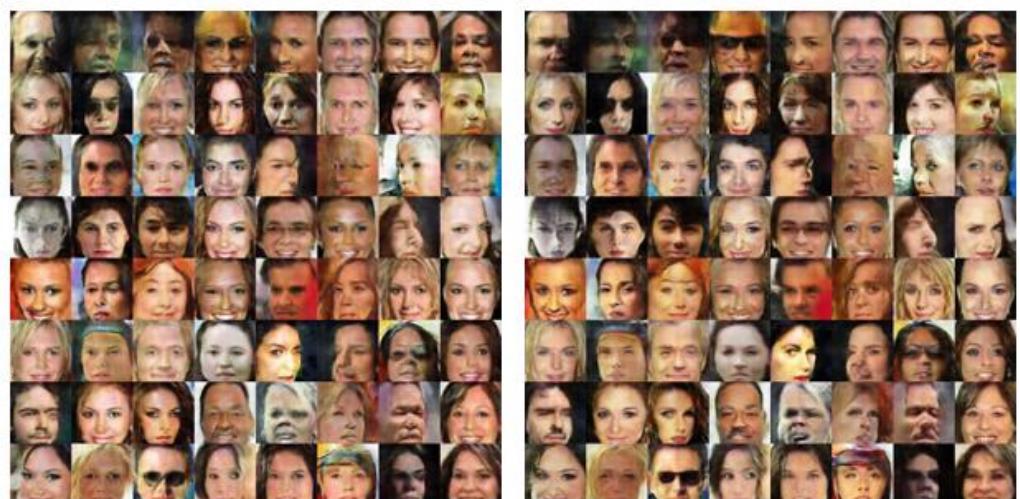




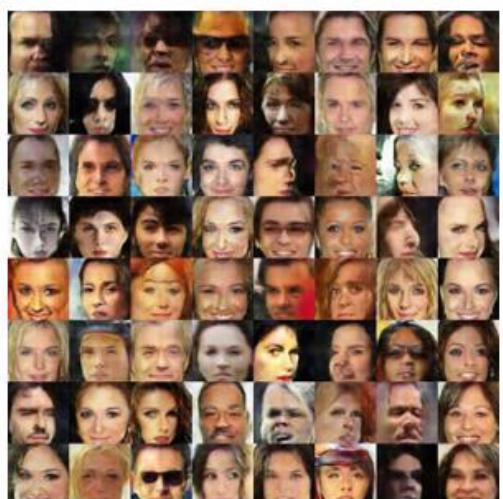
Epoch 2



Epoch 3



Epoch 4



Epoch 5

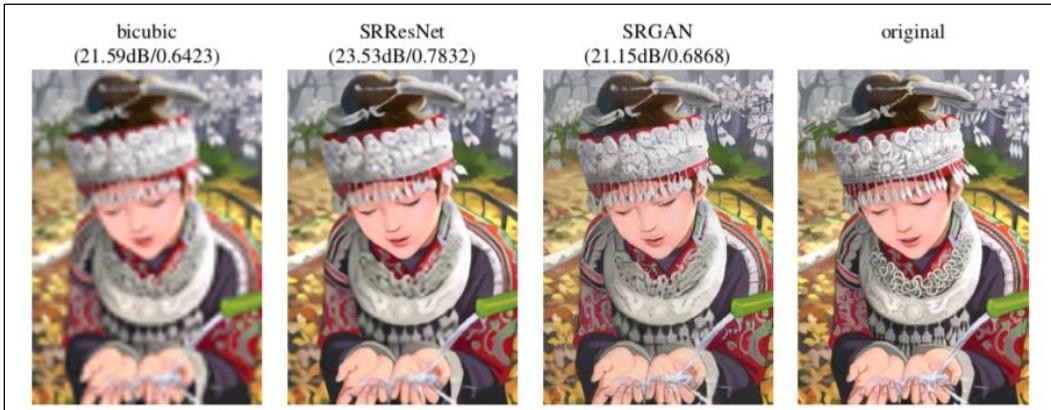
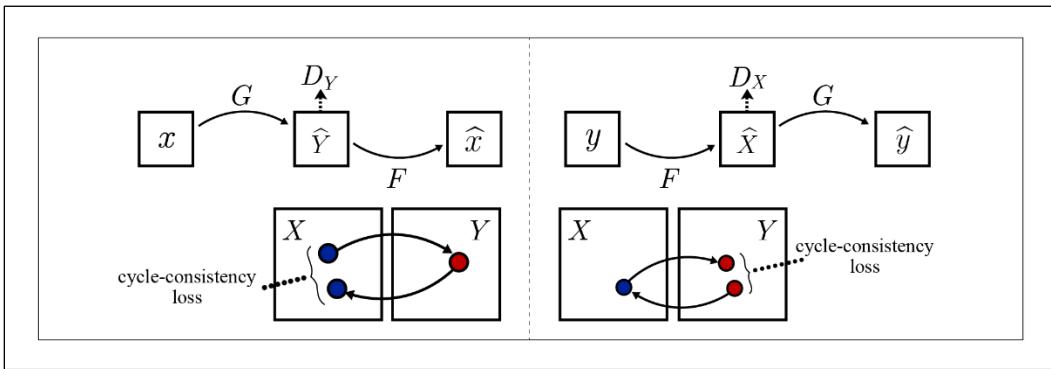
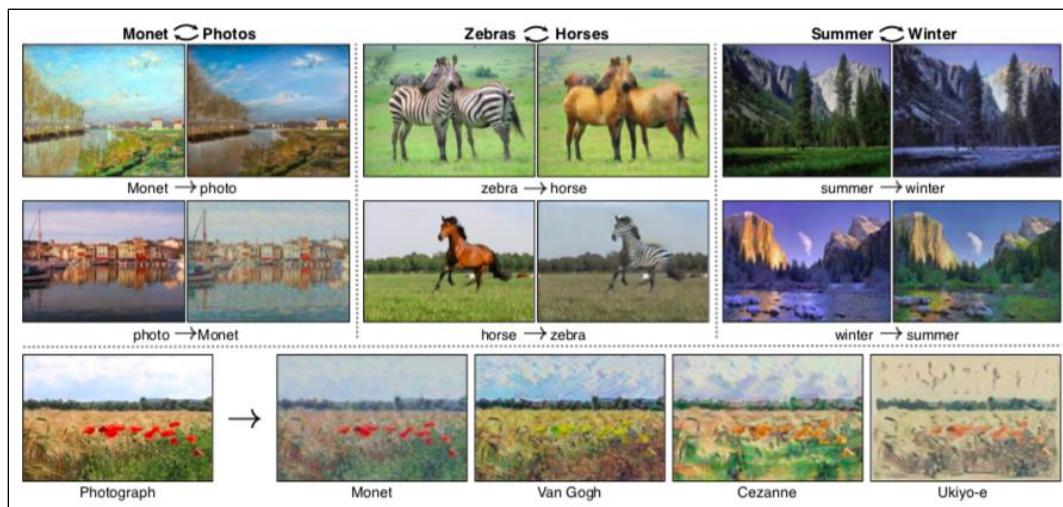
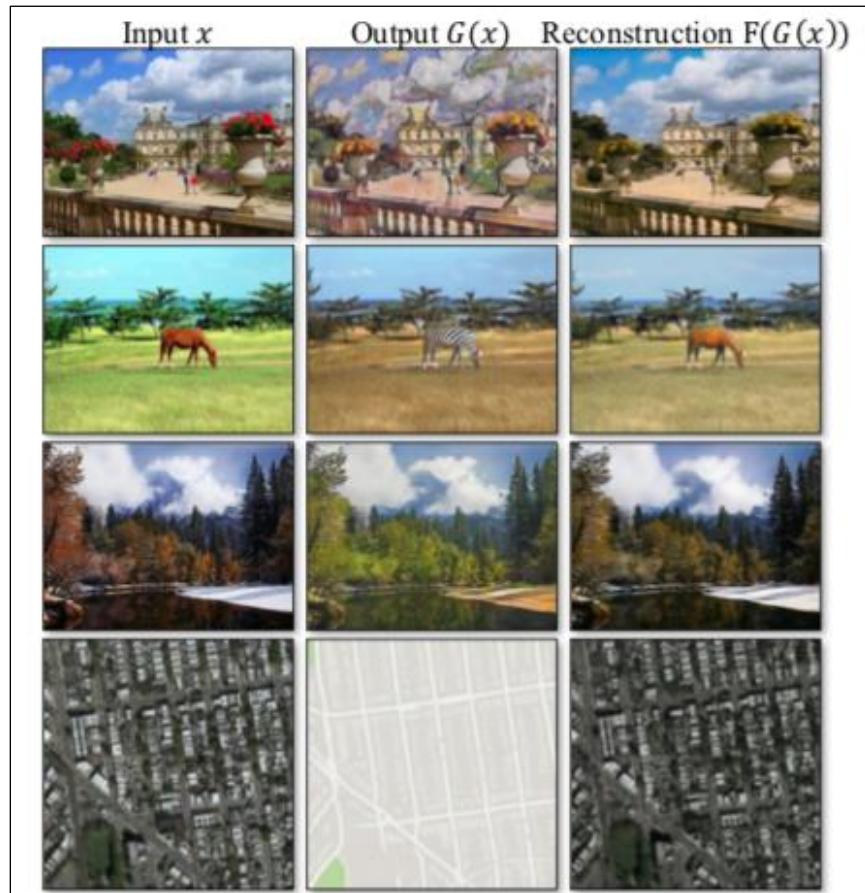
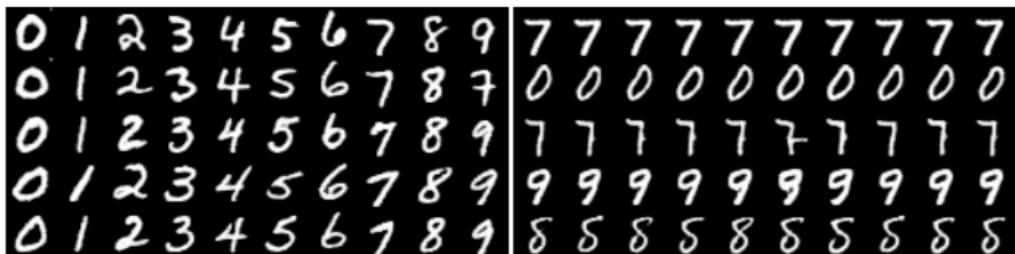
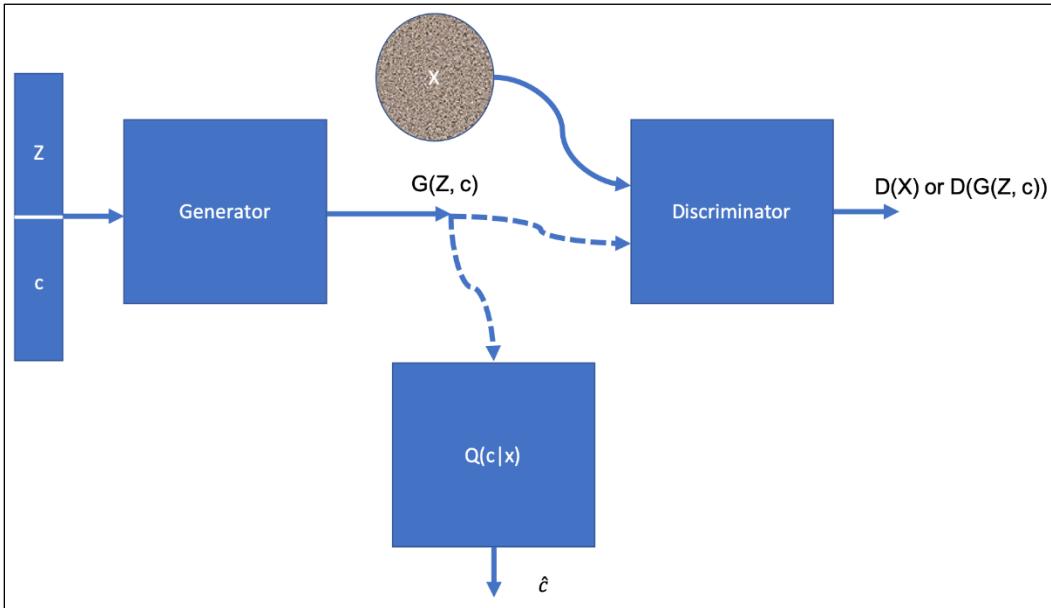


Figure 2: From left to right: bicubic interpolation, deep residual network optimized for MSE, deep residual generative adversarial network optimized for a loss more sensitive to human perception, original HR image. Corresponding PSNR and SSIM are shown in brackets. [4× upscaling]

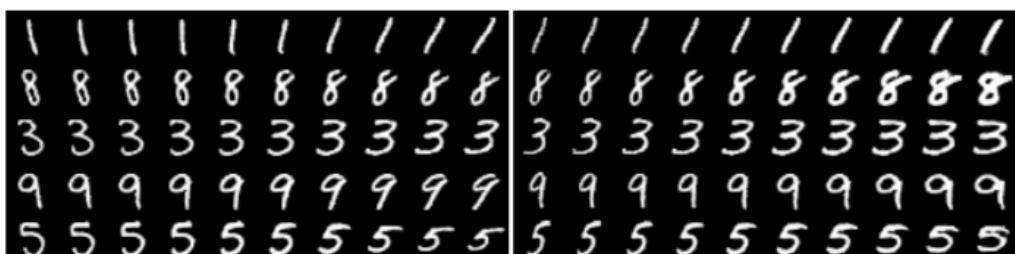






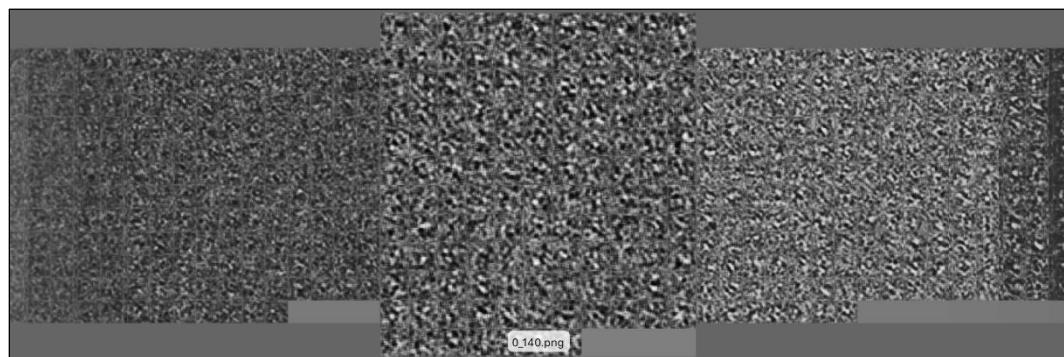
(a) Varying c_1 on InfoGAN (Digit type)

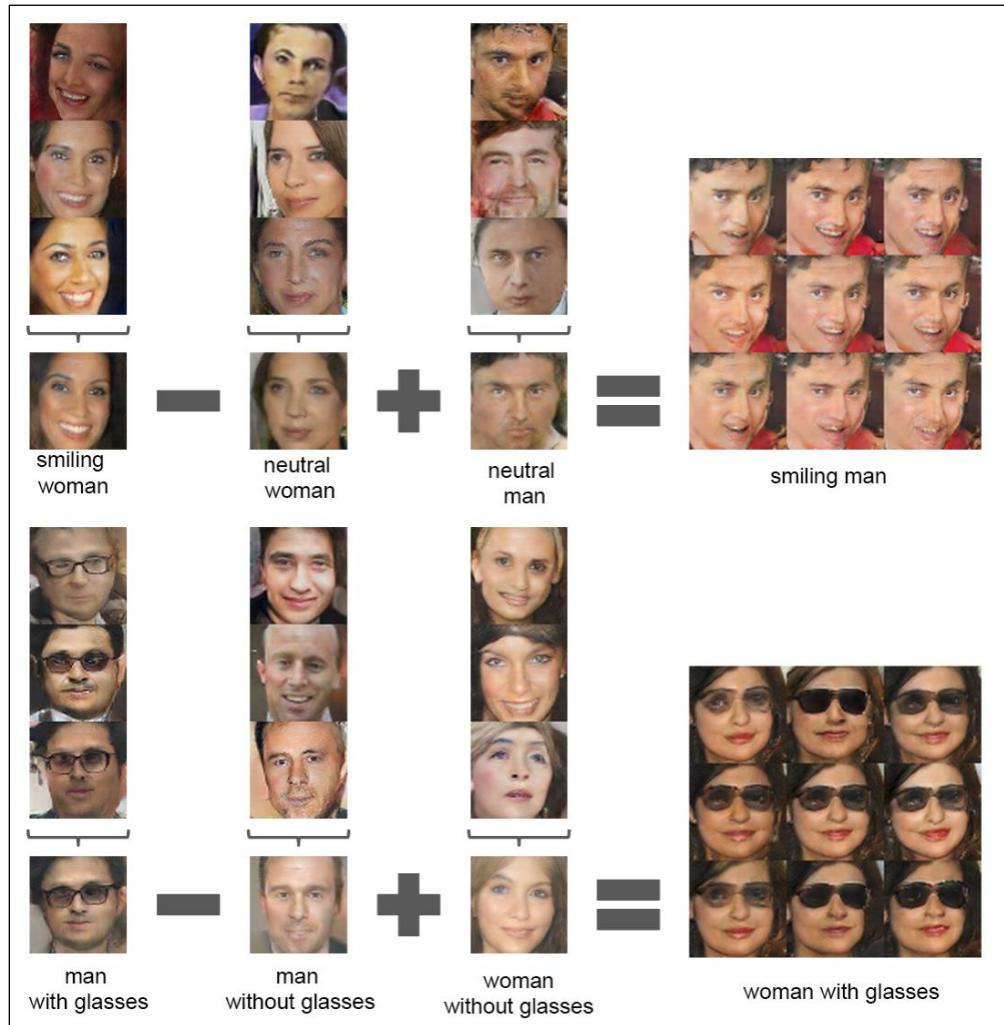
(b) Varying c_1 on regular GAN (No clear meaning)



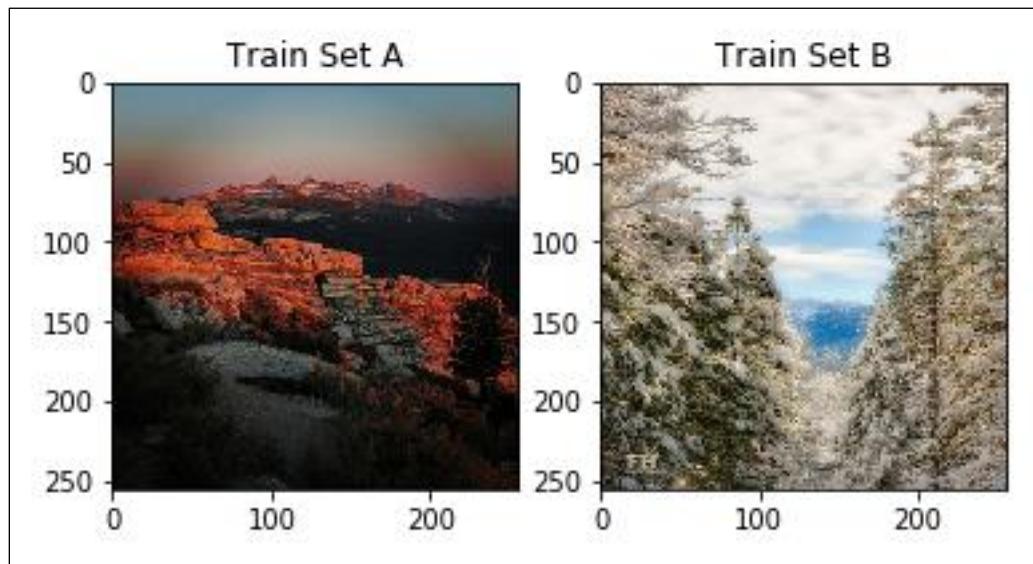
(c) Varying c_2 from -2 to 2 on InfoGAN (Rotation)

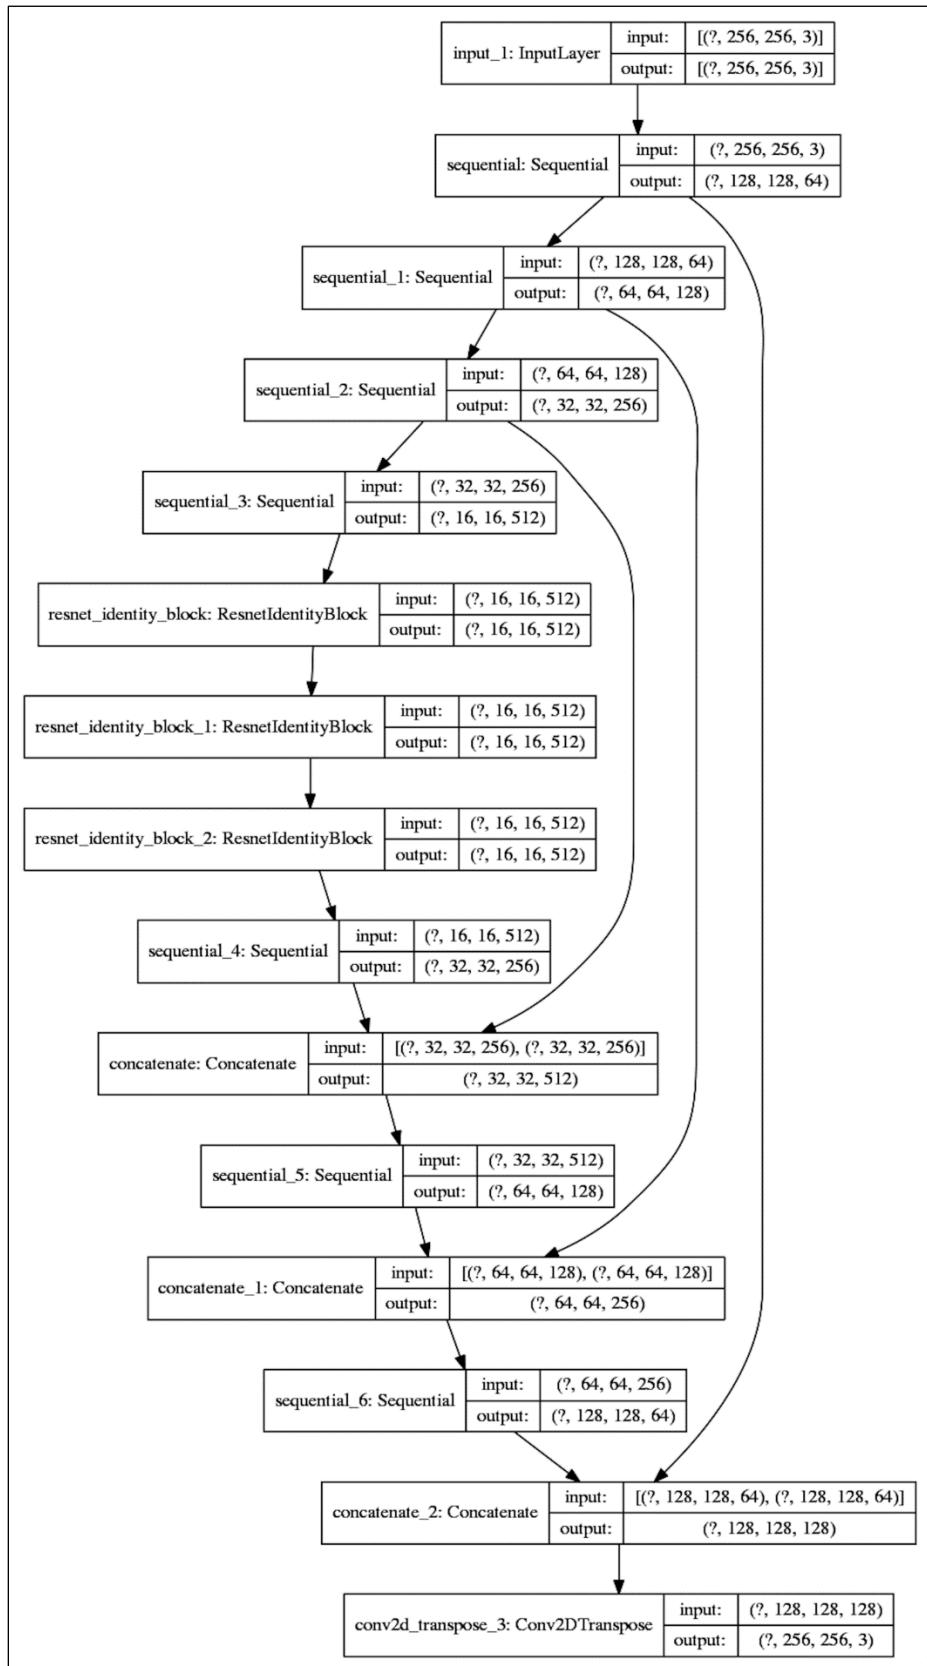
(d) Varying c_3 from -2 to 2 on InfoGAN (Width)

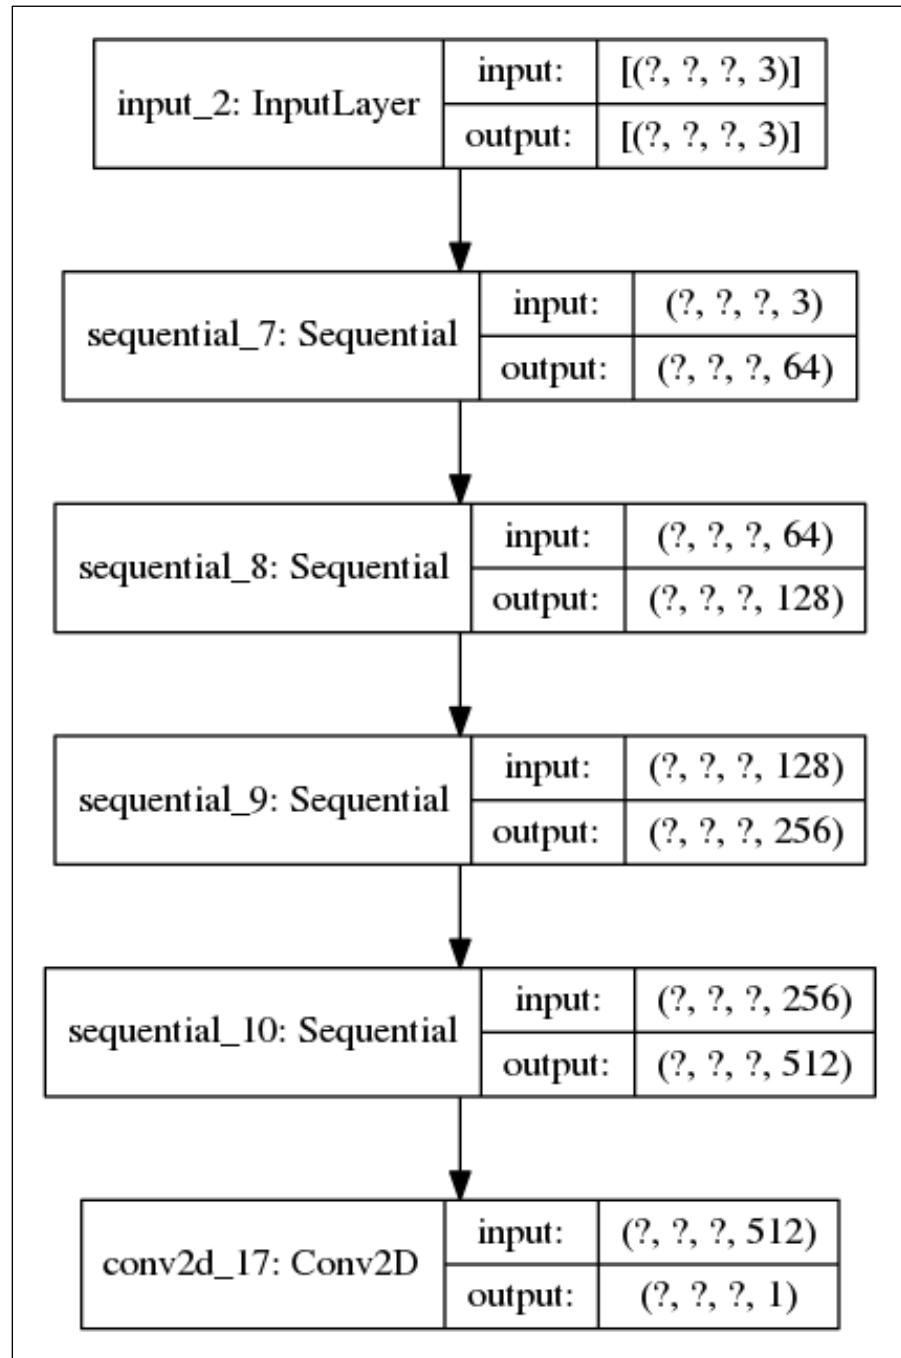


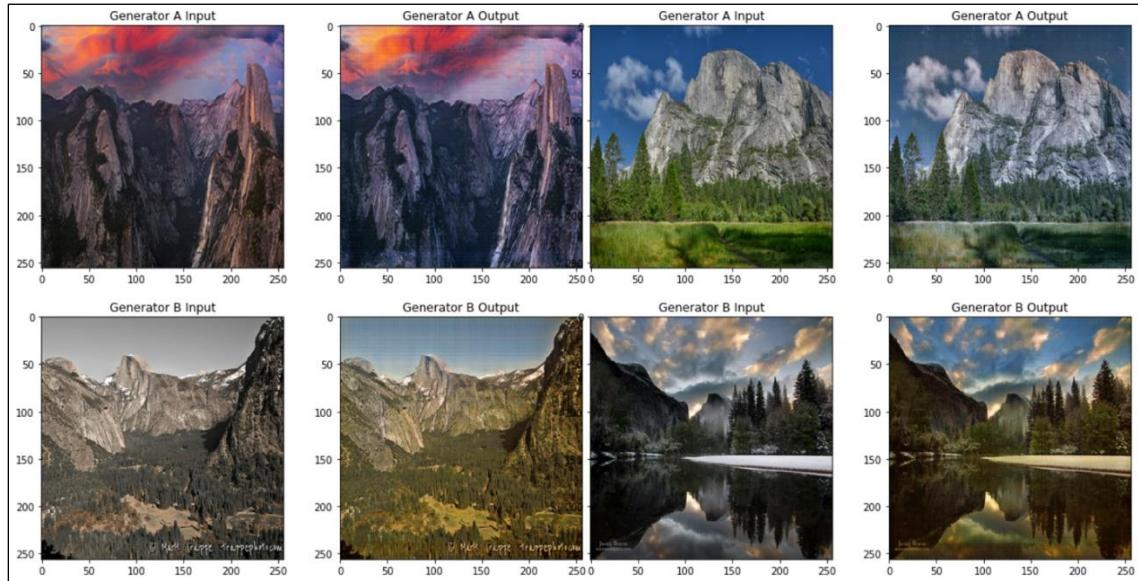
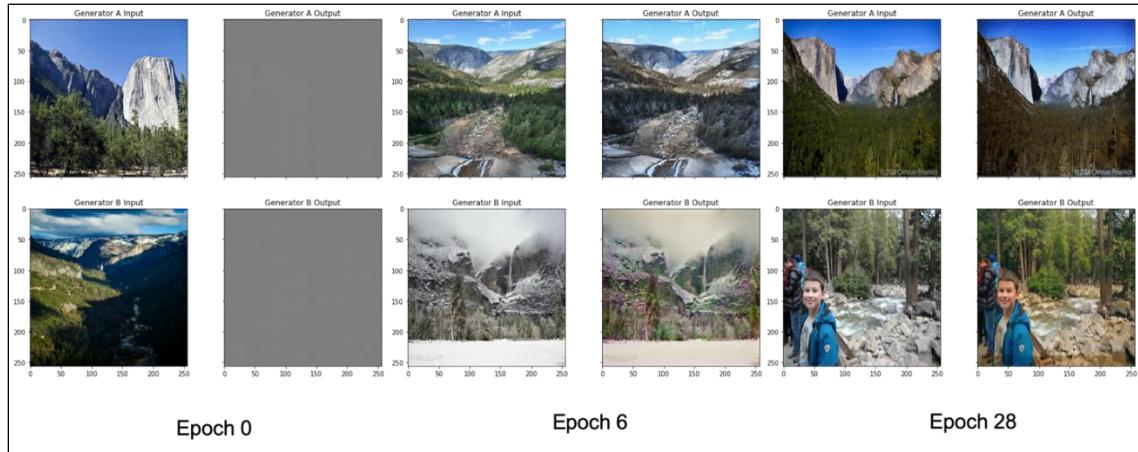




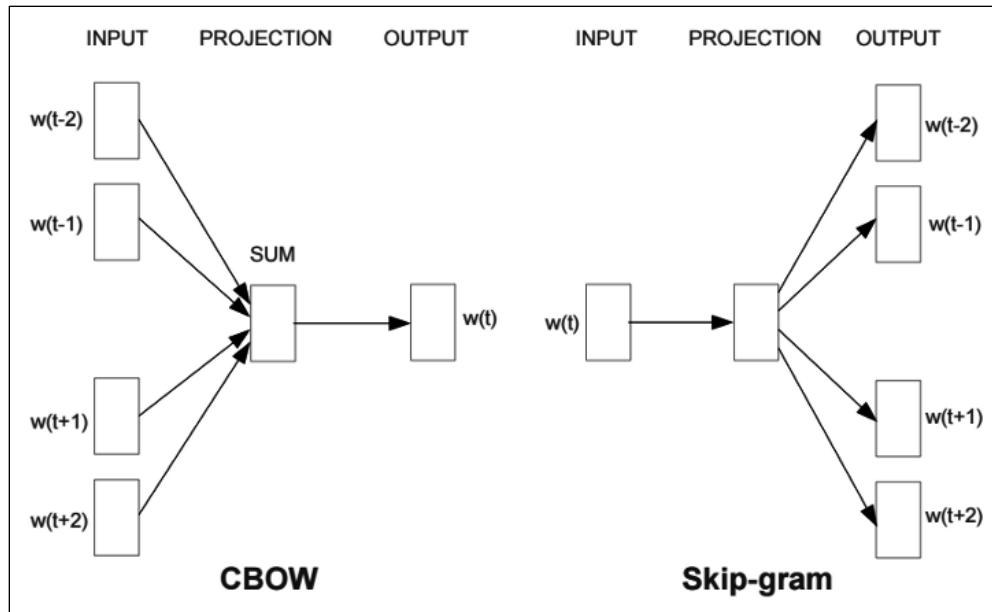
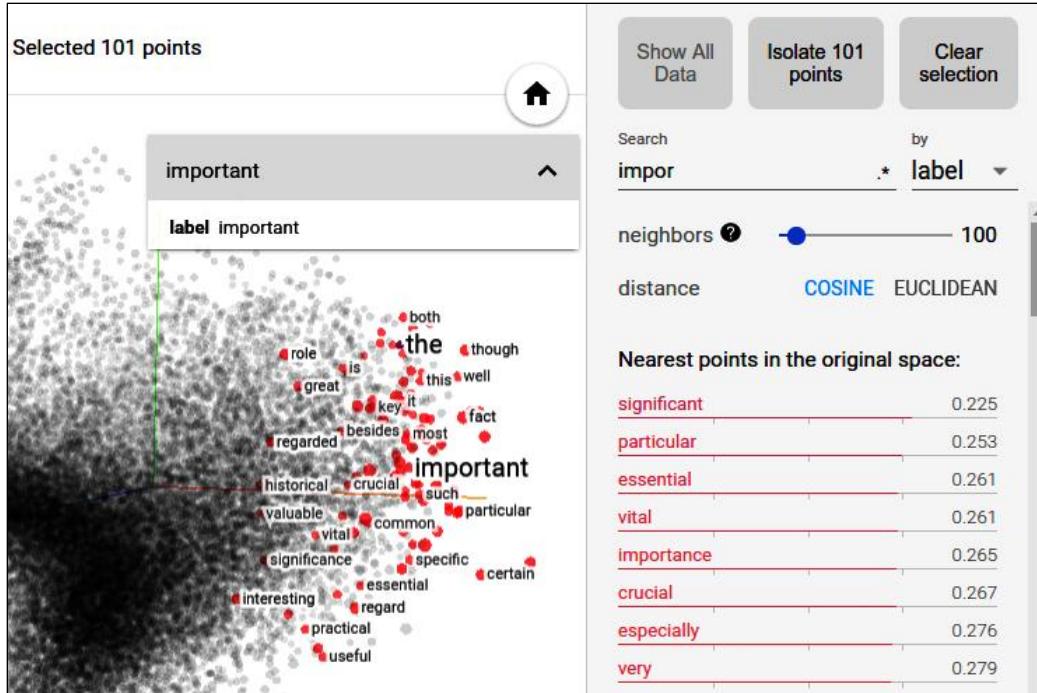


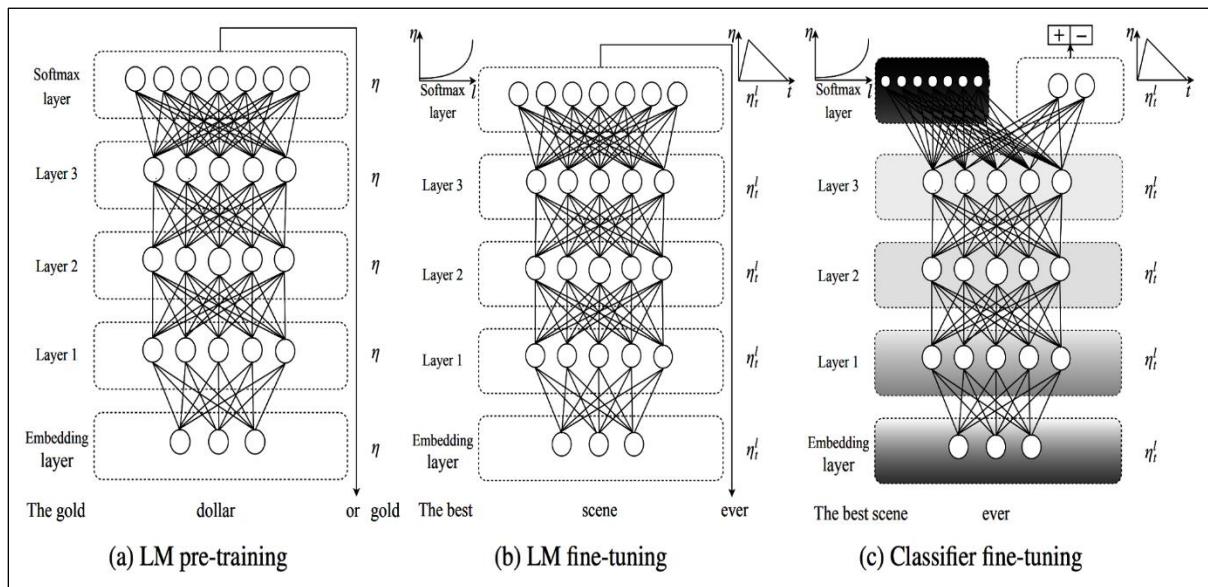
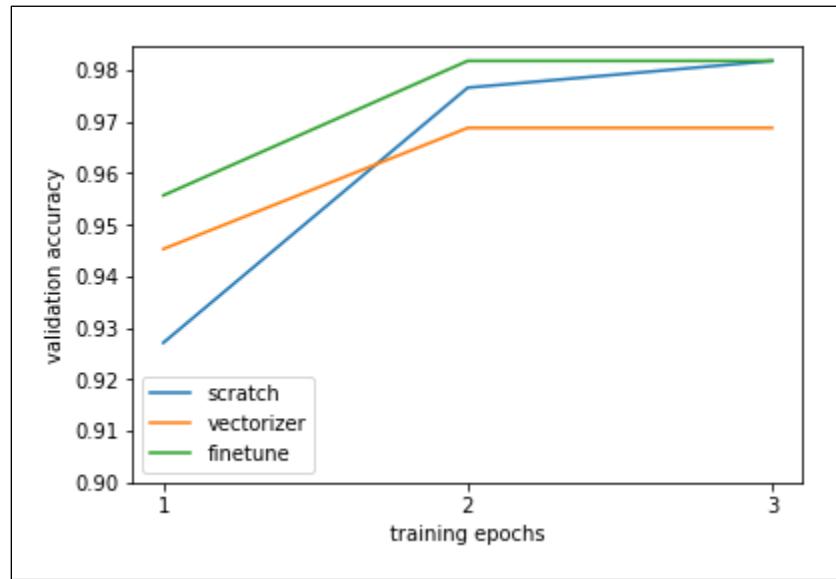




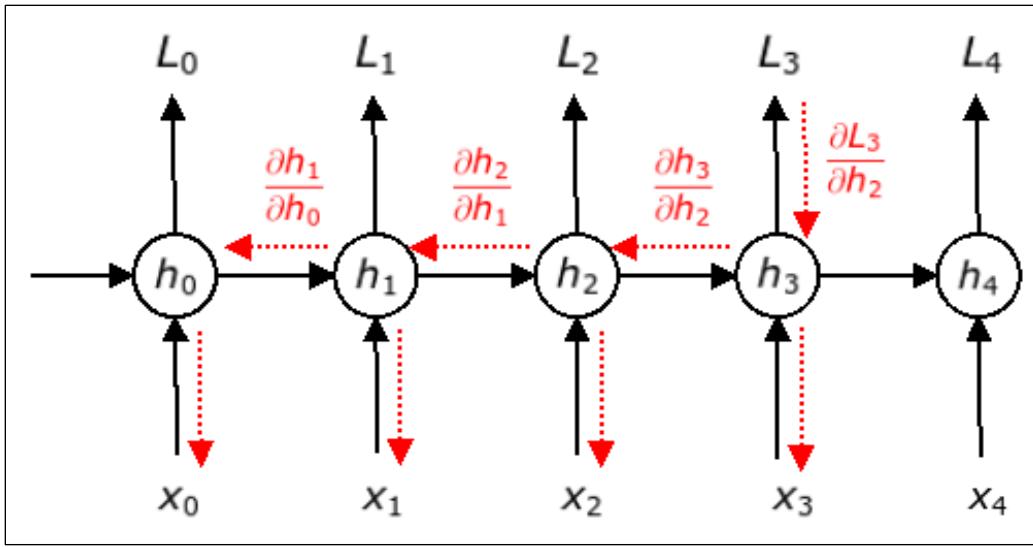
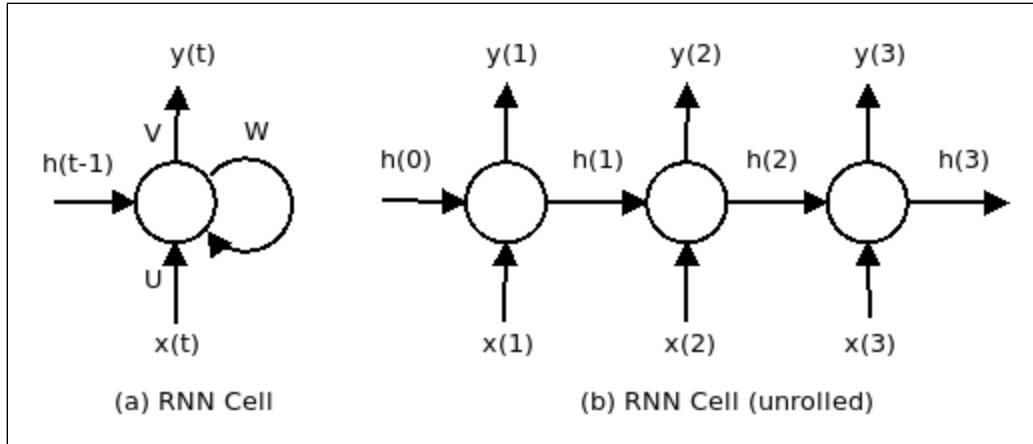


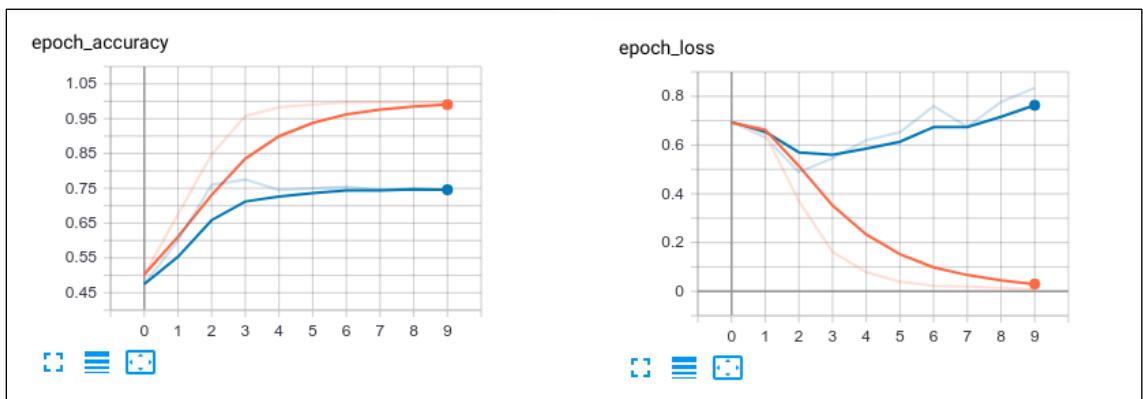
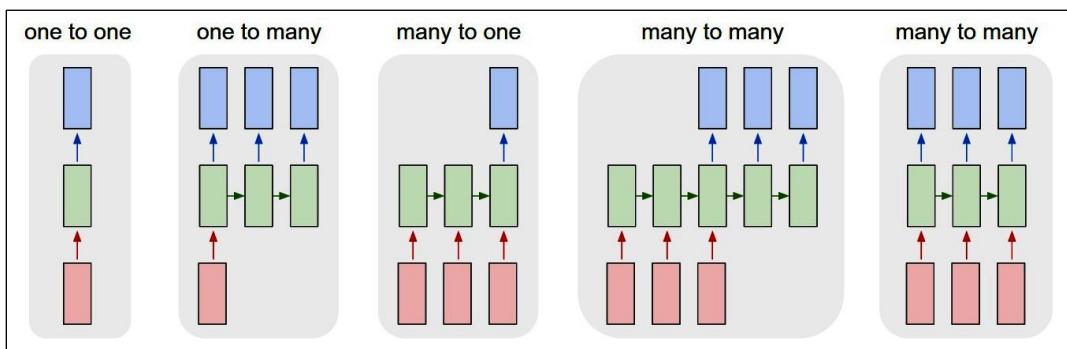
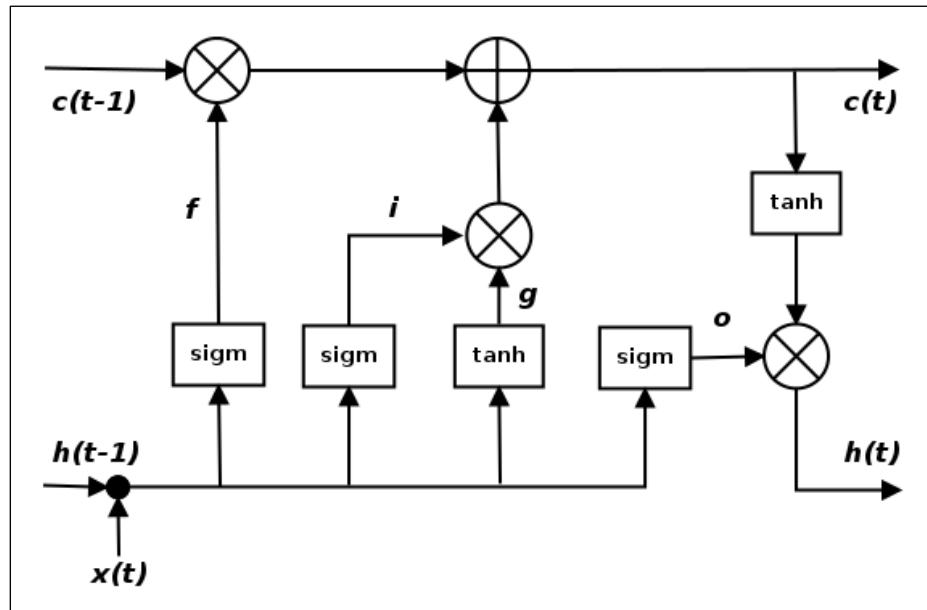
Chapter 7: Word Embeddings

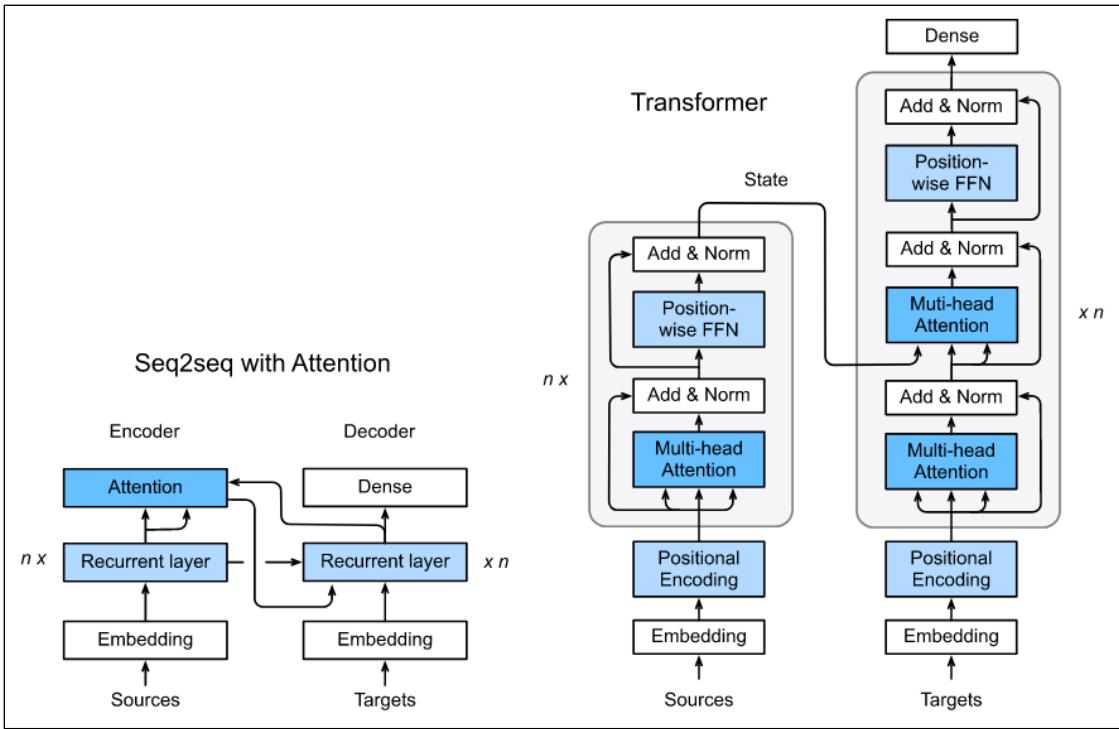
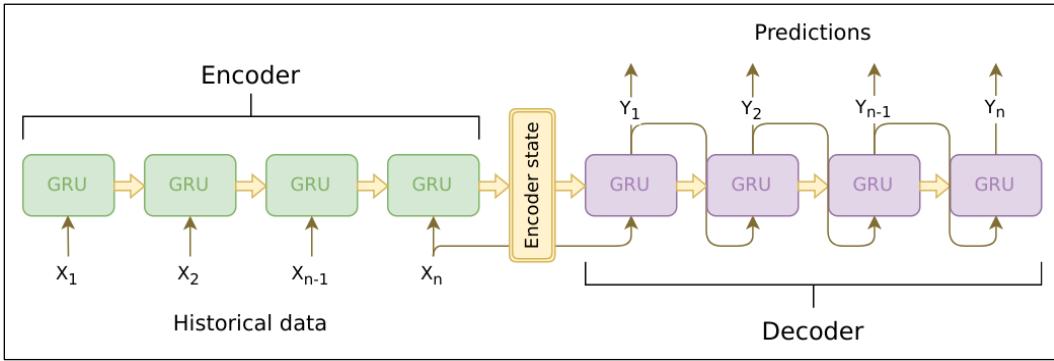




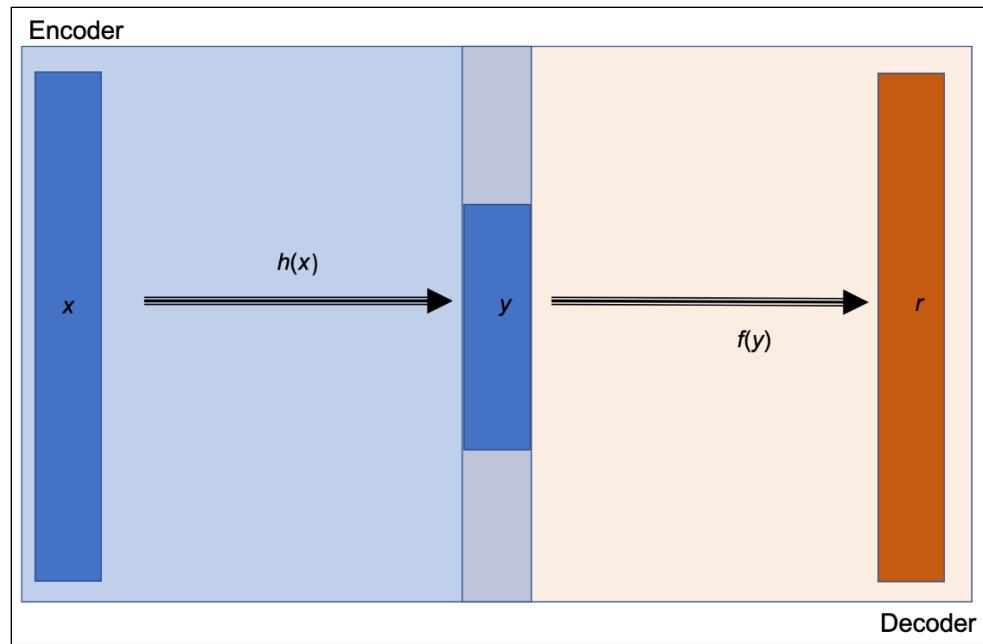
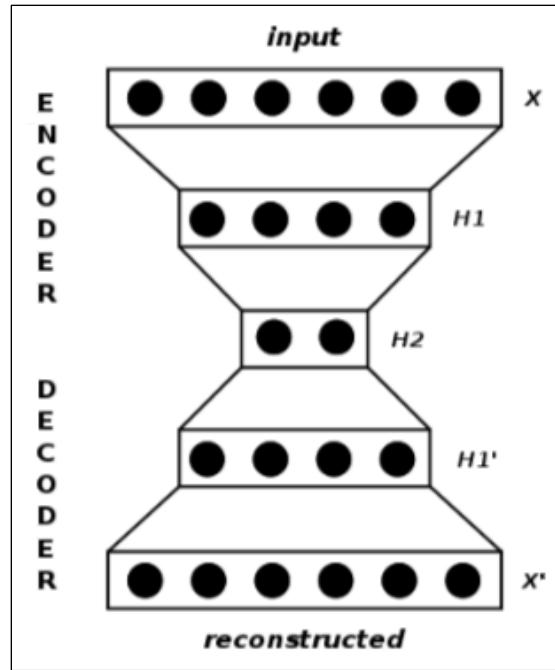
Chapter 8: Recurrent Neural Networks

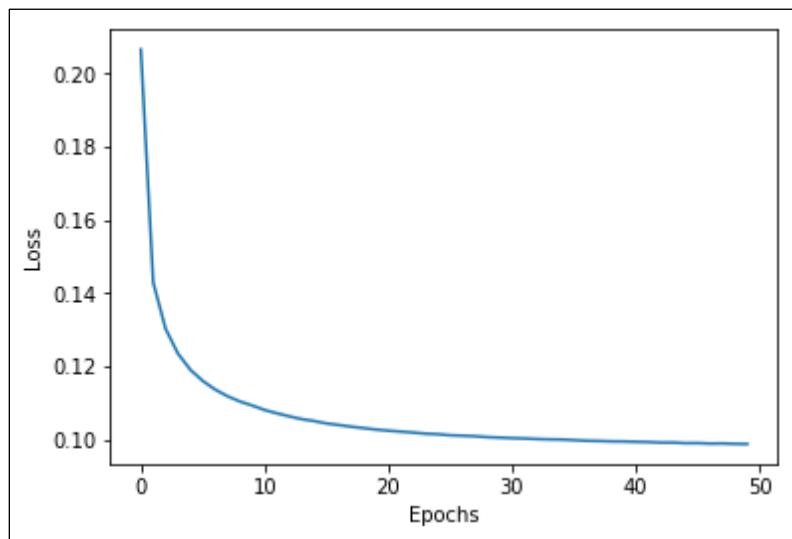
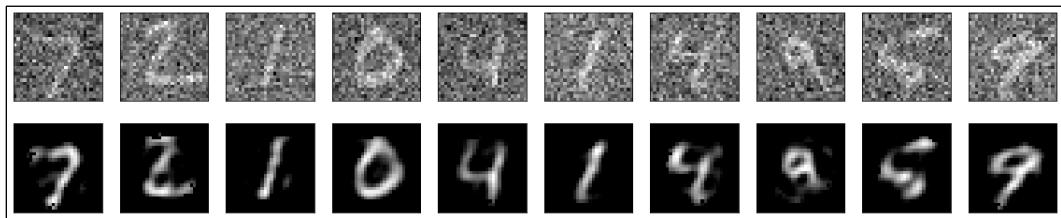
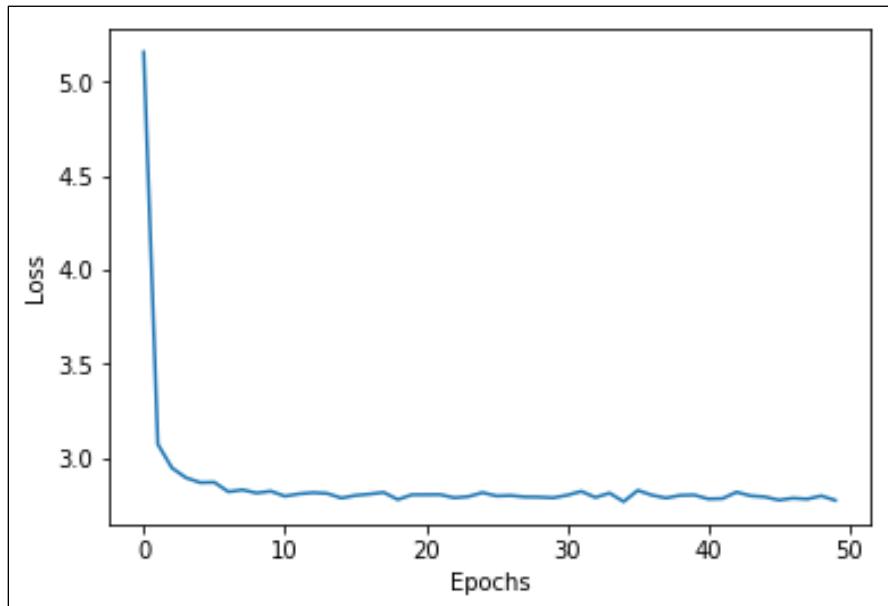


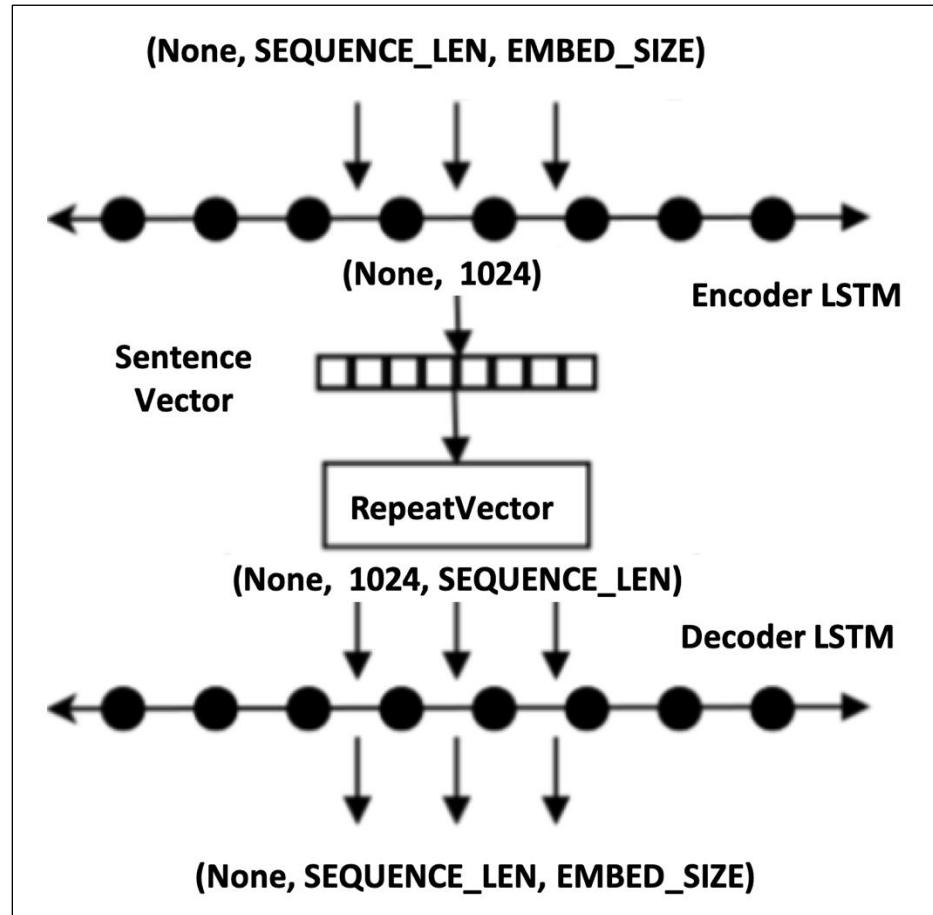
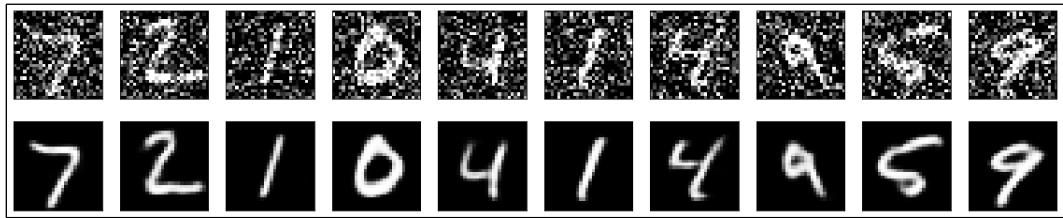


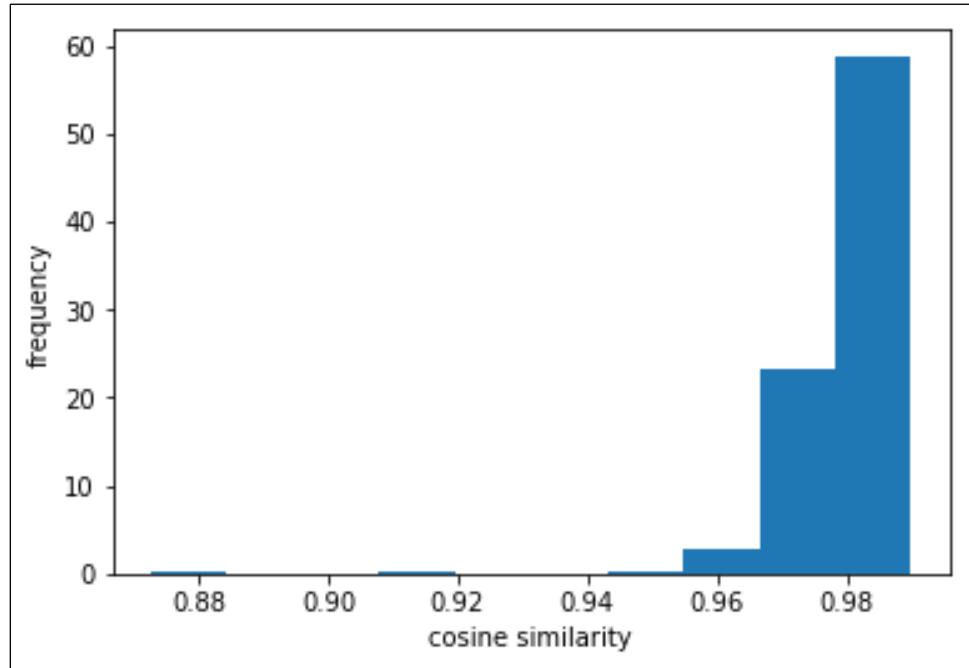
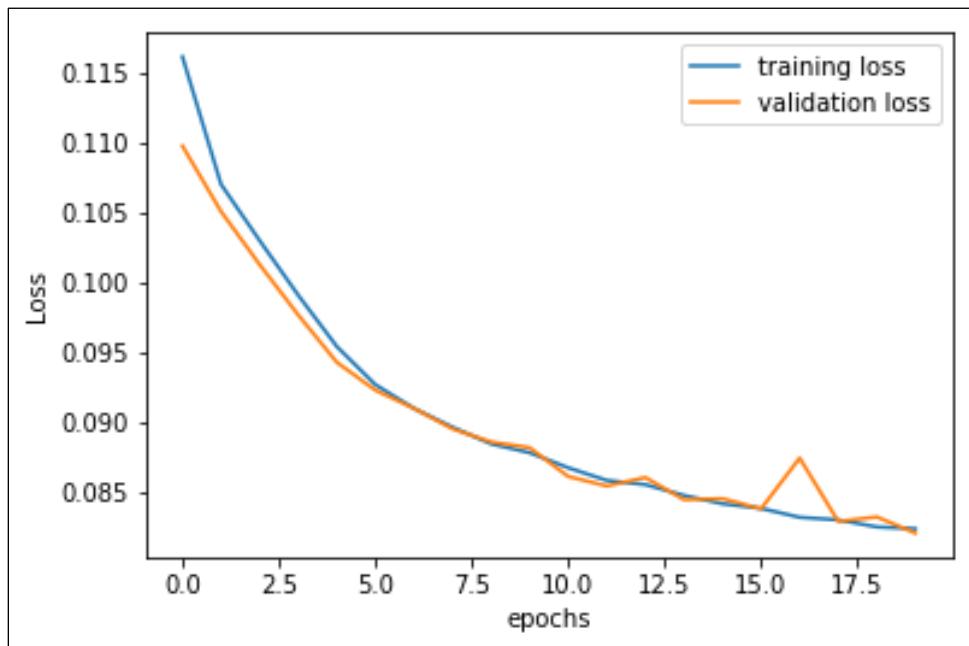


Chapter 9: Autoencoders

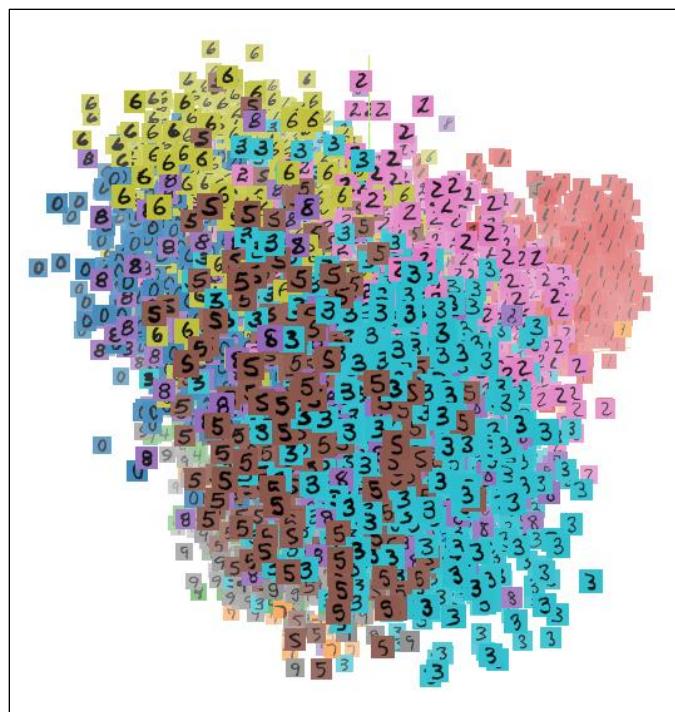
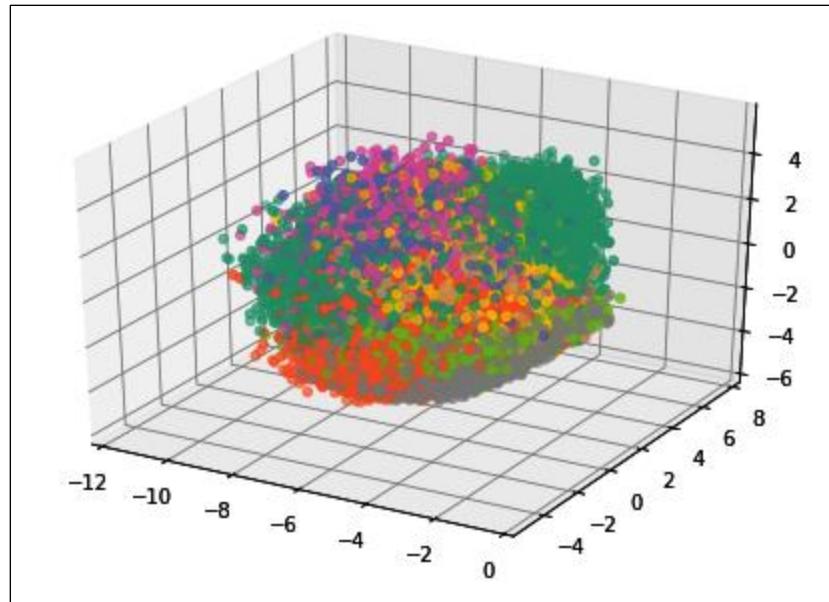


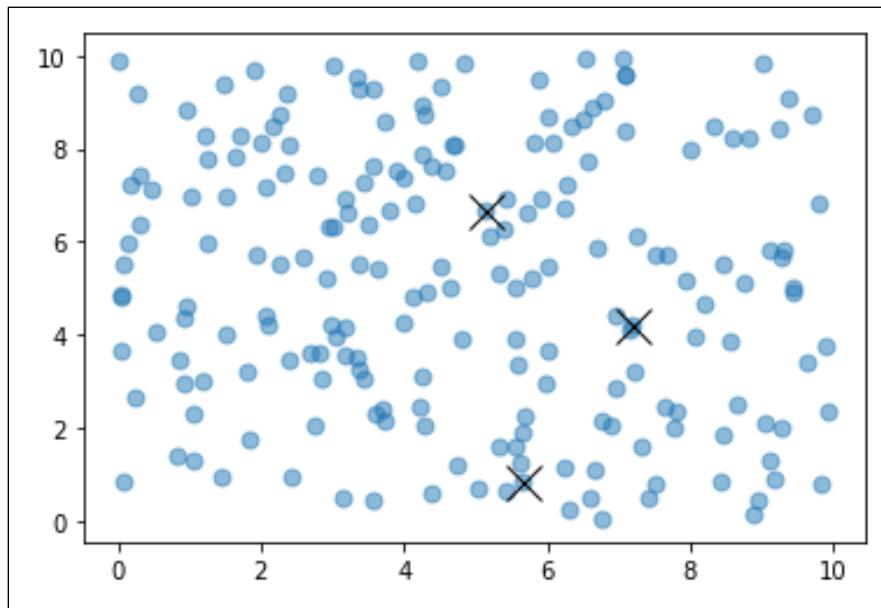
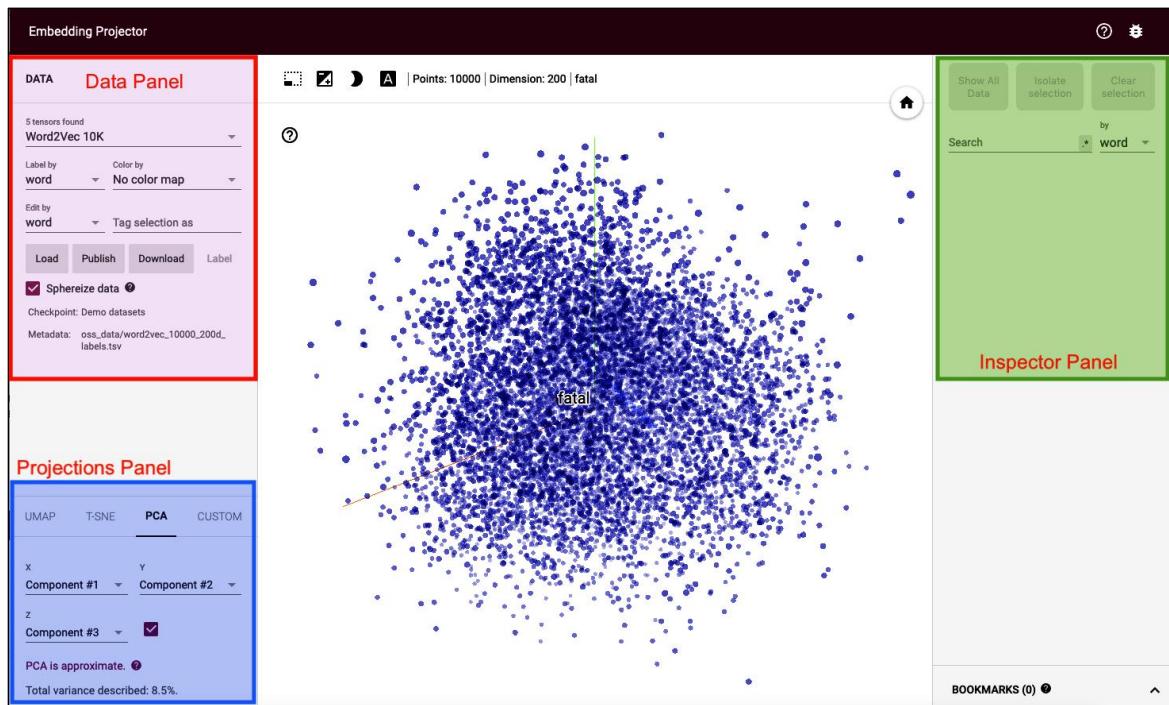


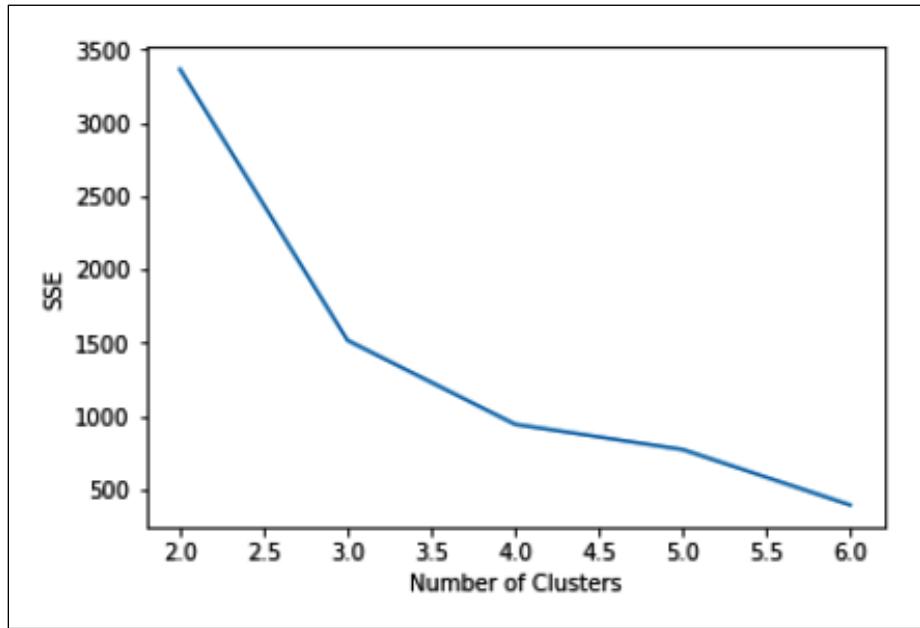
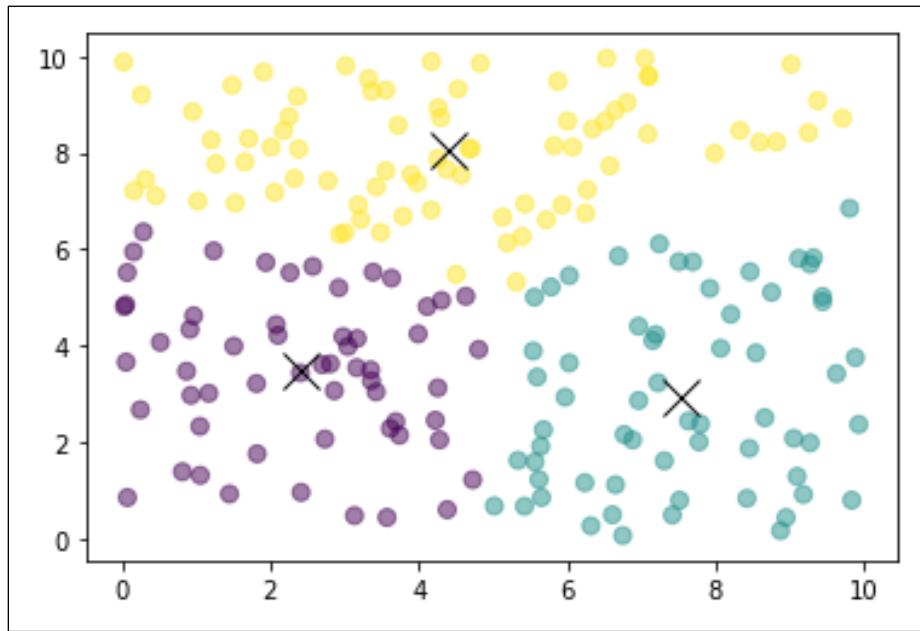


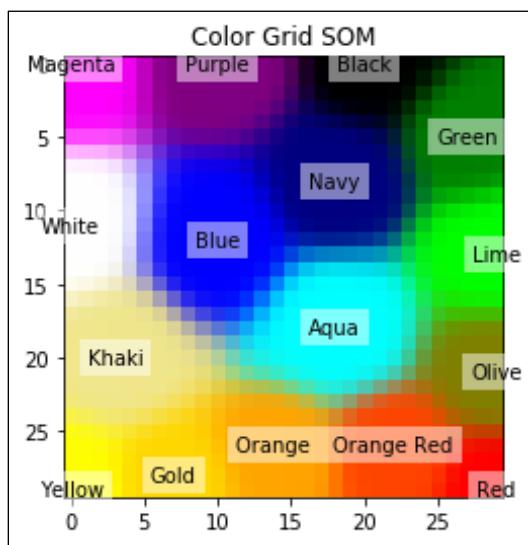
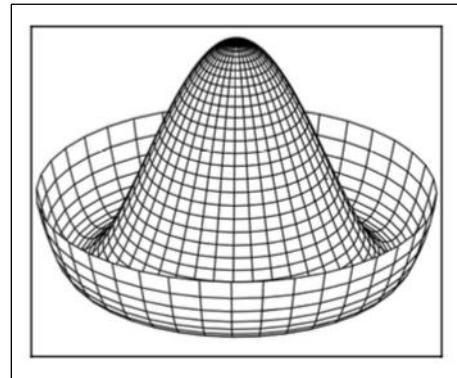
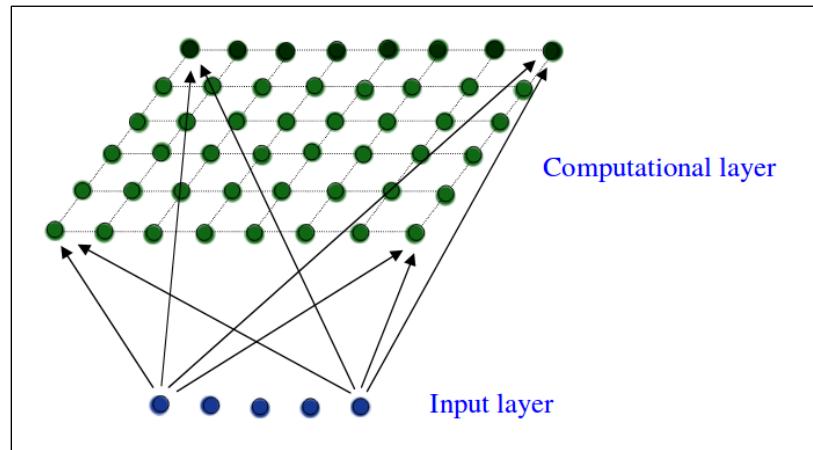


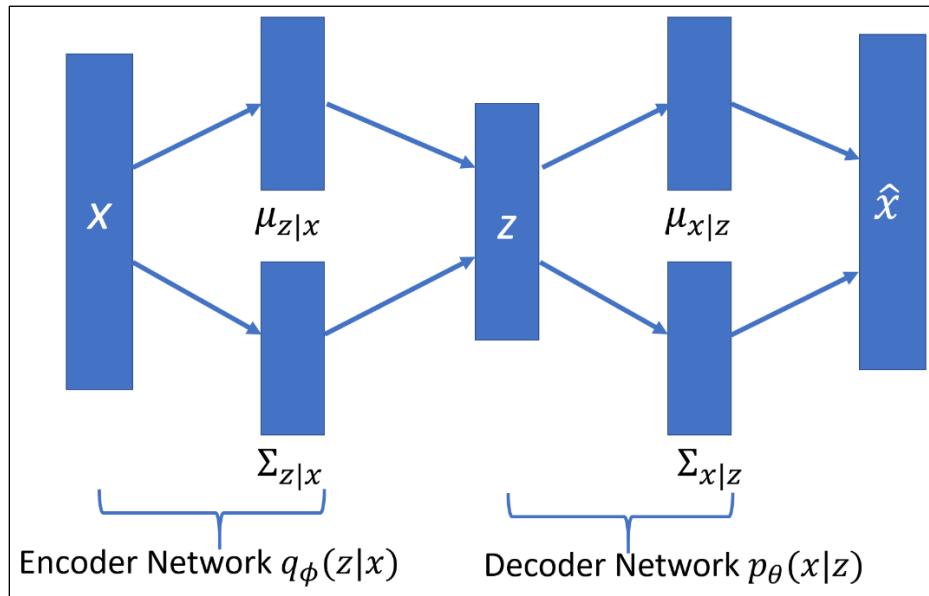
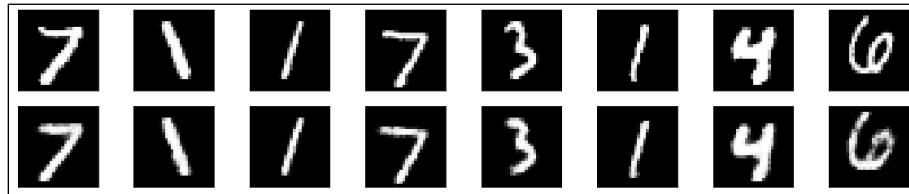
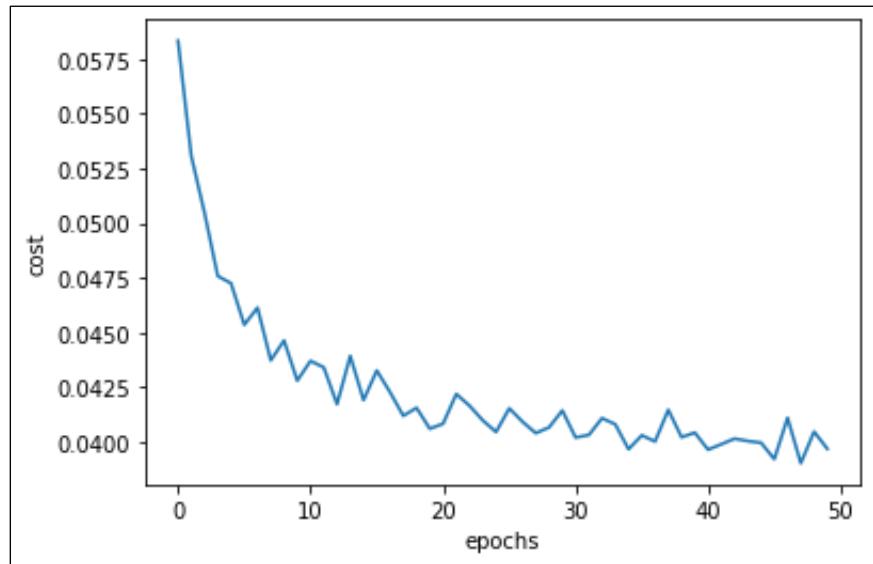
Chapter 10: Unsupervised Learning

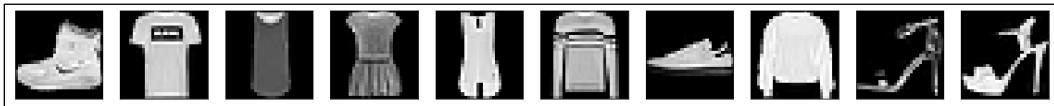






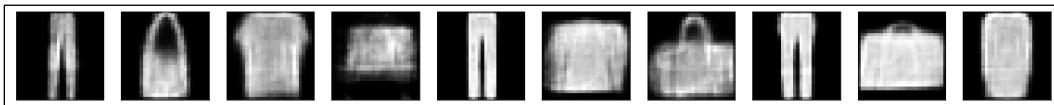




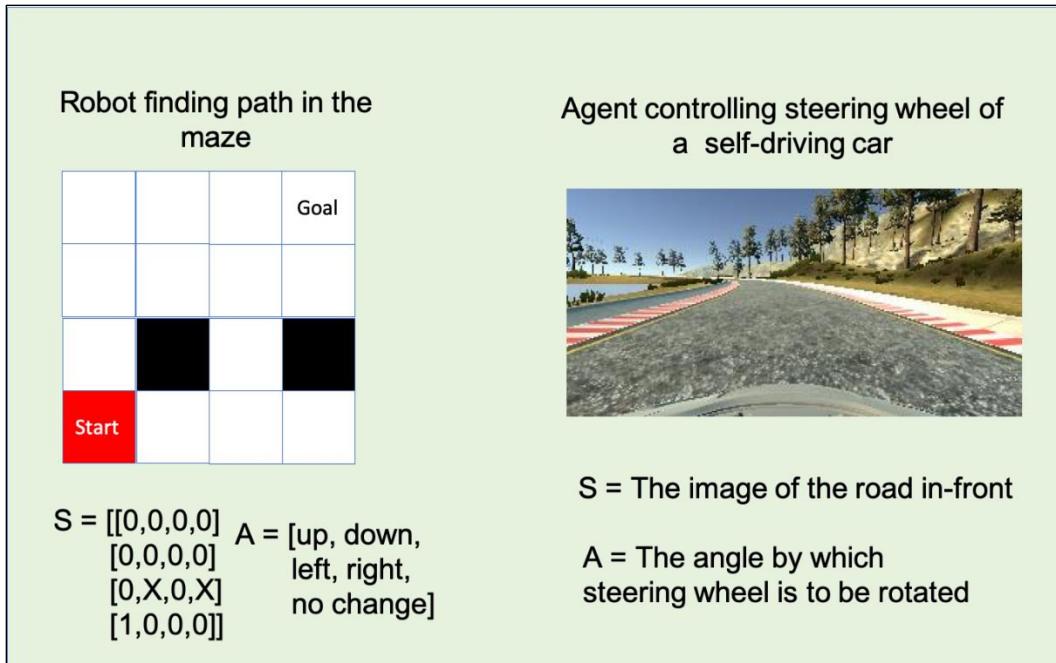
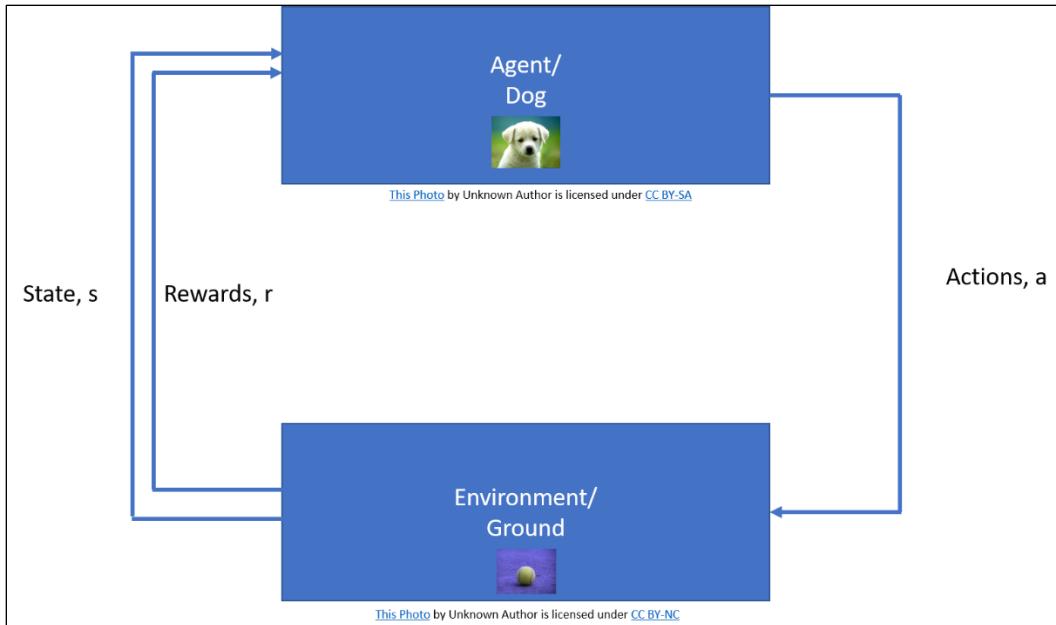


Model: "vae"

Layer (type)	Output Shape	Param #
<hr/>		
dense (Dense)	multiple	401920
dense_1 (Dense)	multiple	5130
dense_2 (Dense)	multiple	5130
dense_3 (Dense)	multiple	5632
dense_4 (Dense)	multiple	402192
<hr/>		
Total params: 820,004		
Trainable params: 820,004		
Non-trainable params: 0		



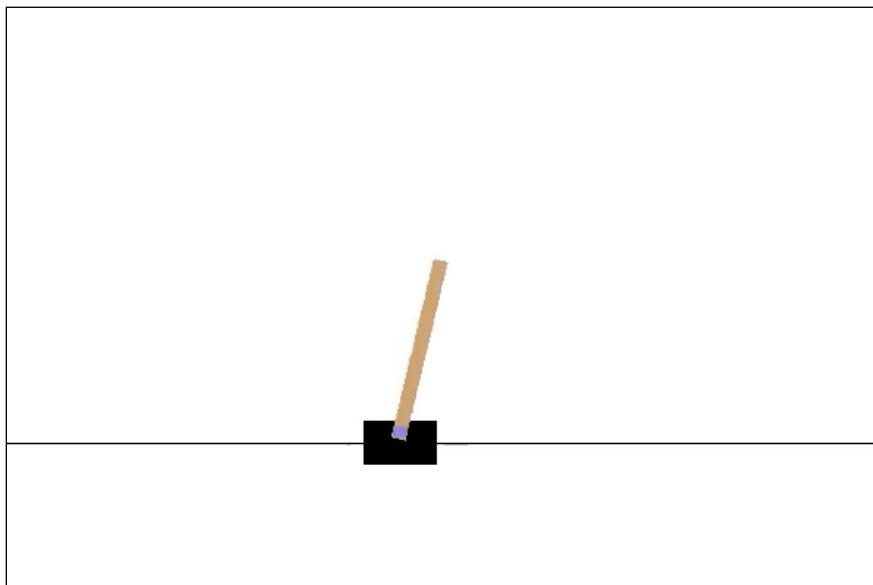
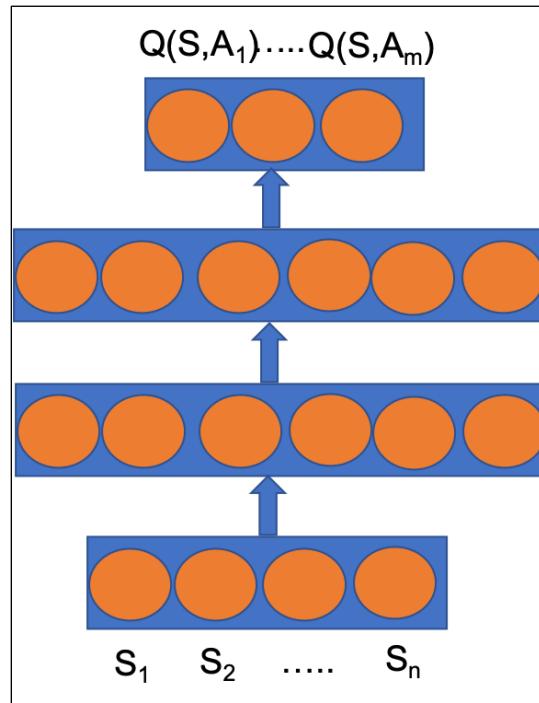
Chapter 11: Reinforcement Learning



-3	-2	-1	0
-4	-3	-2	-1
-5		-3	
-6	-5	-4	-5

Each box has the value function:
Number of steps needed to reach goal (green box)

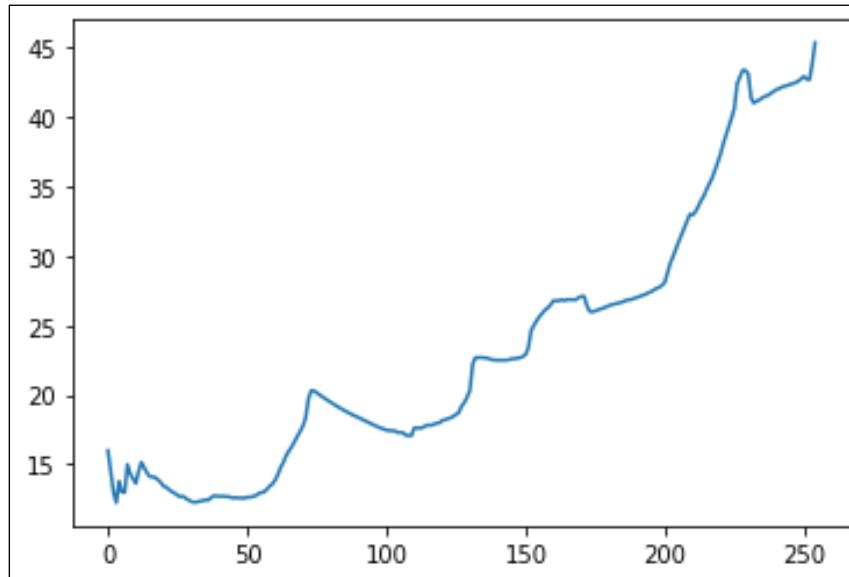


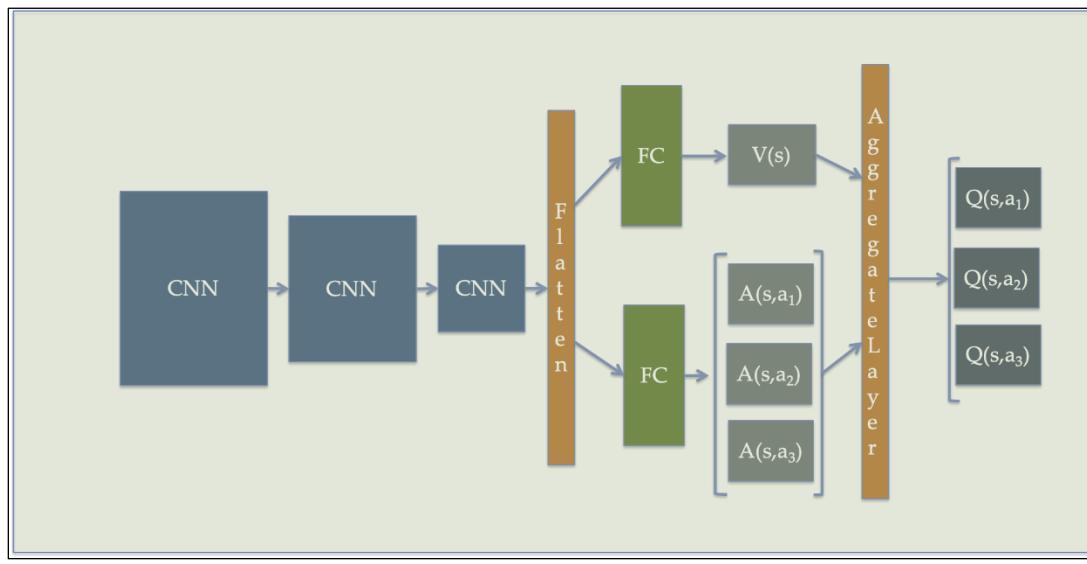
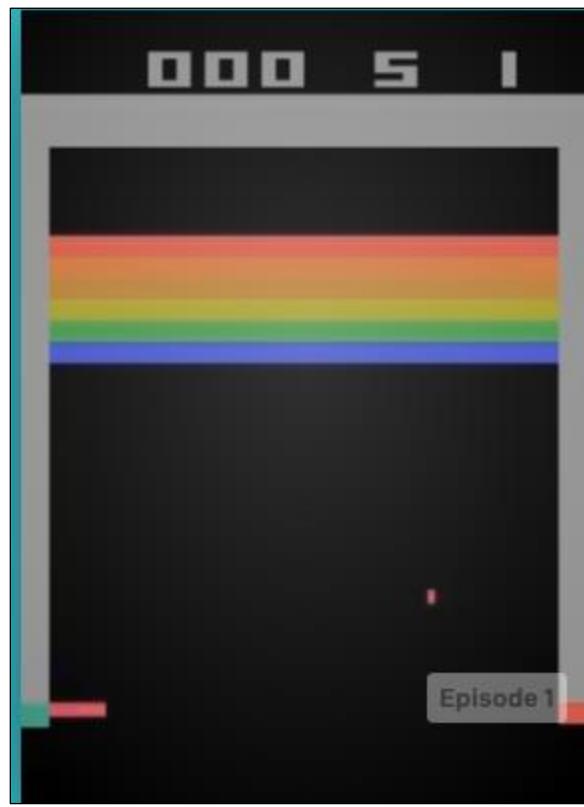


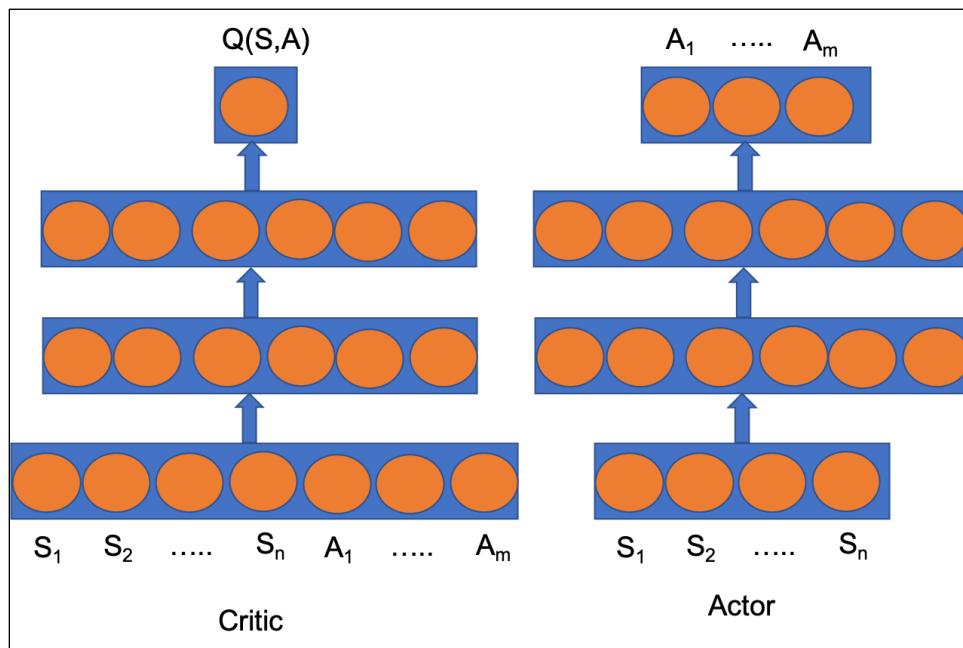
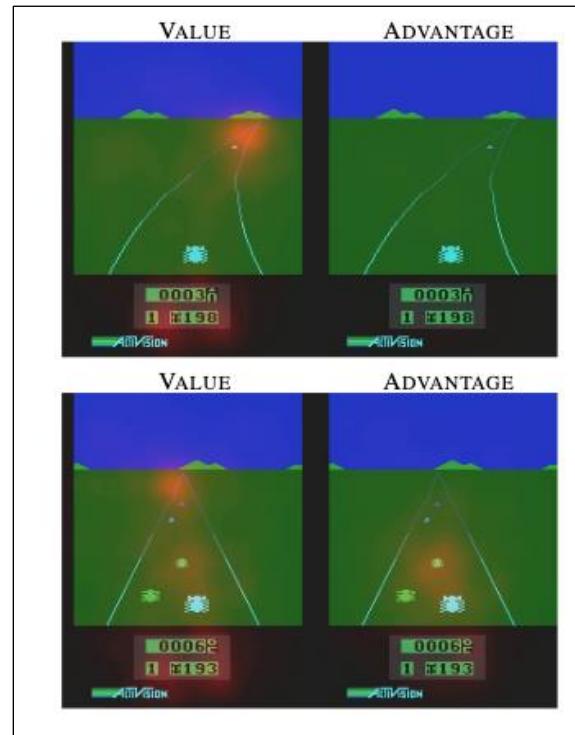
Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 24)	120
dense_1 (Dense)	(None, 48)	1200
dense_2 (Dense)	(None, 2)	98
<hr/>		
Total params: 1,418		
Trainable params: 1,418		
Non-trainable params: 0		

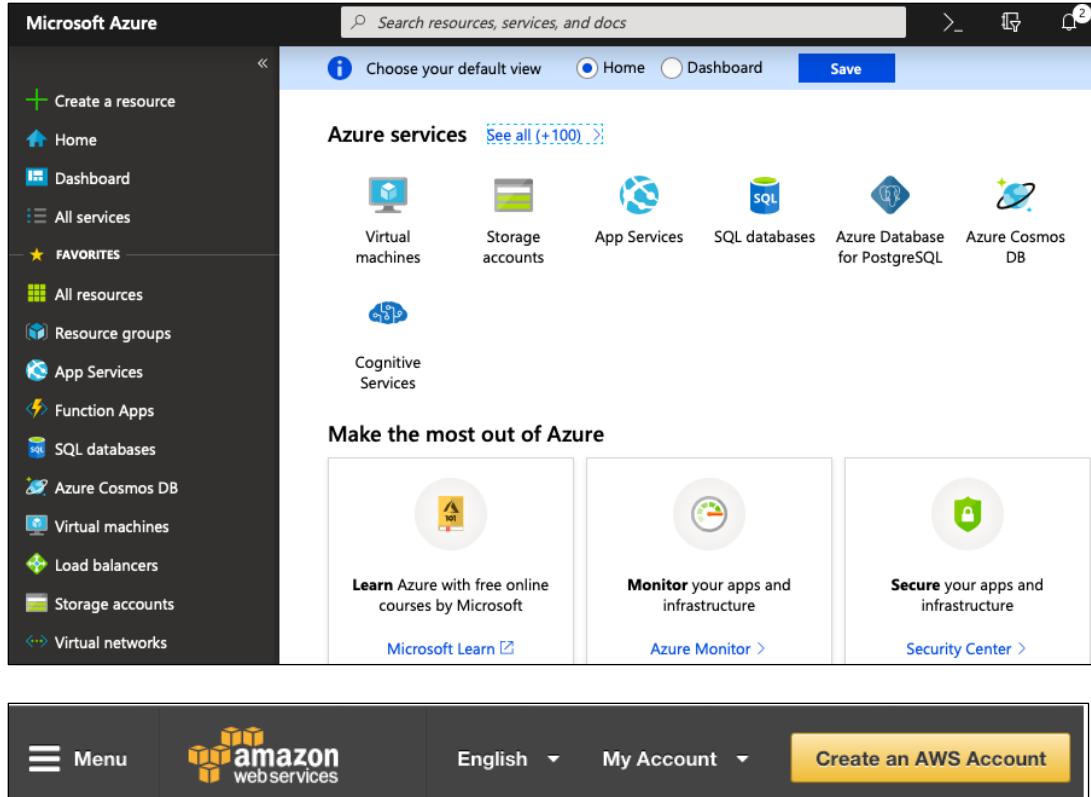
[Episode 0] – Mean survival time over last 100 episodes was 16.0 ticks.
[Episode 100] – Mean survival time over last 100 episodes was 17.47 ticks.
[Episode 200] – Mean survival time over last 100 episodes was 28.1 ticks.
Ran 254 episodes. Solved after 154 trials ✓



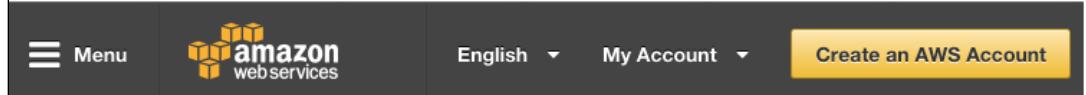




Chapter 12: TensorFlow and Cloud



The screenshot shows the Microsoft Azure portal homepage. On the left, there's a dark sidebar with various navigation options like 'Create a resource', 'Home', 'Dashboard', and 'All services'. Below these are sections for 'FAVORITES' (Virtual machines, Storage accounts, App Services, SQL databases, Azure Database for PostgreSQL, Azure Cosmos DB), 'All resources' (Resource groups, App Services, Function Apps, SQL databases, Azure Cosmos DB, Virtual machines, Load balancers, Storage accounts, Virtual networks), and 'Virtual networks'. The main content area has a search bar at the top. It features a 'Choose your default view' button, a 'Home' radio button (which is selected), and a 'Dashboard' radio button. A 'Save' button is also present. Below this, there's a section titled 'Azure services' with a link to 'See all (+100)'. It lists several services with icons: Virtual machines, Storage accounts, App Services, SQL databases, Azure Database for PostgreSQL, and Azure Cosmos DB. There's also a separate icon for Cognitive Services. At the bottom, there's a section titled 'Make the most out of Azure' with three cards: 'Learn Azure with free online courses by Microsoft' (with a Microsoft Learn link), 'Monitor your apps and infrastructure' (with an Azure Monitor link), and 'Secure your apps and infrastructure' (with a Security Center link).



The screenshot shows the Amazon Web Services (AWS) homepage. At the top, there's a dark header with a 'Menu' icon, the 'amazon webservices' logo, language and account dropdowns ('English', 'My Account'), and a prominent yellow 'Create an AWS Account' button.

The screenshot shows the AWS Management Console homepage. At the top, there is a navigation bar with the AWS logo, a "Services" dropdown, a "Resource Groups" dropdown, a bell icon for notifications, and account-specific options for "AK", "Oregon", and "Support".

AWS Management Console

AWS services

Find Services
You can enter names, keywords or acronyms.

All services

- Compute**
 - EC2
 - Lightsail [\[?\]](#)
 - ECR
 - ECS
 - EKS
 - Lambda
 - Batch
 - Elastic Beanstalk
 - Serverless Application Repository
- Storage**
 - S3
 - EFS
 - FSx
 - S3 Glacier
 - Storage Gateway
 - AWS Backup
- Database**
 - RDS
- Machine Learning**
 - Amazon SageMaker
 - Amazon Comprehend
 - Amazon Forecast
 - Amazon Lex
 - Amazon Machine Learning
 - Amazon Personalize
 - Amazon Polly
 - Amazon Rekognition
 - Amazon Textract
 - Amazon Transcribe
 - Amazon Translate
 - AWS DeepLens
 - AWS DeepRacer
- Analytics**
 - Athena
 - EMR
 - CloudSearch
 - Elasticsearch Service
 - Kinesis
 - QuickSight [\[?\]](#)

Access resources on the go

 Access the Management Console using the AWS Console Mobile App. [Learn more \[\\[?\\]\]\(#\)](#)

Explore AWS

AWS IQ
Complete your AWS projects faster with help from AWS Certified third-party experts. [Get started \[\\[?\\]\]\(#\)](#)

Stream Live re:Invent Keynotes and Launches, Dec 2 – 6
Hear from AWS leaders, and learn about new products. [Sign up \[\\[?\\]\]\(#\)](#)

Amazon RDS
Set up, operate, and scale your relational database in the cloud. [Learn more \[\\[?\\]\]\(#\)](#)

EC2 Spot Instances
Run fault-tolerant workloads on Spot Instances and save up to 90% on compute. [Learn more \[\\[?\\]\]\(#\)](#)

Google Cloud Platform My First Project

Pins appear here

Home Marketplace Billing APIs & Services Support IAM & admin Getting started Security COMPUTE App Engine Compute Engine Kubernetes Engine Cloud Functions Cloud Run STORAGE Bigtable

DASHBOARD ACTIVITY CUSTOMISE

Project info

Project name: My First Project
Project ID: grand-brook-257114
Project number: 63092496236

ADD PEOPLE TO THIS PROJECT

Go to project settings

Resources

This project has no resources

APIs

Requests (requests/sec)

No data is available for the selected time frame.

17:15 17:30 17:45 18:00

Go to APIs overview

Trace

No trace data from the last 7 days

Get started with Stackdriver Trace

Getting started

Explore and enable APIs

Google Cloud Platform status

Google Kubernetes Engine incident no.19012
We are investigating an issue with Google Kubernetes Engine where some nodes in recently upgraded clusters (see affected versions) may be experiencing elevated numbers of kernel panics
Began at 2019-11-04 (11:46:04)
All times are US/Pacific
Data provided by status.cloud.google.com

Go to Cloud status dashboard

Error Reporting

No sign of any errors. Have you set up Error Reporting?

Learn how to set up Error Reporting

News

How to calculate distances on a map with the Maps JavaScript API
2 days ago

Kubernetes development, simplified—Skaffold is now GA
2 days ago

Updating App Engine with more new runtimes: Nodejs 12, Go 1.13, PHP 7.3 and Python 3.8

The screenshot shows the Google Cloud Platform dashboard for a project named "My First Project". The left sidebar lists various services: Home, Marketplace, Billing, APIs & Services, Support, IAM & admin, Getting started, Security, COMPUTE (App Engine, Compute Engine, Kubernetes Engine, Cloud Functions, Cloud Run), and STORAGE (Bigtable). The main content area is divided into sections: "Project info" (Project name: My First Project, Project ID: grand-brook-257114, Project number: 63092496236), "APIs" (Shows a timeline from 17:15 to 18:00 with no data available), "Resources" (Shows no resources), "Trace" (Shows no trace data from the last 7 days), and "Getting started" (Links to explore and enable APIs). On the right, there are three cards: "Google Cloud Platform status" (Shows a recent incident for Google Kubernetes Engine), "Error Reporting" (Shows no errors), and "News" (Shows three recent articles: "How to calculate distances on a map with the Maps JavaScript API", "Kubernetes development, simplified—Skaffold is now GA", and "Updating App Engine with more new runtimes: Nodejs 12, Go 1.13, PHP 7.3 and Python 3.8").

The screenshot shows the AWS EC2 Dashboard. The left sidebar contains navigation links for EC2 Dashboard, Instances, Images, Elastic Block Store, Network & Security, and Load Balancing. The main content area displays resource statistics and a 'Create Instance' section with a prominent 'Launch Instance' button highlighted by a red oval. The right sidebar includes sections for Account Attributes, Additional Information, and AWS Marketplace.

EC2 Dashboard

- Events
- Tags
- Reports
- Limits

INSTANCES

- Instances
- Launch Templates
- Spot Requests
- Savings Plans
- Reserved Instances
- Dedicated Hosts
- Scheduled Instances
- Capacity Reservations

IMAGES

- AMIs
- Bundle Tasks

ELASTIC BLOCK STORE

- Volumes
- Snapshots
- Lifecycle Manager

NETWORK & SECURITY

- Security Groups
- Elastic IPs
- Placement Groups
- Key Pairs
- Network Interfaces

LOAD BALANCING

- Load Balancers
- Target Groups

Resources

You are using the following Amazon EC2 resources in the US East (N. Virginia) region:

0 Running Instances	0 Elastic IPs
0 Dedicated Hosts	0 Snapshots
0 Volumes	0 Load Balancers
0 Key Pairs	1 Security Groups
0 Placement Groups	

Learn more about the latest in AWS Compute from AWS re:Invent by viewing the [EC2 Videos](#).

Create Instance

To start using Amazon EC2 you will want to launch a virtual server, known as an Amazon EC2 instance.

Launch Instance

Note: Your instances will launch in the US East (N. Virginia) region

Migrate a Machine

Use CloudEndure Migration to simplify, expedite, and automate large-scale migrations from physical, virtual, and cloud-based infrastructure to AWS.

[Get started with CloudEndure Migration](#)

Service Health

Service Status:

- US East (N. Virginia): ✓ Operating normally

Availability Zone Status:

- us-east-1a: Availability zone is operating normally
- us-east-1b: Availability zone is operating normally
- us-east-1c: Availability zone is operating normally

Scheduled Events

US East (N. Virginia):

No events

Account Attributes

Supported Platforms

- VPC

Default VPC

- vpc-721ba808

Console experiments

Settings

Additional Information

[Getting Started Guide](#)

[Documentation](#)

[All EC2 Resources](#)

[Forums](#)

[Pricing](#)

[Contact Us](#)

AWS Marketplace

Find free software trial products in the AWS Marketplace from the [EC2 Launch Wizard](#). Or try these popular software:

CloudEndure Migration

By Amazon Web Services

Rating ★★★★★

[View all Infrastructure Software](#)

Matillion ETL for Amazon Redshift

By Matillion

Rating ★★★★★

\$1.37 to \$5.48/hr for software + AWS usage fees

[View all Business Software](#)

Google Cloud Platform My First Project

Compute Engine VM instances

VM instances

Instance groups Instance templates Sole-tenant nodes Disks Snapshots Images TPUs Committed use discounts Metadata Health checks Zones Network endpoint groups Operations Security scans Settings Marketplace

Compute Engine VM instances

Compute Engine lets you use virtual machines that run on Google's infrastructure. Create micro-VMs or larger instances running Debian, Windows or other standard images. Create your first VM instance, import it using a migration service or try the quickstart to build a sample app.

Create Import Take the quickstart

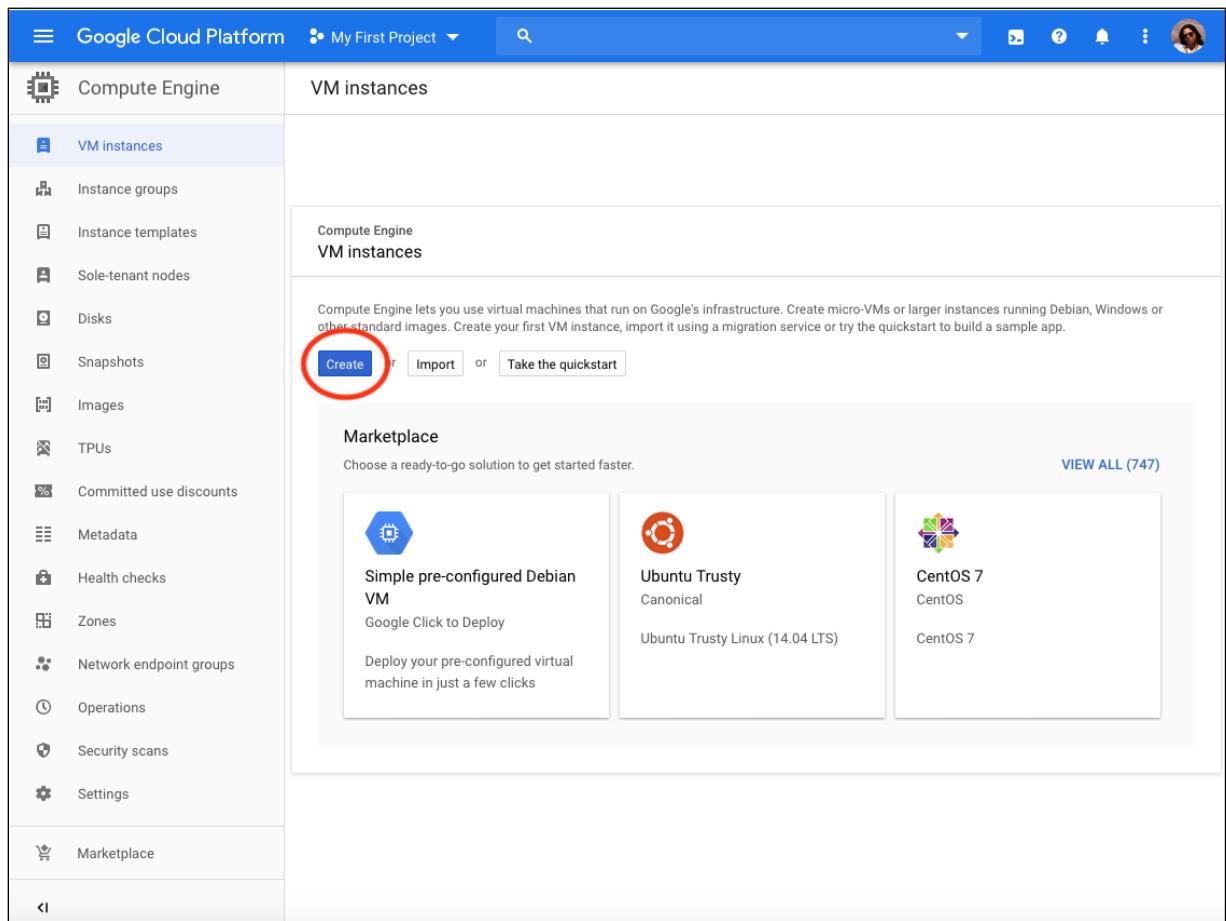
Marketplace

Choose a ready-to-go solution to get started faster. VIEW ALL (747)

Simple pre-configured Debian VM Google Click to Deploy Deploy your pre-configured virtual machine in just a few clicks

Ubuntu Trusty Canonical Ubuntu Trusty Linux (14.04 LTS)

CentOS 7 CentOS 7



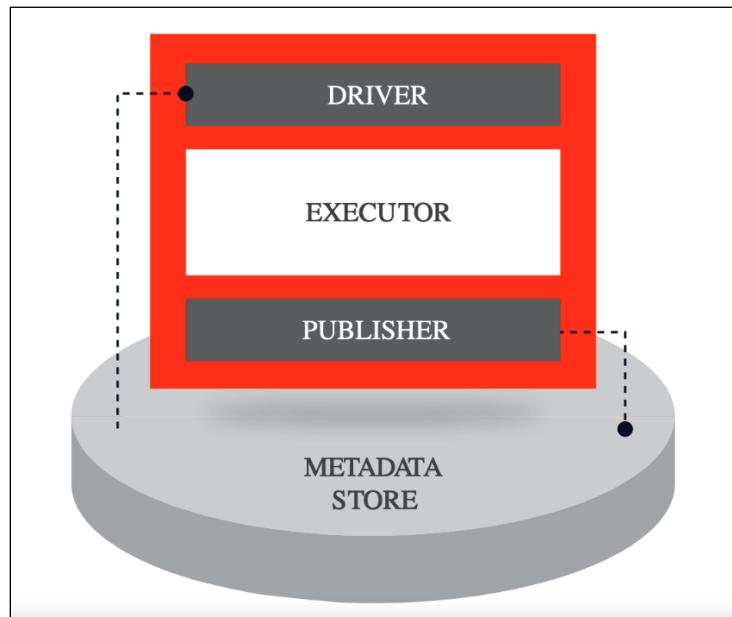
Examples Recent Google Drive GitHub Upload

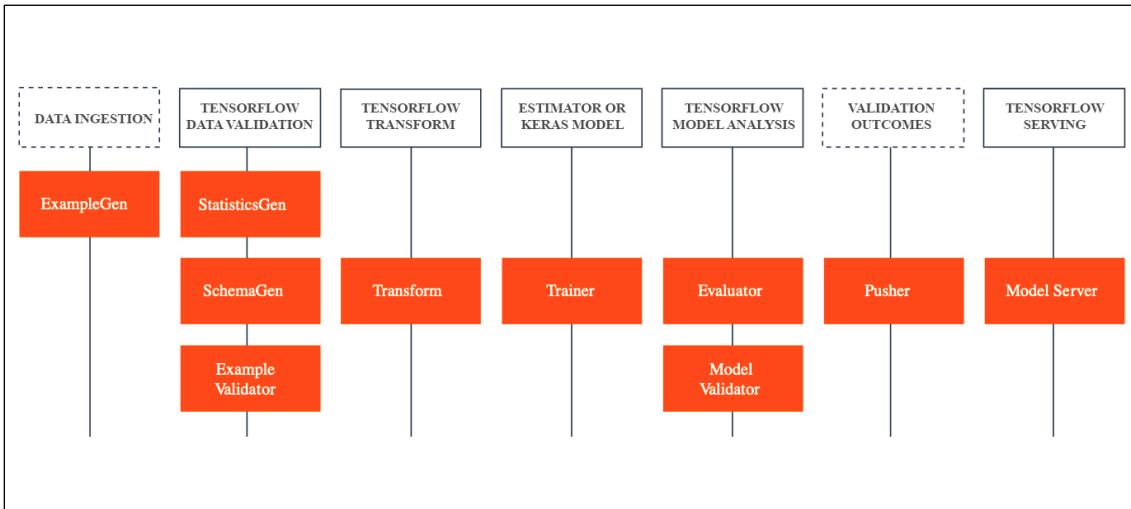
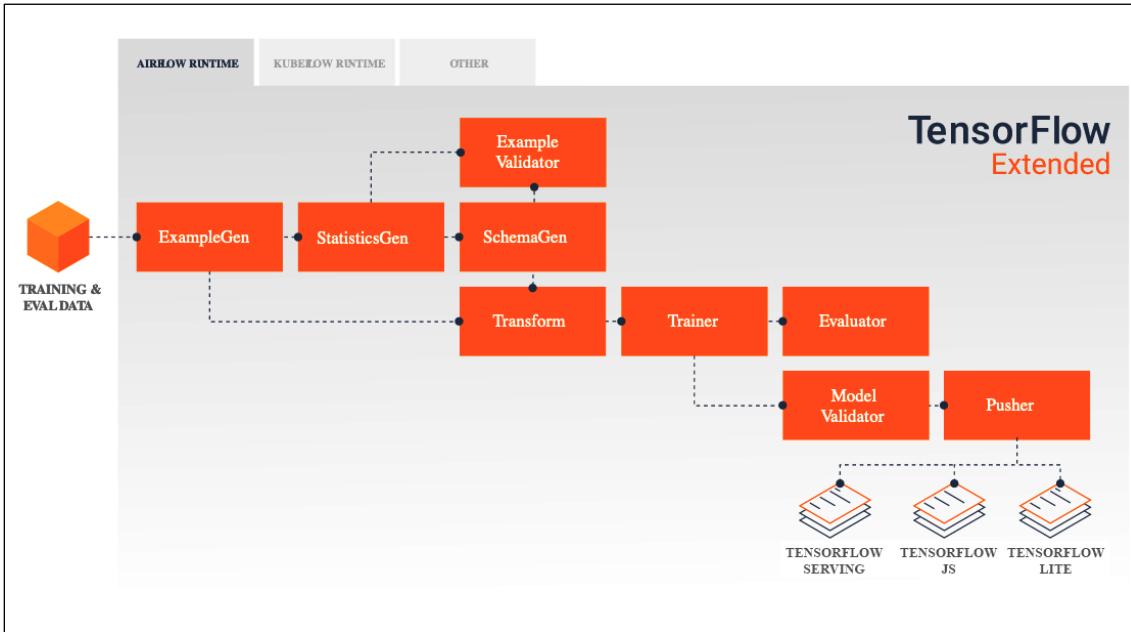
Filter notebooks

Title

 Overview of Colaboratory Features	<input type="checkbox"/>
 Markdown Guide	<input type="checkbox"/>
 Charts in Colaboratory	<input type="checkbox"/>
 External data: Drive, Sheets, and Cloud Storage	<input type="checkbox"/>
 Getting started with BigQuery	<input type="checkbox"/>

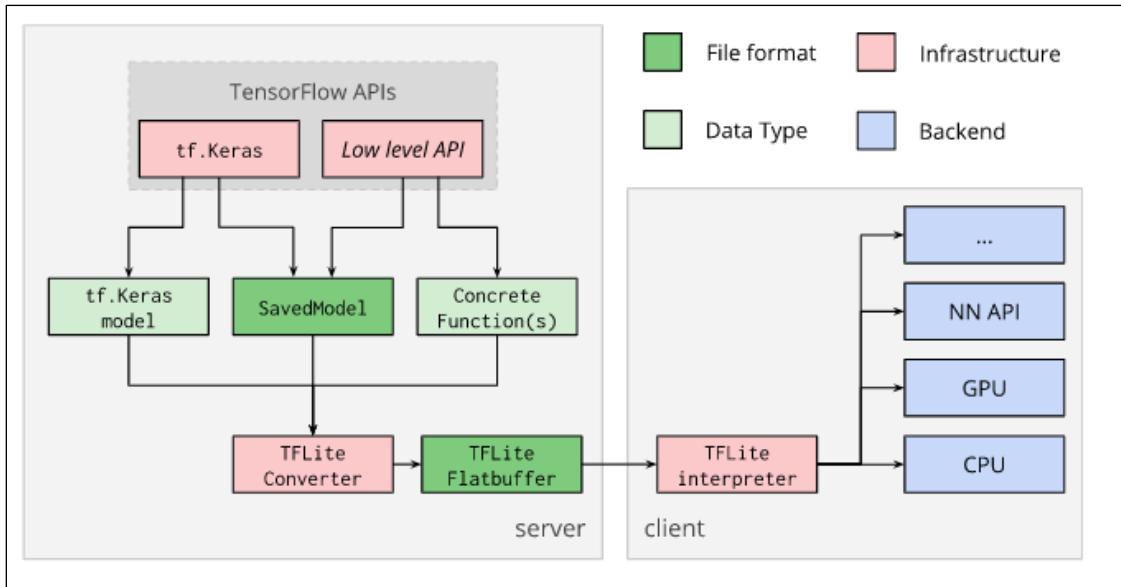
NEW PYTHON 3 NOTEBOOK ▾ CANCEL





Chapter 13: TensorFlow for Mobile and IoT and TensorFlow.js

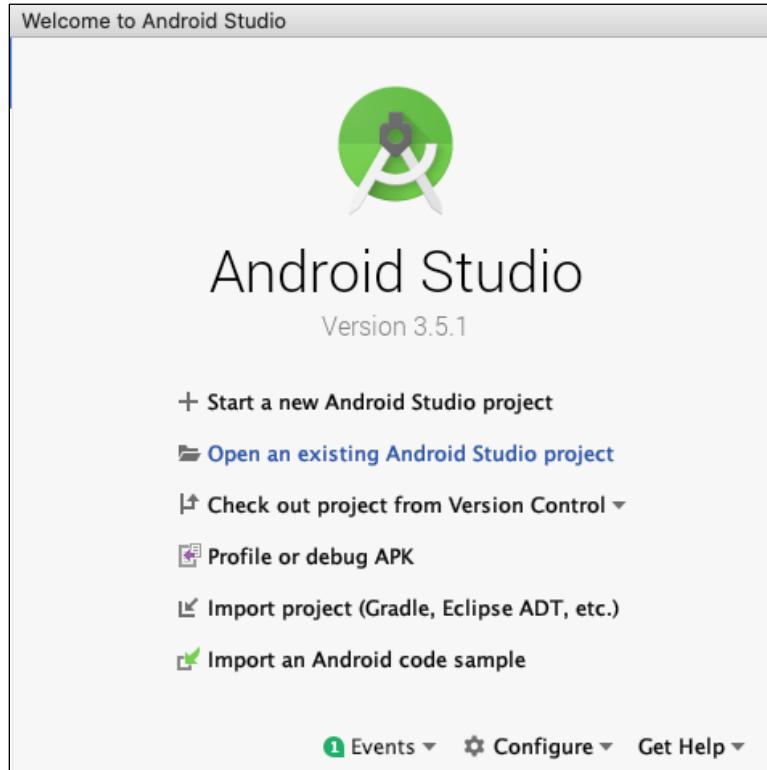
Model	Top-1 Accuracy (Original)	Top-1 Accuracy (Post Training Quantized)	Top-1 Accuracy (Quantization Aware Training)	Latency (Original) (ms)	Latency (Post Training Quantized) (ms)	Latency (Quantization Aware Training) (ms)	Size (Original) (MB)	Size (Optimized) (MB)
Mobilenet-v1-1-224	0.709	0.657	0.70	124	112	64	16.9	4.3
Mobilenet-v2-1-224	0.719	0.637	0.709	89	98	54	14	3.6
Inception_v3	0.78	0.772	0.775	1130	845	543	95.7	23.9
Resnet_v2_101	0.770	0.768	N/A	3973	2868	N/A	178.3	44.9



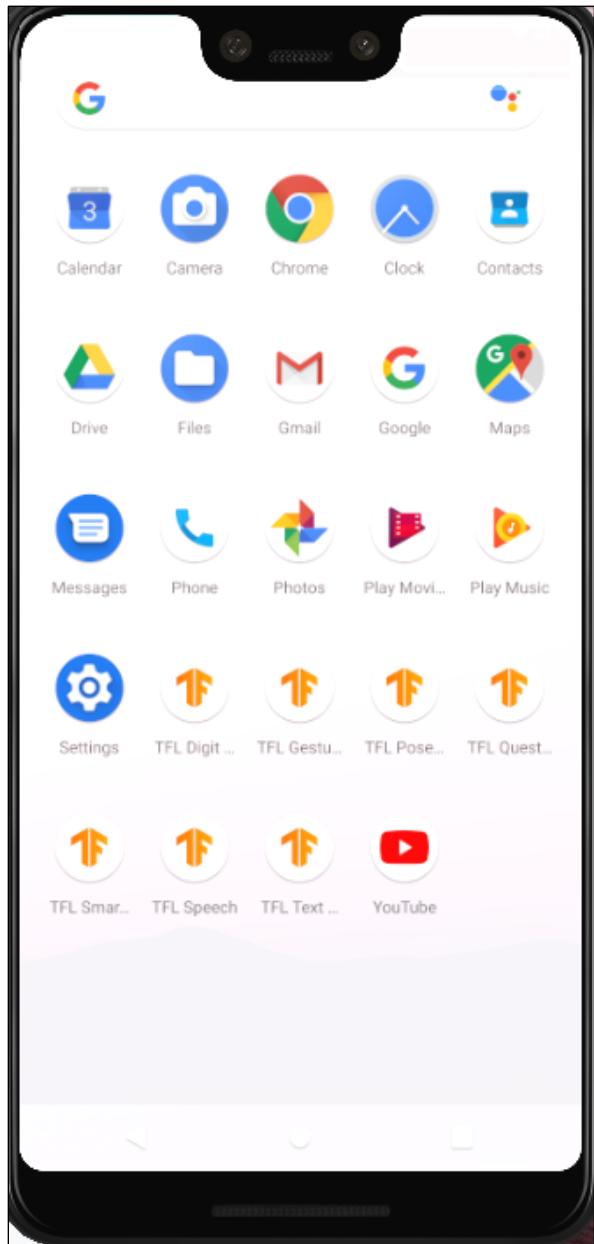
```

From-4590-back-to-2018-to-observe-the-world-before-the-big-fall:~ antonio$ sdkmanager --list
Warning: File /Users/antonio/.android/repositories.cfg could not be loaded.
Installed packages:=====
Path                                Version | Description
-----|-----|-----
add-ons;addon-google_apis-google-24   1       | Google APIs
build-tools;28.0.3                     28.0.3 | Android SDK Build-Tools 28.0.3
build-tools;29.0.2                     29.0.2 | Android SDK Build-Tools 29.0.2
emulator                            29.2.1 | Android Emulator
patcher;v4                           1       | SDK Patch Applier v4
platforms;android-28                  6       | Android SDK Platform 28
platforms;android-29                  3       | Android SDK Platform 29
system-images;android-29;google_apis_playstore;x86 8       | Google Play Intel x86 Atom System Image
tools                               26.1.1 | Android SDK Tools 26.1.1

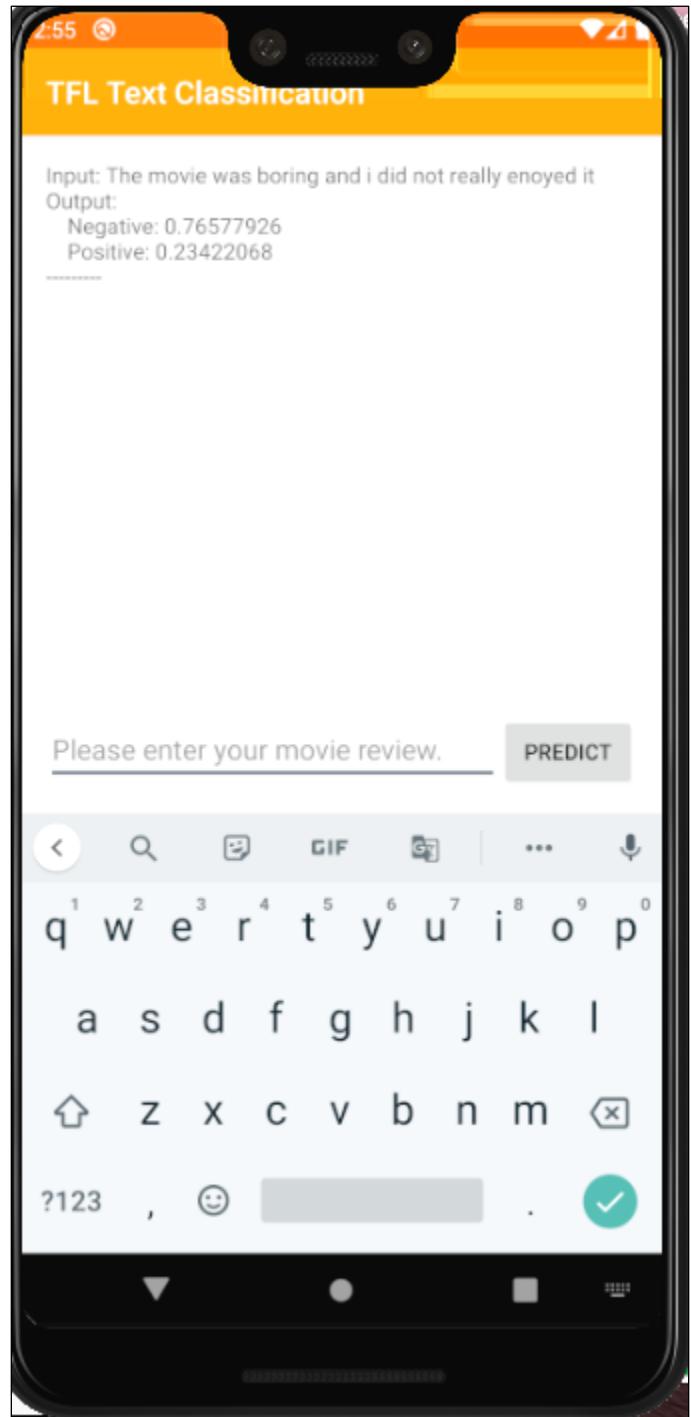
```



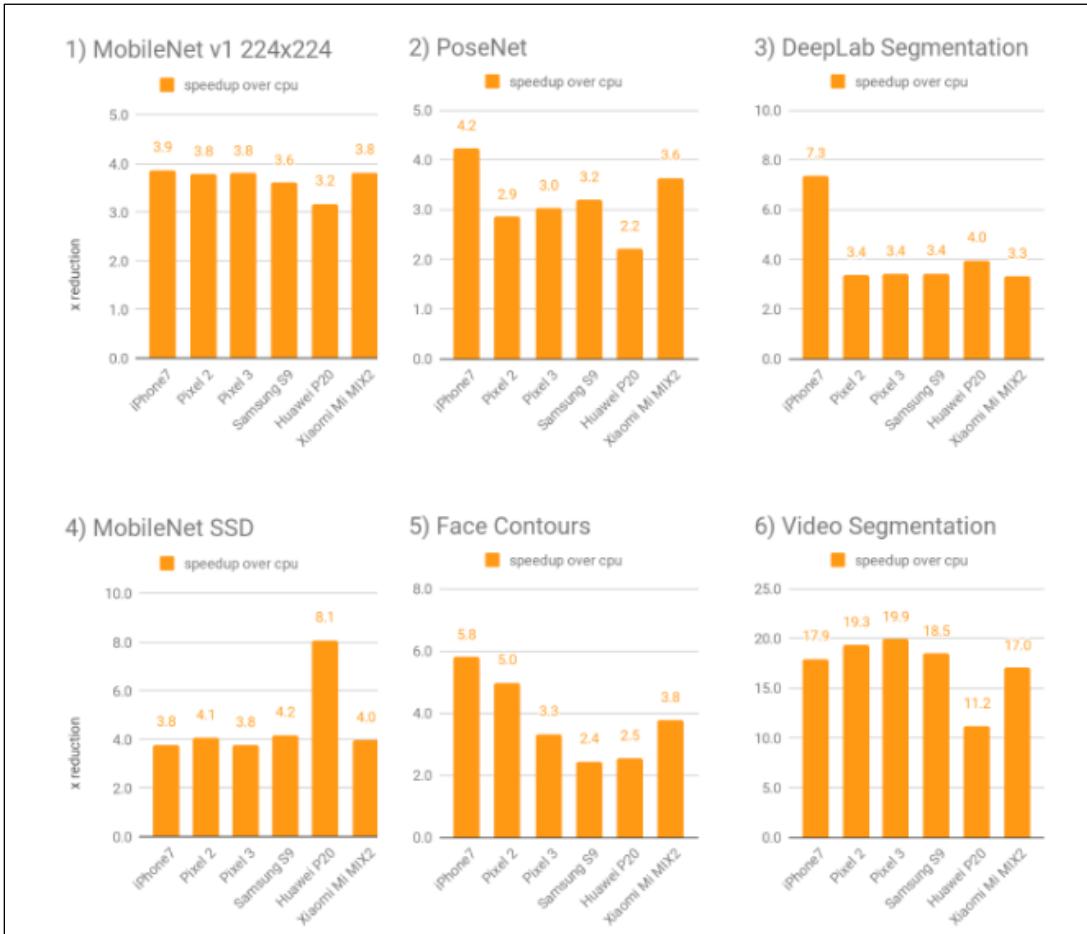
Your Virtual Devices					
Type	Name	Play Store	Resolution	API	Target
Pixel 3 XL API 29			1440 x 2960: 560dpi	29	Android 10.0 (Google...)

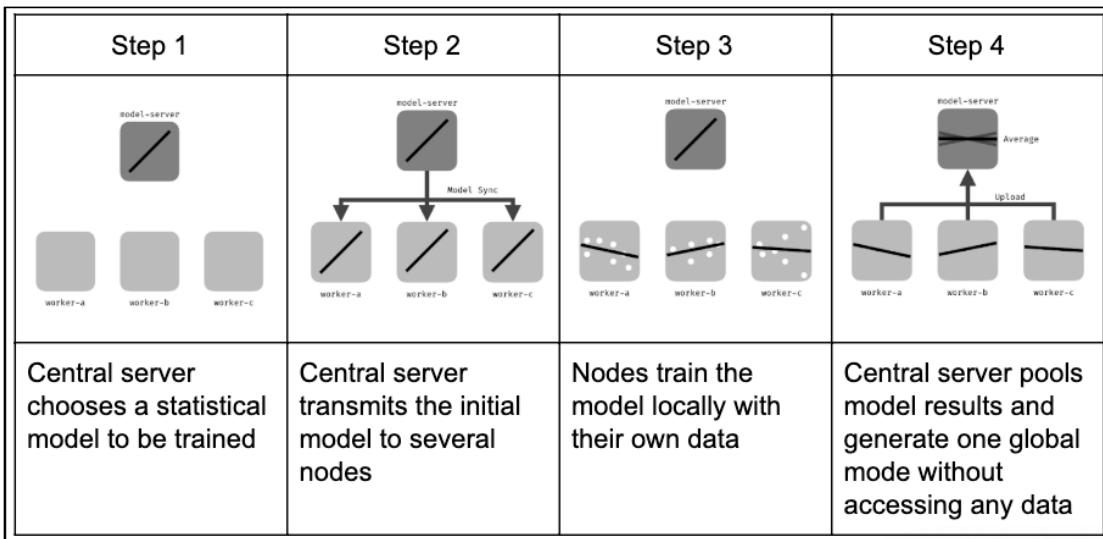
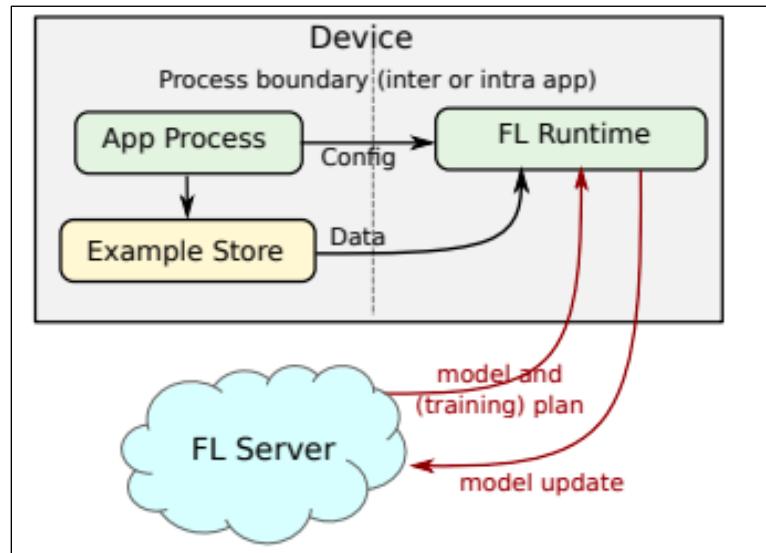


Model name	Model size	Top-1 accuracy	Top-5 accuracy	TF Lite performance
Mobilenet_V1_0.25_128_quant	0.5 Mb	39.5%	64.4%	3.7 ms
Mobilenet_V1_0.25_160_quant	0.5 Mb	42.8%	68.1%	5.5 ms
Mobilenet_V1_0.25_192_quant	0.5 Mb	45.7%	70.8%	7.9 ms
Mobilenet_V1_0.25_224_quant	0.5 Mb	48.2%	72.8%	10.4 ms
Mobilenet_V1_0.50_128_quant	1.4 Mb	54.9%	78.1%	8.8 ms
Mobilenet_V1_0.50_160_quant	1.4 Mb	57.2%	80.5%	13.0 ms
Mobilenet_V1_0.50_192_quant	1.4 Mb	59.9%	82.1%	18.3 ms
Mobilenet_V1_0.50_224_quant	1.4 Mb	61.2%	83.2%	24.7 ms
Mobilenet_V1_0.75_128_quant	2.6 Mb	55.9%	79.1%	16.2 ms
Mobilenet_V1_0.75_160_quant	2.6 Mb	62.4%	83.7%	24.3 ms
Mobilenet_V1_0.75_192_quant	2.6 Mb	66.1%	86.2%	33.8 ms
Mobilenet_V1_0.75_224_quant	2.6 Mb	66.9%	86.9%	45.4 ms
Mobilenet_V1_1.0_128_quant	4.3 Mb	63.3%	84.1%	24.9 ms
Mobilenet_V1_1.0_160_quant	4.3 Mb	66.9%	86.7%	37.4 ms
Mobilenet_V1_1.0_192_quant	4.3 Mb	69.1%	88.1%	51.9 ms
Mobilenet_V1_1.0_224_quant	4.3 Mb	70.0%	89.0%	70.2 ms
Mobilenet_V2_1.0_224_quant	3.4 Mb	70.8%	89.9%	53.4 ms
Inception_V1_quant	6.4 Mb	70.1%	89.8%	154.5 ms
Inception_V2_quant	11 Mb	73.5%	91.4%	235.0 ms
Inception_V3_quant	23 Mb	77.5%	93.7%	637 ms
Inception_V4_quant	41 Mb	79.5%	93.9%	1250.8 ms

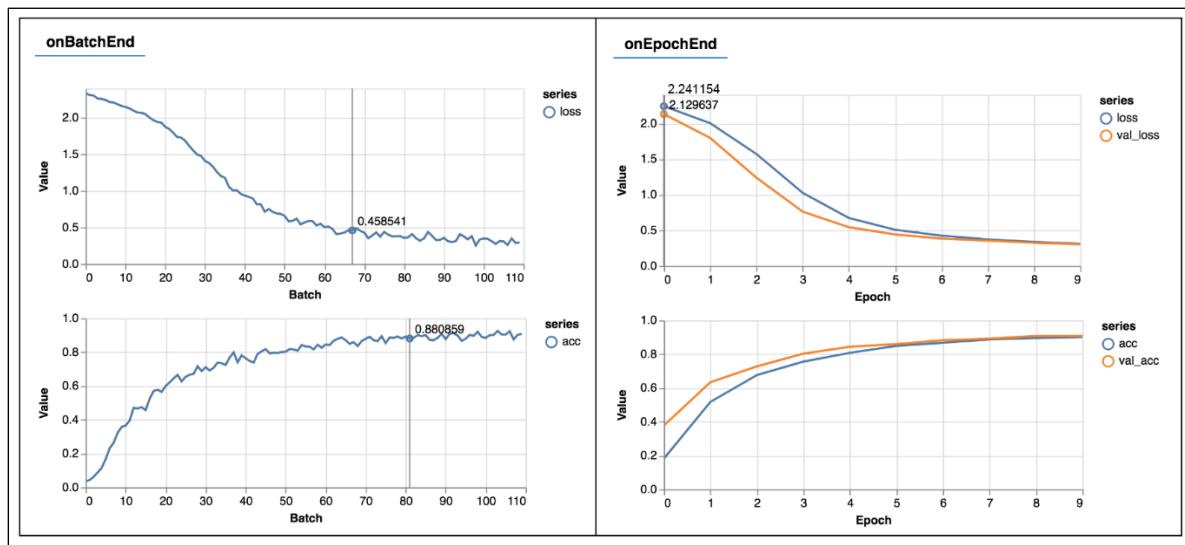


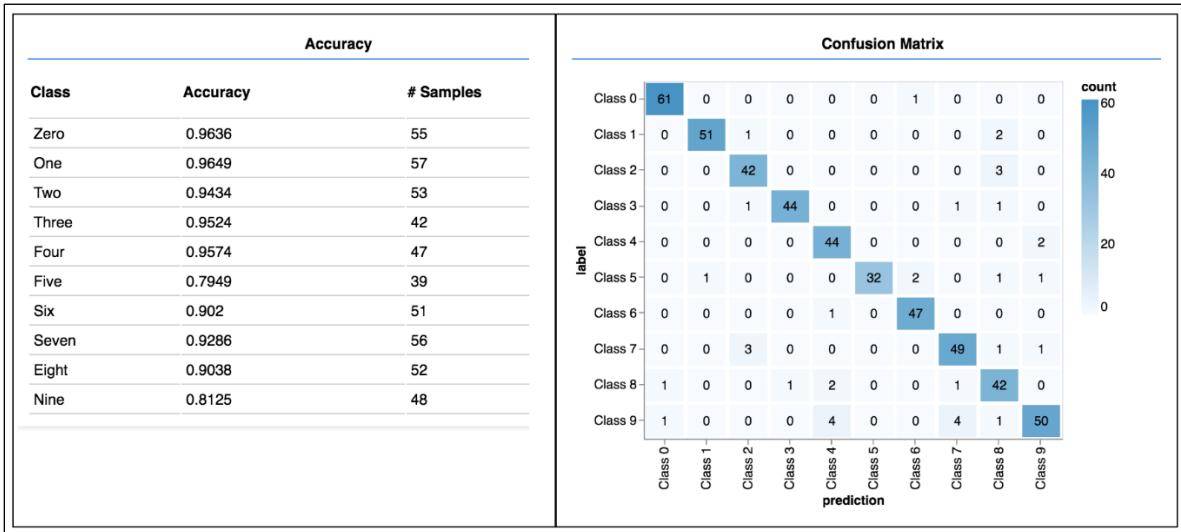




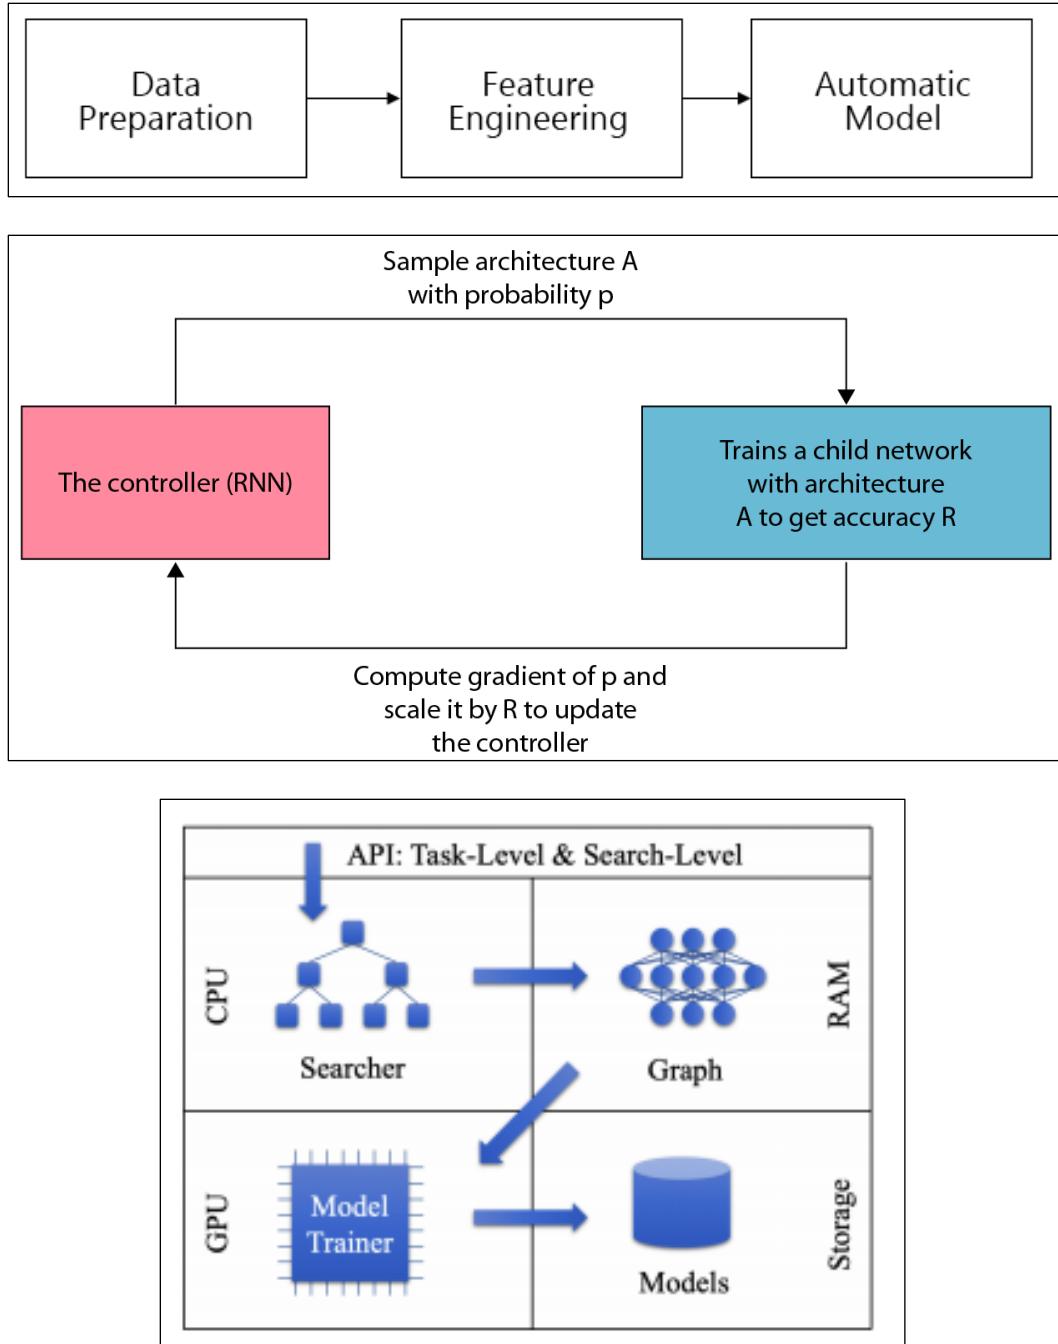


Model Architecture			
Layer Name	Output Shape	# Of Params	Trainable
conv2d_Conv2D1	[batch,24,24,8]	208	true
max_pooling2d_MaxPooling2D1	[batch,12,12,8]	0	true
conv2d_Conv2D2	[batch,8,8,16]	3,216	true
max_pooling2d_MaxPooling2D2	[batch,4,4,16]	0	true
flatten_Flatten1	[batch,256]	0	true
dense_Dense1	[batch,10]	2,570	true





Chapter 14: An introduction to AutoML



Cloud AutoML BETA

Train high-quality custom machine learning models with minimal effort and machine learning expertise.

[Try AutoML](#) ^

[View documentation](#)

AutoML Natural Language
AutoML Translation
AutoML Video Intelligence
AutoML Vision
AutoML Tables

machine learning models

learning products that enables developers expertise to train high-quality models specific to their business needs. It relies on Google's state-of-the-art transfer learning and neural architecture search technology.

Machine Learning

AutoML Tables BETA

Create supervised machine learning models with your tabular data. AutoML Tables supports a variety of data types and problem types (binary and multi-class classification; regression).

Click "Enable API" to turn on the Cloud AutoML API and start using AutoML Tables.

[ENABLING API](#)

Google Cloud Platform		authentica	NEW DATASET			
Tables	Datasets <small>BETA</small>					
Datasets	Name	Dataset source	Total columns	Total rows	Time of creation	Status
	No rows to display					
Models						

Create new dataset

Dataset name *

test_bank_marketing

Use letters, numbers and underscores up to 32 characters.

CANCEL

CREATE DATASET

← test_bank_marketing BETA

IMPORT

SCHEMA

ANALYZE

TRAIN

EVALUATE

PREDICT

Import your data

AutoML Tables uses tabular data that you import to train a custom machine learning model. Your dataset must contain at least one input feature column and a target column. Optional columns can be added to configure parameters like the data split, weights, etc. [Preparing your training data](#)

- Import data from BigQuery
- Select a CSV file from Cloud Storage
- Upload files from your computer

Select a CSV file from Cloud Storage

The bucket containing the CSV must be in the us-central1 region. [CSV formatting](#)

gs:// *



cloud-ml-tables-data/bank-marketing.csv

BROWSE

IMPORT

[test_bank_marketing](#) BETA

IMPORT SCHEMA ANALYZE TRAIN EVALUATE PREDICT

Your data is being imported

Data import can take up to one hour. You can close this window. You'll receive an email when your data is ready to use.

IMPORT	SCHEMA	ANALYZE	TRAIN	EVALUATE	PREDICT
Select a target Select a column to be the target (what you want your model to predict) and add optional parameters like weight and time columns					
Target column ? RESET		Column name ? Data type ? Nullability ? Age Numeric <input checked="" type="checkbox"/> Nullable Job Categorical <input checked="" type="checkbox"/> Nullable MaritalStatus Categorical <input checked="" type="checkbox"/> Nullable Education Categorical <input checked="" type="checkbox"/> Nullable Default Categorical <input checked="" type="checkbox"/> Nullable Balance Numeric <input checked="" type="checkbox"/> Nullable Housing Categorical <input checked="" type="checkbox"/> Nullable Loan Categorical <input checked="" type="checkbox"/> Nullable Contact Categorical <input checked="" type="checkbox"/> Nullable Day Numeric <input checked="" type="checkbox"/> Nullable Month Categorical <input checked="" type="checkbox"/> Nullable Duration Numeric <input checked="" type="checkbox"/> Nullable Campaign Numeric <input checked="" type="checkbox"/> Nullable PDays Numeric <input checked="" type="checkbox"/> Nullable Previous Numeric <input checked="" type="checkbox"/> Nullable POutcome Categorical <input checked="" type="checkbox"/> Nullable Deposit <input checked="" type="checkbox"/> Target Categorical <input checked="" type="checkbox"/> Nullable			
The selected column is categorical data. AutoML Tables will build a classification model, which will predict the target from the classes in the selected column. Learn more					
Additional parameters (Optional) Before continuing, review your dataset schema to make sure each column has the appropriate data type and nullability setting					
CONTINUE					

IMPORT	SCHEMA	ANALYZE	TRAIN	EVALUATE	PREDICT
⚠ Not up to date. Click the "Continue" button on the Schema tab to regenerate statistics.					
All features	17	Feature name ↑	Type	Missing ⓘ	Distinct values ⓘ
		Age	Numeric	0% (0)	77
		Balance	Numeric	0% (0)	7,168
		Campaign	Numeric	0% (0)	48
		Contact	Categorical	0% (0)	3
		Day	Numeric	0% (0)	31
		Default	Categorical	0% (0)	2
		Deposit Target	Categorical	0% (0)	2
		Duration	Numeric	0% (0)	1,573
		Education	Categorical	0% (0)	4
		Housing	Categorical	0% (0)	2
		Job	Categorical	0% (0)	12
		Loan	Categorical	0% (0)	2
		MaritalStatus	Categorical	0% (0)	3
		Month	Categorical	0% (0)	12
		PDays	Numeric	0% (0)	559
		POutcome	Categorical	0% (0)	4
		Previous	Numeric	0% (0)	41
Rows per page: 50 ▾ 1 – 17 of 17 ⏪ ⏩					

IMPORT SCHEMA ANALYZE **TRAIN** EVALUATE PREDICT

Train your model

Model name * test_bank_marketi_20190913073044

Training budget

Enter a number between 1 and 72 for the maximum number of node hours to spend training your model. If your model stops improving before then, AutoML Tables will stop training and you'll only be charged for the actual node hours used. [Training pricing guide](#)

Budget * 1 maximum node hour 

Input feature selection

By default, all other columns in your dataset will be used as input features for training (excluding target, weight, and split columns).

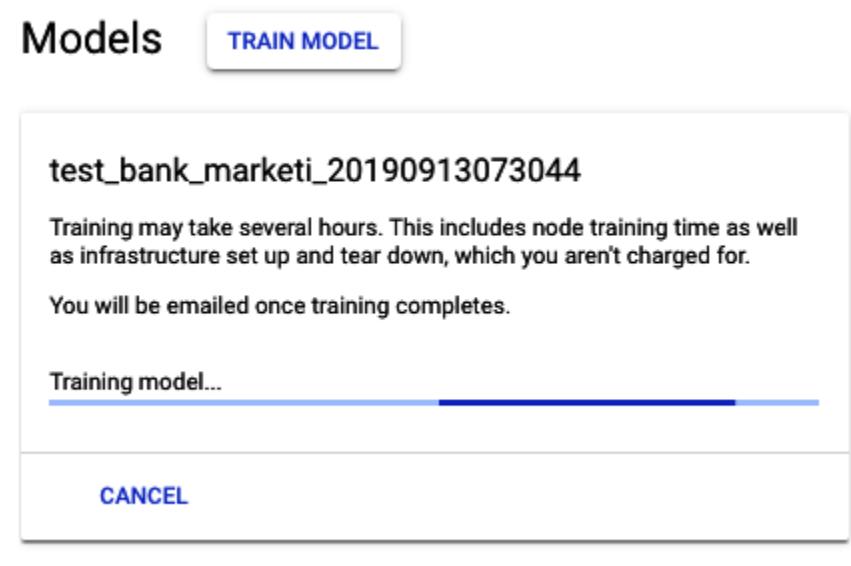
16 feature columns * All columns selected 

Summary

Model type: Binary classification model
Data split: Automatic
Target: Deposit
Input features: 16 features
Rows: 45,211 rows

Advanced options ▾

TRAIN MODEL



AutoML Tables finished training model "test_bank_marketi_20190913073044"

AutoML Tables <noreply-automl-tables@google.com>
to me ▾

Hello AutoML Tables Customer,

AutoML Tables finished training model "test_bank_marketi_20190913073044".

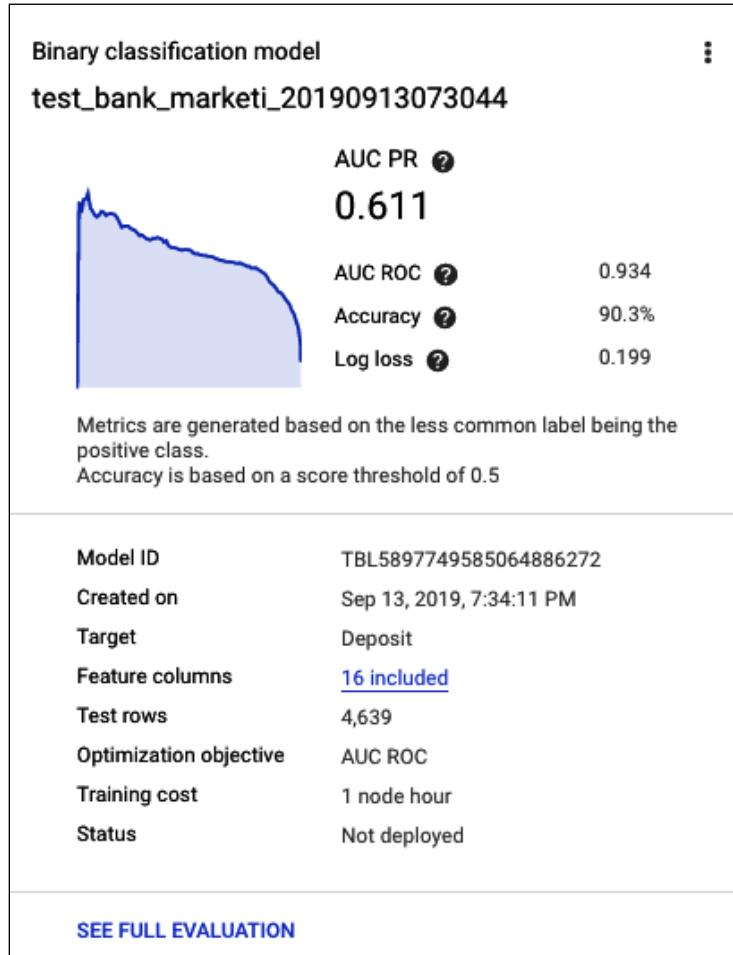
Additional Details:

Resource Name:
projects/655848112025/locations/us-central1/models/TBL897749585064886272

Operation State: Succeeded

To continue your progress, go back to your model using
<https://console.cloud.google.com/automl-tables/datasets/TBL8775197903233744896/train?project=655848112025>

Sincerely,
The Google Cloud AI Team

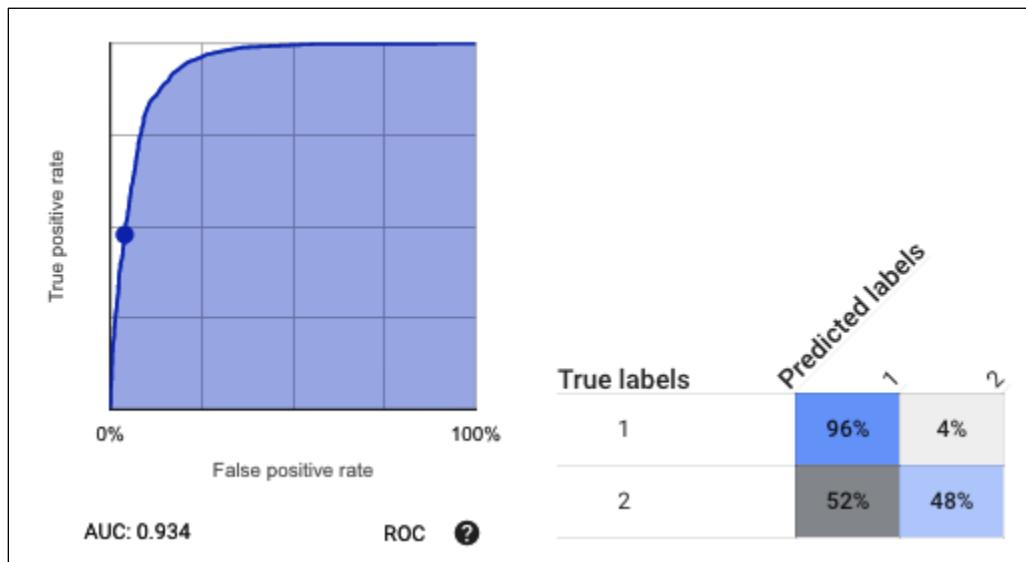


Model
test_bank_marketi_20190913073044 ▾

Binary classification model
Sep 13, 2019, 7:34:11 PM
Training cost: 1 node hour

Target	Feature columns	Optimized for	AUC PR ?	AUC ROC ?	Accuracy ?	Log loss ?
Deposit	16 included 4,639 test rows	AUC ROC	0.611	0.934	90.3%	0.199

Metrics are generated using the least-common class as the positive class. Accuracy based on score threshold of 0.5



IMPORT SCHEMA ANALYZE TRAIN EVALUATE **PREDICT**

BATCH PREDICTION ONLINE PREDICTION

Model test_bank_marketi_20190913073044

Deploying model...

Execute the request

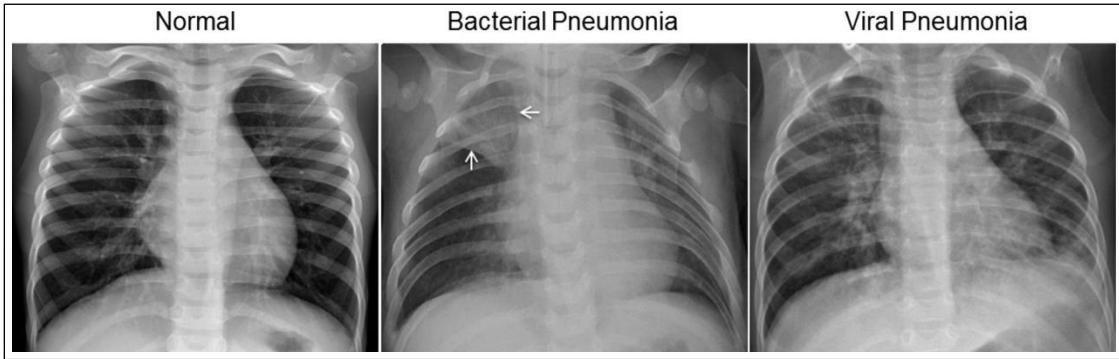
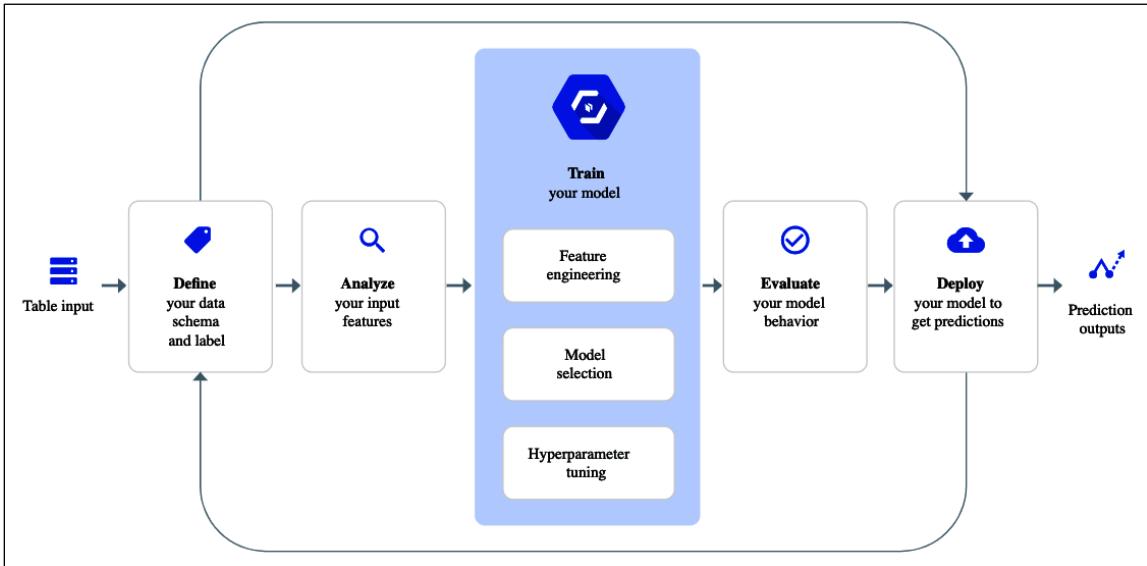
```
$ curl -X POST -H "Content-Type: application/json" \
-H "Authorization: Bearer $(gcloud auth application-default print-access-token)" \
https://automl.googleapis.com/v1beta1/projects/655848112025/locations/us-central1/models/test_bank_marketi_20190913073044/instances:predict \
-d @request.json
```

Access your model through a REST API

request.json

```
{  
  "payload": {  
    "row": {  
      "values": [  
        "39",  
        "admin.",  
        "married",  
        "secondary",  
        "no",  
        "70",  
        "yes",  
        "no",  
        "cellular",  
        "31",  
        "jul",  
        "13",  
        "11",  
        "-1",  
        "0",  
        "unknown"  
      ],  
      "columnSpecIds": [  
        "3086500662981165056",  
        "8274647433711976448",  
        "4815882919891435520",  
        "204196901464047616",  
        "5968804424498282496",  
        "3230615851057020928",  
        "7842301869484408832",  
        "2077694346450173952",  
        "4383537355663867904",  
        "6689380364877561856",  
        "8995223374091255808",  
        "7121725929105129472",  
        "2510039910677741568",  
        "5392343672194859008",  
        "780657653767471104",  
        "3662961415284588544"  
      ]  
    }  
  }  
}
```

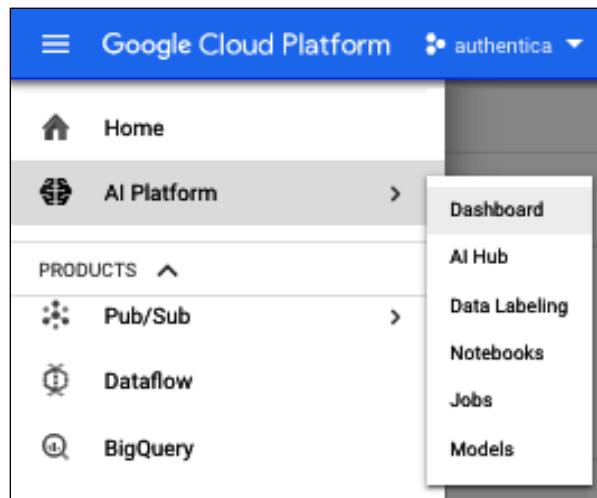
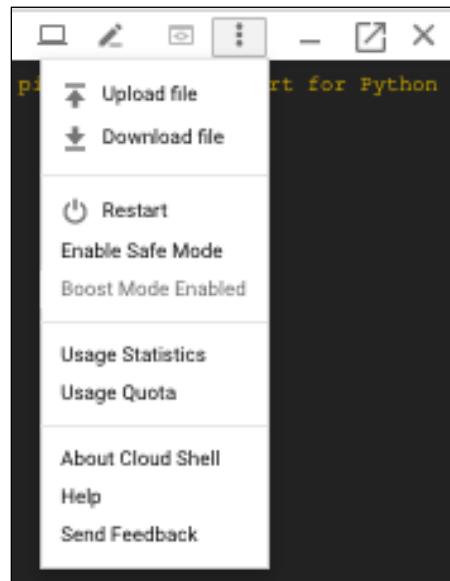
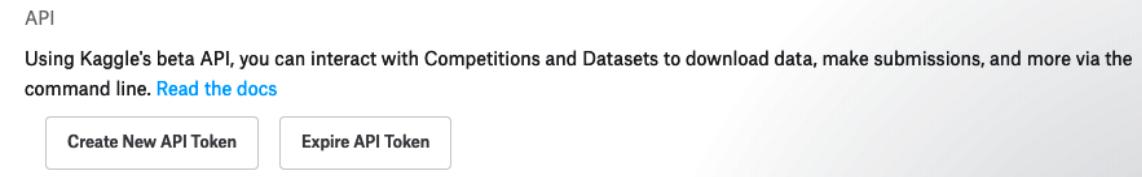
IMPORT	SCHEMA	ANALYZE	TRAIN	EVALUATE	PREDICT																																																																		
Predict label Deposit				Prediction result 1 Confidence score: 0.999 2 Confidence score: 0.001																																																																			
<table border="1"> <thead> <tr> <th>Feature column name</th> <th>Data type</th> <th>Status ↓</th> <th colspan="3">Value</th> </tr> </thead> <tbody> <tr> <td>Age</td> <td>Numeric</td> <td>Required</td> <td colspan="3">39</td> </tr> <tr> <td>Balance</td> <td>Numeric</td> <td>Required</td> <td colspan="3">70</td> </tr> <tr> <td>Campaign</td> <td>Numeric</td> <td>Required</td> <td colspan="3">11</td> </tr> <tr> <td>Contact</td> <td>Categorical</td> <td>Required</td> <td colspan="3">cellular</td> </tr> <tr> <td>Day</td> <td>Numeric</td> <td>Required</td> <td colspan="3">31</td> </tr> <tr> <td>Default</td> <td>Categorical</td> <td>Required</td> <td colspan="3">no</td> </tr> <tr> <td>Duration</td> <td>Numeric</td> <td>Required</td> <td colspan="3">13</td> </tr> <tr> <td>Education</td> <td>Categorical</td> <td>Required</td> <td colspan="3">secondary</td> </tr> <tr> <td>Housing</td> <td>Categorical</td> <td>Required</td> <td colspan="3">yes</td> </tr> <tr> <td>Job</td> <td>Categorical</td> <td>Required</td> <td colspan="3">admin.</td> </tr> </tbody> </table>						Feature column name	Data type	Status ↓	Value			Age	Numeric	Required	39			Balance	Numeric	Required	70			Campaign	Numeric	Required	11			Contact	Categorical	Required	cellular			Day	Numeric	Required	31			Default	Categorical	Required	no			Duration	Numeric	Required	13			Education	Categorical	Required	secondary			Housing	Categorical	Required	yes			Job	Categorical	Required	admin.		
Feature column name	Data type	Status ↓	Value																																																																				
Age	Numeric	Required	39																																																																				
Balance	Numeric	Required	70																																																																				
Campaign	Numeric	Required	11																																																																				
Contact	Categorical	Required	cellular																																																																				
Day	Numeric	Required	31																																																																				
Default	Categorical	Required	no																																																																				
Duration	Numeric	Required	13																																																																				
Education	Categorical	Required	secondary																																																																				
Housing	Categorical	Required	yes																																																																				
Job	Categorical	Required	admin.																																																																				
Rows per page: 10 ▾ 1 – 10 of 16 < >																																																																							
PREDICT RESET																																																																							



AutoML Vision	
<h3>Image Classification <small>BETA</small></h3> <p>Train a custom model to classify images, then deploy it to the cloud or on the edge. Learn more</p>	<h3>Object Detection <small>BETA</small></h3> <p>Train a custom model to detect objects in an image with bounding boxes and labels, then deploy it to the cloud or on the edge. Learn more</p>
<p>→ Get started</p>	<p>→ Get started</p>
Vision API	
<h3>Vision API</h3> <p>Use Google's pre-trained models to assign labels to images and classify them into millions of predefined categories. Detect objects and faces, read printed and handwritten text, and more.</p>	<h3>Vision Product Search</h3> <p>Use Google's pre-trained models to create engaging mobile experiences that match user photos to items in your product catalog and return a list of visually similar results.</p>
<p>→ View docs</p>	<p>→ View docs</p>

The screenshot shows the AutoML Vision interface. At the top, there is a navigation bar with the AutoML Vision logo, a 'BETA' label, and a 'NEW DATASET' button. Below the navigation bar, there is a sidebar with three icons: a list icon, a lightbulb icon, and a search icon. The main content area is titled 'Datasets'. At the bottom of the screen, there is a blue footer bar with several icons: a gear, a question mark, a bell, a three-dot menu, and a user profile picture. A 'Activate Cloud Shell' button is also present in the footer.

```
[root@cloudshell-... authentica-de791]# 81kB 10.3MB/s
Building wheels for collected packages: kaggle, python-slugify
Building wheel for kaggle (setup.py) ... done
Created wheel for kaggle: filename=kaggle-1.5.5-cp27-none-any.whl size=71896 sha256=7ddd36303fe62d7aa432f69c2f622947ac56981f8176840bb0a1f17ae150led2
Stored in directory: /root/.cache/pip/wheels/db/6a/80/6cd1892ab9b9b13633db3c74e16cba4e17ec2700f51541f06
Building wheel for python-slugify (setup.py) ... done
Created wheel for python-slugify: filename=python_slugify-3.0.3-py2.py3-none-any.whl size=4789 sha256=7ee1c37428cce2b858d7a191fdc0c694b872d7c0541c614d1a074de5dcbbfb2b
Stored in directory: /root/.cache/pip/wheels/0f/96/ca/05fb01165975402d1e37f8dd346df0d0c39e1d0761bd17bb
Successfully built kaggle python-slugify
Installing collected packages: urllib3, python-dateutil, tqdm, text-unidecode, python-slugify, kaggle
  Found existing installation: urllib3 1.25.3
    Uninstalling urllib3-1.25.3:
      Successfully uninstalled urllib3-1.25.3
Successfully installed kaggle-1.5.5 python-dateutil-2.8.0 python-slugify-3.0.3 text-unidecode-1.2 tqdm-4.35.0 urllib3-1.24.3
a_gulli@cloudshell-... authentica-de791]# sudo pip install kaggle
```



NEW INSTANCE **REFRESH** **START** **STOP**

Customize instance

R 3.6
R 3.6 and key libraries pre-installed

Python
Python 2 and 3 with Pandas, SciKit Learn and other key packages pre-installed

TensorFlow 1.14
TensorFlow 1.14 pre-installed with support for Keras

TensorFlow 2.0 [EXPERIMENTAL]
TensorFlow 2.0 pre-installed with support for Keras

Pytorch 1.1
PyTorch 1.1 pre-installed

RAPIDS XGboost [EXPERIMENTAL]
XGboost optimized for NVIDIA GPUs

CUDA 10.1
Optimized for NVIDIA GPUs

	Instance name	Region	Environment	Machine type	GPUs	Permission	Labels
<input type="checkbox"/>	<input checked="" type="radio"/> tensorflow-20190914-091341	us-west1-b		4 vCPUs, 15 GB RAM	None	Service account	No labels

instance name

tensorflow-20190914-091341 **OPEN JUPYTERLAB**

Kernel Git Tabs Settings Help

Launcher

Create dataset

Dataset name
chestXrays ?

Import images

To build a custom model, you first need to import a set of images to train it. Generally the more images the better. Each image should be categorized with a label (labels are essential for telling the model how to identify an image).

Processed images will be stored on Cloud Storage.

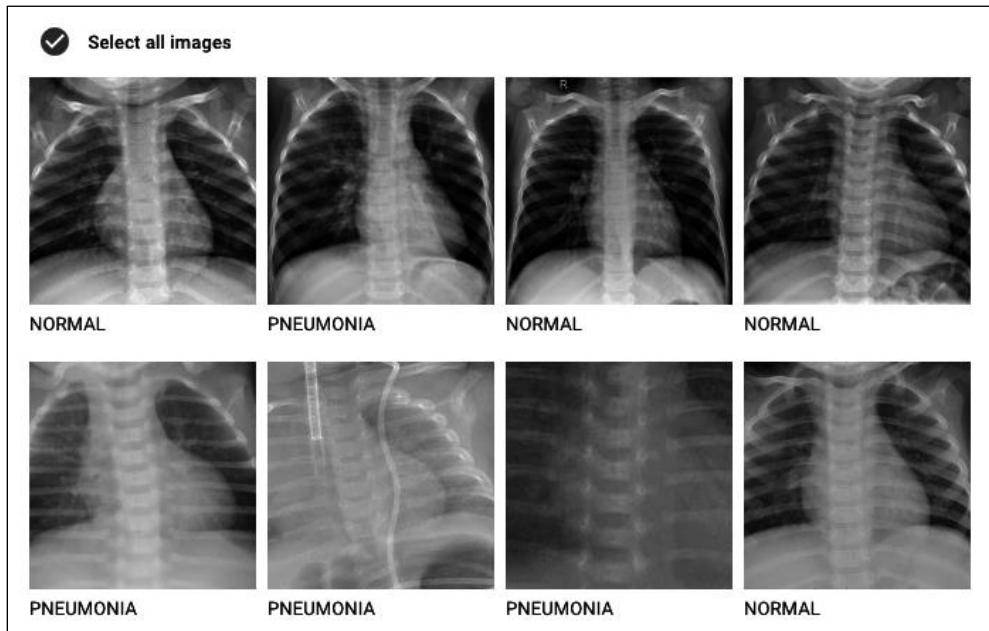
Upload images from your computer ?
Supports JPG, PNG, ZIP.

SELECT FILES

Select a CSV file on Cloud Storage ?
The **CSV file** should be a list of paths to your images on GCS and their labels, if available.
gs://authentica-de791-vcm/data.csv

Import images later
In the next step, you can add images and label them

IMAGES	TRAIN	EVALUATE	PREDICT
Importing images			
<div style="height: 10px; background-color: blue; width: 100%;"></div>			
CANCEL			



IMAGES	TRAIN	EVALUATE	PREDICT												
You have enough images to start training															
At least 100 images are currently assigned to each label. Learn more															
<table border="1"> <tr> <td>0</td> <td>100</td> <td>500</td> <td>1,000</td> </tr> <tr> <td>NORMAL</td> <td colspan="3">1340</td> </tr> <tr> <td>PNEUMONIA</td> <td colspan="3">3850</td> </tr> </table>				0	100	500	1,000	NORMAL	1340			PNEUMONIA	3850		
0	100	500	1,000												
NORMAL	1340														
PNEUMONIA	3850														
Your images will be automatically split into training and test sets , so you can evaluate your model's performance. Unlabeled images will not be used.															
Training images	4160														
Validation images	544														
Test images	486														
<input type="button" value="START TRAINING"/>															

Train new model

Model name

chestXrays_v20190914150213

Model type

Cloud-hosted

Host your model on Google Cloud for online predictions.

Edge

Download your model for offline/mobile use. Typically has lower accuracy than Cloud-hosted models.

Training budget

Your model's accuracy generally depends on how long you allow it to train, and the quality of your dataset. Your model automatically stops training when it stops improving. You pay only for the node hours used.

1 node hour (free*)



Data summary

5190 labeled images, 2 labels

* Your first node hour is free, for up to 10 models each month. [Pricing guide](#)

CANCEL

START TRAINING

IMAGES

TRAIN

EVALUATE

PREDICT

Training vision classification model

Training can take 15 minutes to several hours or more, depending on the compute hours assigned. In the meantime, you can close this window. You will be emailed once training completes.

[CANCEL](#)

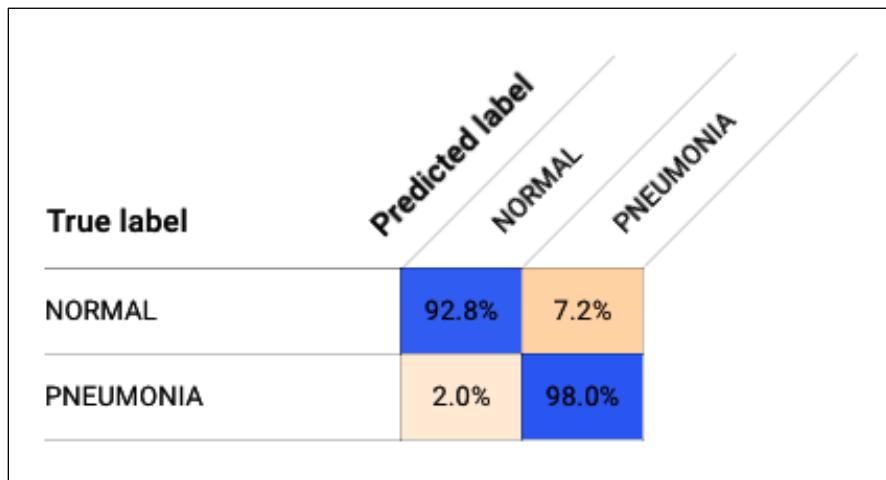
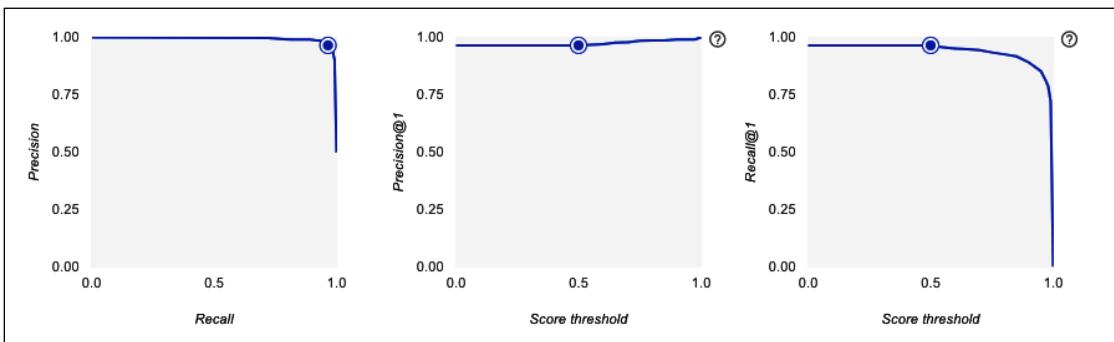
Models [TRAIN NEW MODEL](#)

chestXrays_v20190914150213

Created Sep 14, 2019 1 compute hour	Analyzed 5190 images 2 labels, 486 test images	Avg precision ⓘ 0.992	Precision ⓘ 96.502%	Recall ⓘ 96.502%
--	---	---------------------------------	-------------------------------	----------------------------

Precision and recall are based on a score threshold of 0.5

[SEE FULL EVALUATION](#) [RESUME TRAINING](#) ⓘ [...](#)



Dataset

Chest X-Ray Images (Pneumonia)

5,863 images, 2 categories

Paul Mooney · updated a year ago (Version 2)

Data Kernels (155) Discussion (12) Activity Metadata Download (1 GB) New Notebook :

Public Your Work Favorites Sort by Hotness

Outputs Languages Types Tags Search notebooks

154 @ Beating everything with Depthwise Convolution 1y ago · GPU · advanced, deep learning

50 @ Best Score on Kaggle , 96 % Recall 1y ago · cnn, image data, transfer learning

ARTIFICIAL INTELLIGENCE

- Data Labeling
- AI Platform
- Natural Language
- Tables
- Talent Solution
- Translation
- Vision
- Video Intelligence

Natural Language products

AutoML Text Classification BETA

Build a machine learning model to classify content into a custom set of categories. [Learn more](#)

→ Launch app

AutoML Sentiment Analysis BETA

Build a machine learning model to analyze attitudes within text. [Learn more](#)

→ Launch app

AutoML Entity Extraction BETA

Build a machine learning model to recognize a custom set of entities within text. [Learn more](#)

→ Get started

Cloud Natural Language API

Use Google's proven pre-trained model for general content classification; sentiment analysis; entity recognition; and more.

→ View API docs

Dataset name
happiness ?

Objective



Single-label classification
Predict the **one** correct label that you want assigned to a document.



Multi-label classification
Predict all the correct labels that you want assigned to a document.



Sentiment analysis
Understand the overall sentiment expressed in a block of text.

Upload a CSV file from your computer ?

The CSV file should be a list of GCS paths (or the text itself) and their labels, if available.

SELECT FILES

Upload text items from your computer ?

Supports .TXT, .ZIP.

happiness.csv X

SELECT FILES

Please select at least one file to upload.

Select a CSV file on Cloud Storage ?

The **CSV file** should be a list of GCS paths (or the text itself) and their labels, if available.

`gs://authentica-de791-lcm/`

Import text items later

Build your set of text items, and label directly in the workspace.

CREATE DATASET **CANCEL**

TEXT ITEMS	TRAIN	EVALUATE	PREDICT
All texts	12663	Type to filter text items...	
Labeled	12663	<input type="checkbox"/> Text	Label
Unlabeled	0	<input type="checkbox"/> I finished all of my work by the end of the day.	achievement
<input type="text"/> Type to filter... ::		<input type="checkbox"/> An event that made me happy in the past 24 hours is getting free breakfast.	enjoy_the_moment
achievement	3931	<input type="checkbox"/> When I managed to get my custom PC up and running for the first time.	achievement
affection	4337	<input type="checkbox"/> My mother flew out of town to visit our family in KS. I was so happy to see her off on the plane and I could feel the joy she must have felt upon her way out there.	affection
bonding	1584	<input type="checkbox"/> > Nowadays, happiness is a fuzzy concept and can mean many different things to many people. Part of the challenge of a science of happiness is to identify different concepts o...	enjoy_the_moment
enjoy_the_moment	1380	<input type="checkbox"/> I was given a free dessert at a restaurant.	enjoy_the_moment
exercise	196	<input type="checkbox"/> I was nominated for an award.	achievement
leisure	986		
nature	249		
Add label			

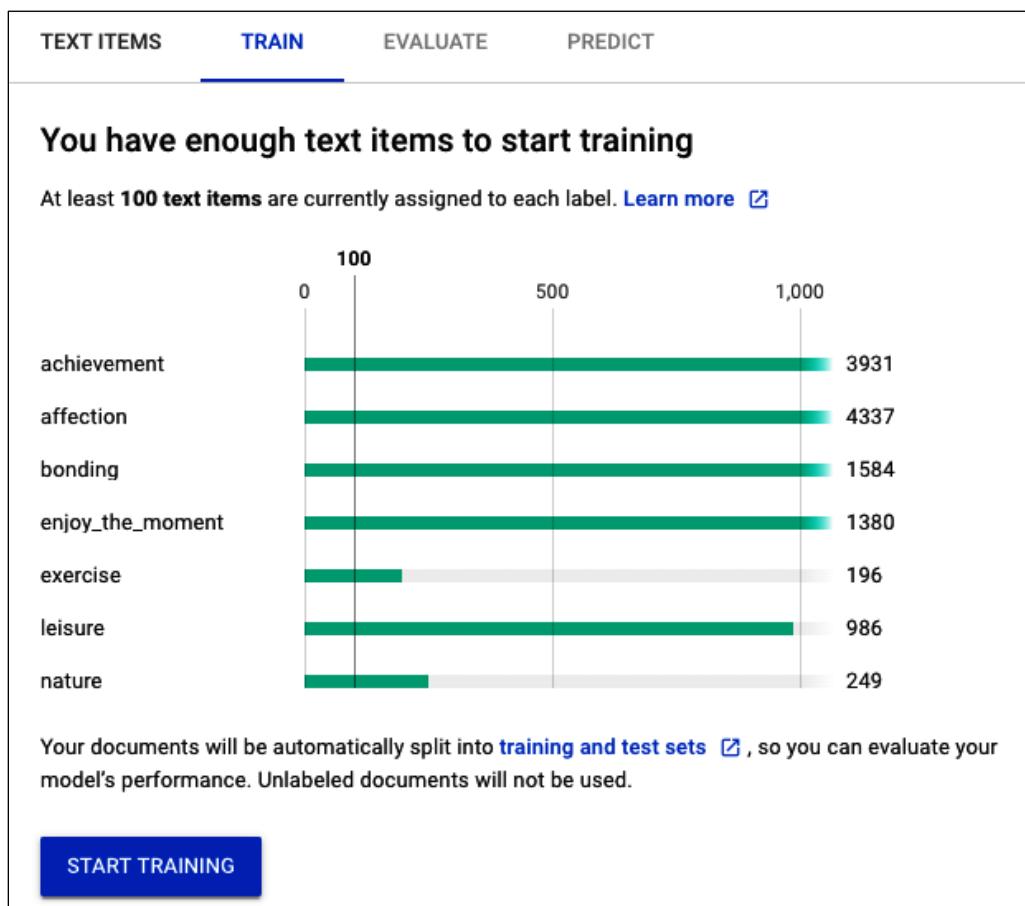
Train new model

Model name
happiness_v20190914210031

Data summary
12663 labeled text items, 7 labels

You will be emailed when training completes. [Pricing guide](#)

[CANCEL](#) [START TRAINING](#)



TEXT ITEMS TRAIN EVALUATE PREDICT

Training text model

Training can take several hours or more, depending on the complexity of your dataset. In the meantime, you can close this window. You will be emailed once training completes.

[CANCEL](#)

Models [TRAIN NEW MODEL](#)

happiness_v20190914210031			
Created Sep 15, 2019 02:10 AM	Analyzed 12663 text items 7 labels, 1266 test text items	Avg precision ⓘ 0.94	Precision ⓘ 87.582%
Recall ⓘ 84.123%			

Precision and recall are based on a score threshold of 0.5

[SEE FULL EVALUATION](#)

Translation products

AutoML Translation BETA

Build on top of Google's powerful Translation API with the words, phrases, and idioms that matter most to you. No machine learning experience needed. [Learn more](#)

→ Get started

Translation

Datasets

+ CREATE DATASET

Dashboard

Filter datasets

Datasets

Name Source Target Total pairs Last updated Status

No AutoML Translation datasets created yet.

Models

Create dataset

Dataset name * Use letters, numbers and underscores up to 32 characters.

Translate from ... *

Translate to ... *

CANCEL **CREATE**

[dataset_1568519781600](#) [VIEW STATS](#) [EXPORT DATA](#)

[IMPORT](#) [SENTENCES](#) [TRAIN](#) [PREDICT](#)

Select files to import

To build a custom model, you first need to import a set of sentence pairs to train it. Generally the more sentence pairs, the better. TSV and TMX files are currently supported. You can add more files later. [More Importing tips](#)

- Upload files from your computer
- Select files on Cloud Storage
- Use separate files for training, validation, and testing (advanced)

Upload files from your computer

Your files will be automatically split into training, test, and validation sets. If you have more than 100,000 sentence pairs, use the separate files option.

Maximum 500 files per import. Uploaded files will be stored on Cloud Storage. [?](#)

en-es.tsv 1 file [X](#)

[SELECT FILES](#)

Destination on Cloud Storage <gs://authentica-de791-lcm/translate/> [BROWSE](#)

[dataset_1568519781600](#) [VIEW STATS](#) [EXPORT DATA](#)

[IMPORT](#) [SENTENCES](#) [TRAIN](#) [PREDICT](#)

Translation (EN → ES)

All sentences	8,720	Filter table	
Training	6,976	Source Suggestions based on your search and browsing history Visually similar images on the web	Target Sugerencias basadas en tu historial de búsqueda y navegación Imágenes similares de la Web
Validation	872	Tayeb Salih's 88th Birthday Ehud Manor's 74th Birthday	88 aniversario del nacimiento de Tayeb Salih 74. ^a aniversario del nacimiento de Ehud Manor
Testing	872	Enter blog names or URLs, separated by commas. Is this place good for groups?	Escriba los nombres de los blogs o las URL, separados por comas. ¿Es un buen lugar para grupos?
Filter file		Most Recent YouTube Session See results in-app	Sesión de YouTube más reciente Ver resultados en la aplicación
Auto Split		Suggestions based on your search history Is this an auto body shop?	Sugerencias en función de tu historial de búsqueda ¿Es un taller de chapa y pintura?
en-es.tsv	8,720	Administrative log data for your projects	Administrar datos de registro de tus proyectos

Train new model

Model name *
dataset_156851978_20190915060931

Base model
Google NMT

Data summary
6976 training pairs, 872 validation pairs, 872 testing pairs

You will be emailed when training completes. See the [Pricing guide](#) for details about training time and cost.

START TRAINING CANCEL

← dataset_1568519781600 [VIEW STATS](#) [EXPORT DATA](#)

IMPORT SENTENCES **TRAIN** PREDICT

START TRAINING

dataset_156851978_20190915060931

Training may take several hours. You will be emailed once training completes.

Running: Training model

CANCEL

[dataset_1568519781600](#) [VIEW STATS](#) [EXPORT DATA](#)

IMPORT SENTENCES TRAIN PREDICT

Model dataset_156851978_20190915060931

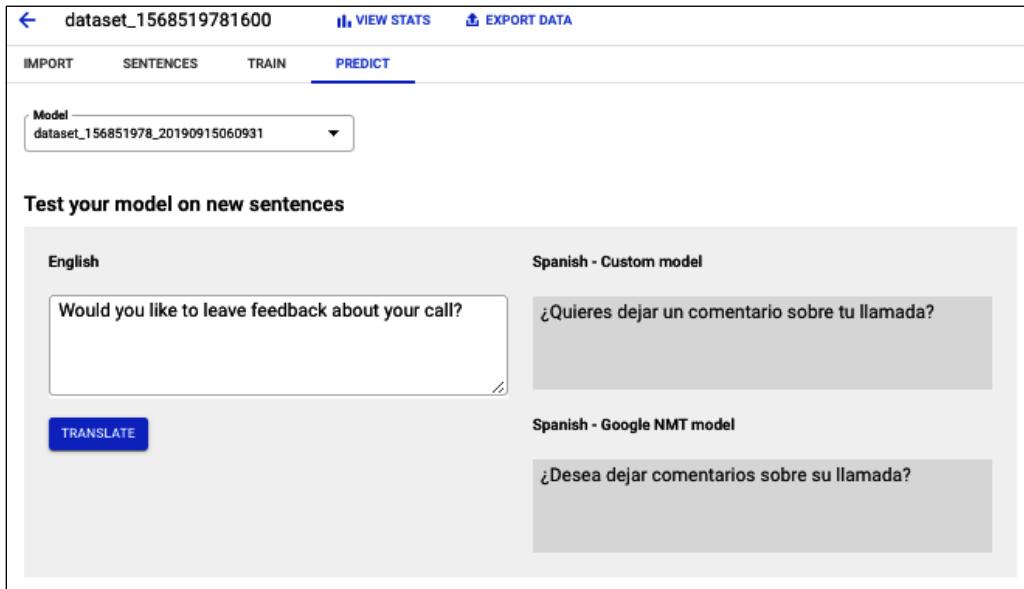
Test your model on new sentences

English	Spanish - Custom model
Would you like to leave feedback about your call?	¿Quieres dejar un comentario sobre tu llamada?

Spanish - Google NMT model

¿Desea dejar comentarios sobre su llamada?
--

TRANSLATE



Use your AutoML model

You can now translate using your custom translation model. (Note: You will need a [service account](#))

REST API PYTHON

request.json

```
{  
  "source_language_code": "en",  
  "target_language_code": "es",  
  "model": "projects/655848112025/locations/us-central1/models/TRL2303314132469809152",  
  "contents": "YOUR_SOURCE_CONTENT"  
}
```

Execute the request

```
$ curl -X POST \  
-H "Authorization: Bearer $(gcloud auth application-default print-access-token) \  
-H "Content-Type: application/json; charset=utf-8" \  
https://translation.googleapis.com/v3beta1/projects/655848112025/locations/us-central1:translateText \  
-d @request.json
```

Video Intelligence Products

AutoML Video Intelligence BETA

Train a custom video model using your own videos. No machine learning experience required. [Learn more](#)

→ Get started

Import videos

AutoML Video Intelligence uses your videos to train a custom machine learning model.

[Learn more about preparing your data.](#)

- Upload labels in your CSV, or upload un-labeled videos, and use our labeling tool.
- At least 100 video segments per label is recommended.
- Processed videos will be stored on Cloud Storage. Standard pricing applies.

Select a CSV file on Cloud Storage

The CSV file should contain paths to your train, test, and/or unassigned CSV files. Videos must be .MOV, .MPEG4, .MP4, or .AVI. [Learn more](#).

Example CSV:

```
TRAIN,gs://domestic-animals-vcm/horses/videos/train.csv  
TEST,gs://domestic-animals-vcm/horses/videos/test.csv
```

gs://*

automl-video-demo-data/hmdb_split1_mp4.csv

BROWSE

CONTINUE

Datasets BETA						CREATE DATASET
	Name	Objective	Total videos	Labeled videos	Last updated	Status
No AutoML video datasets created yet.						

[untitled_1568526765835](#) BETA

[IMPORT](#) [VIDEOS](#) [TRAIN](#) [EVALUATE](#) [TEST & USE](#)

Importing videos

This can take several minutes or more. You will be emailed when import is complete.

IMPORT	VIDEOS	TRAIN	EVALUATE	TEST & USE
All videos	5,062			
Labeled	5,062			
Unlabeled	0			
Filter labels				
Annotations ▾				
brush_hair	100			
cartwheel	100			
catch	100			
chew	100			
clap	99			
Filter videos				
		<div style="display: flex; justify-content: space-around;"> <div style="text-align: center;"> <p>hit(1)</p> </div> <div style="text-align: center;"> <p>golf(1)</p> </div> </div>		
		<div style="display: flex; justify-content: space-around;"> <div style="text-align: center;"> <p>pullup(1)</p> </div> <div style="text-align: center;"> <p>stand(1)</p> </div> </div>		

IMPORT	VIDEOS	TRAIN	EVALUATE	TEST & USE
Labels	Video segments		Train	Test
brush_hair	 100		70	30
cartwheel	 100		70	30
catch	 100		70	30
chew	 100		70	30
clap	 99		70	29
climb	 97		70	27
climb_stairs	 100		70	30
dive	 100		70	30
draw_sword	 100		70	30

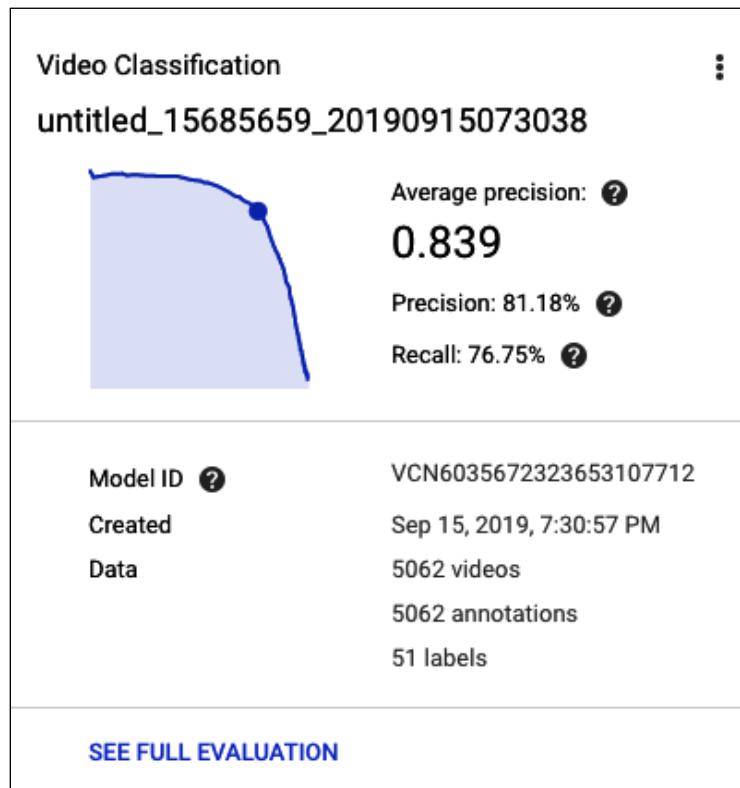
Train new model

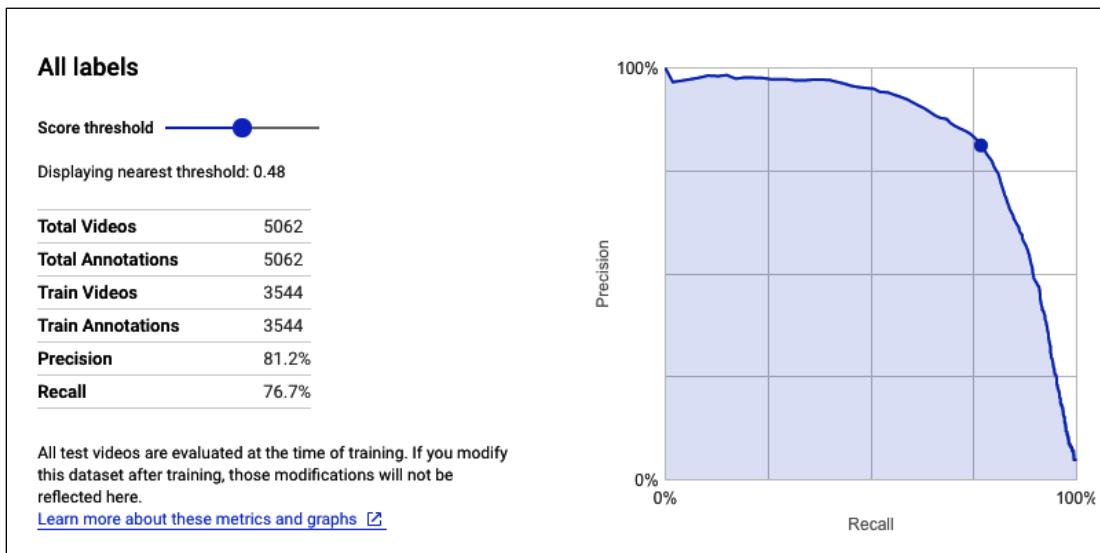
Model name *
untitled_15685659_20190915073038

Data Summary
5062 videos
51 labels

Training budget
You only pay for hours used. If your model stops improving, training will stop. [Training pricing guide](#)

START TRAINING **CANCEL**





Confusion matrix

This table helps you understand where misclassifications occur (which labels get "confused" with each other). The top three misclassifications per label are shown here.

True Label ↑	Correct Prediction	Confused with...		
brush_hair	90%	wave : 6.67%	sit : 3.33%	
cartwheel	86.67%	flic_flac : 10%	handstand : 3.33%	
catch	96.67%	jump : 3.33%		
chew	90%	drink : 6.67%	eat : 3.33%	
clap	89.66%	throw : 6.9%	pick : 3.45%	
climb	100%			
climb_stairs	73.33%	run : 13.33%	walk : 6.67%	climb : 3.33%
dive	76.67%	climb : 10%	fall_floor : 6.67%	somersault : 3.33%

IMPORT	VIDEOS	TRAIN	EVALUATE	TEST & USE
<p>Model <input type="text" value="untitled_15685659_20190915073038"/> BROWSE</p>				
<h2>Test your model</h2> <p>Create a batch prediction request with a CSV. Each row in your CSV should be a Cloud Storage file path to a video, and start/end time. Your model will output prediction results as a CSV. Learn more</p> <p>Batch prediction pricing is based on the compute resources used to generate your results. See pricing guide</p> <p>Input CSV ? gs:// * <input checked="" type="checkbox"/> automl-video-demo-data/hmdb_split1_test_gs_predict.csv BROWSE</p> <p>Results Bucket ? gs:// * <input checked="" type="checkbox"/> authentica-de791-lcm/videos/ BROWSE</p> <p>Where your prediction results are sent</p> <p>GET PREDICTIONS</p>				



0:00:00 / 0:00:00

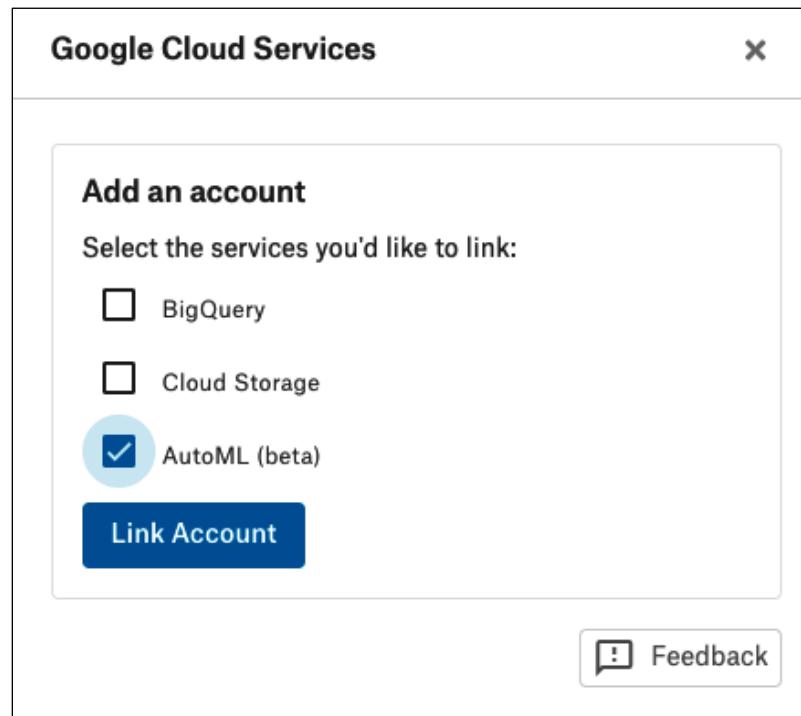
Confidence: 0.50 PREDICT ON: 1 SECOND INTERVAL SHOT SEGMENT

Predicted Results	00:00	00:00	00:01	00:01	00:02
ride_horse	0.993				

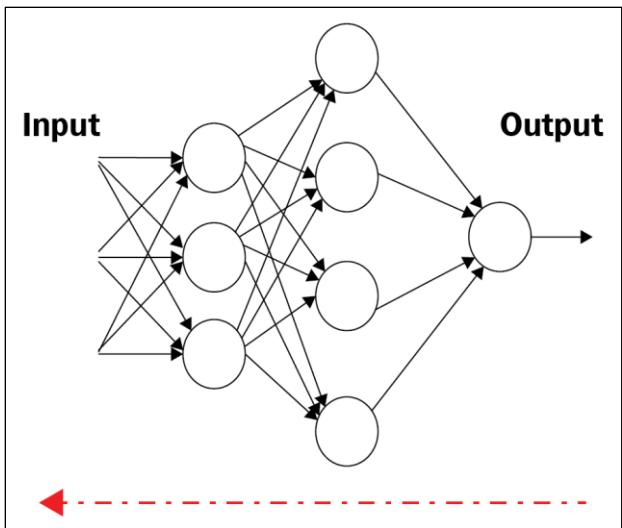
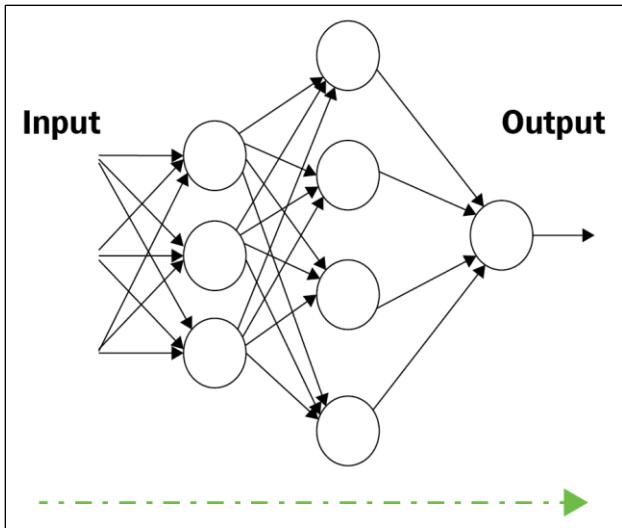
SKU	Product	SKU ID	Usage	Cost	One time credits	Discounts	Subtotal
AutoML Content Classification Model Training Operations	Cloud Natural Language API	41FE-745B-850A	3.32 hour	\$9.95	\$0.00	—	\$9.95
AutoML Tables Deployment	Cloud AutoML	3FEA-6ED1-5D9F	1,562,005,950 mebibyte second	\$2.12	\$0.00	—	\$2.12
N1 Predefined Instance Core running in Americas	Compute Engine	2E27-4F75-95CD	35.17 hour	\$1.11	\$0.00	—	\$1.11
N1 Predefined Instance Ram running in Americas	Compute Engine	6C71-E844-38BC	131.88 gibibyte hour	\$0.56	\$0.00	—	\$0.56
Class A Request Regional Storage	Cloud Storage	4DBF-185F-A415	11,336 count	\$0.03	\$0.00	—	\$0.03
AutoML Tables Online Prediction	Cloud AutoML	F664-8B0D-F8BE	0 hour	\$0.00	\$0.00	—	\$0.00
Network Internet Egress from Americas to China	Compute Engine	9DE9-9092-B3BC	0 gibibyte	\$0.00	\$0.00	—	\$0.00
AutoML Image Classification Model Training First Compute Hours	Cloud Vision API	B018-CE2C-1DF5	1 count	\$0.00	\$0.00	—	\$0.00
AutoML Tables Training	Cloud AutoML	3B5C-4F27-B029	1 hour	\$19.32	-\$19.32	—	\$0.00
Class B Request Regional Storage	Cloud Storage	7870-D10B-2763	641 count	\$0.00	\$0.00	—	\$0.00

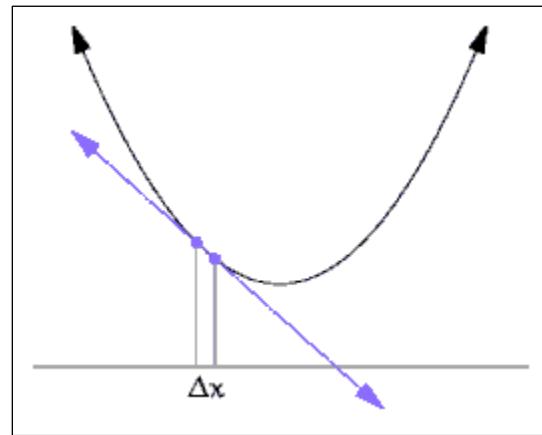
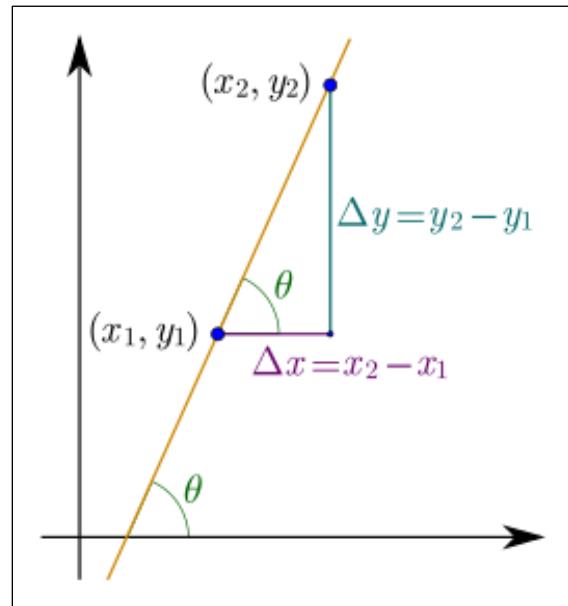
The screenshot shows a Jupyter Notebook interface with a blue 'k' logo in the top-left corner. The title bar reads "kernel5fb6788945 Draft saved". The menu bar includes File, Edit, Insert, Run, Add-ons, and Help. A toolbar below has buttons for New, Open, Run All, and a plus sign. The main area has a code cell labeled "In[1]:" containing Python code. A tooltip for "Google Cloud Services" is visible. The status bar at the bottom right shows "Draft Session (0m)" and resource usage (CPU, RAM, HDD).

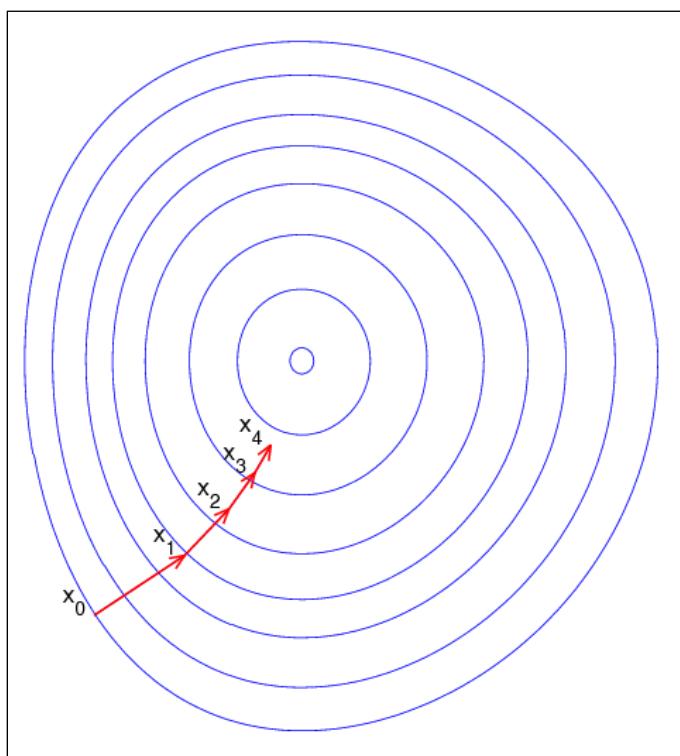
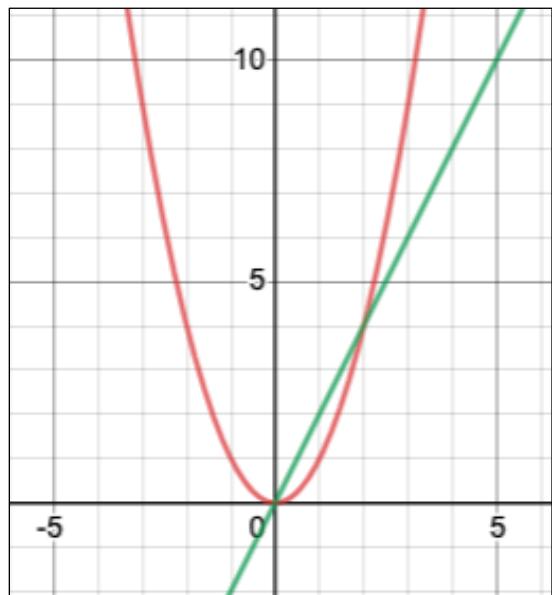
```
# This Python 3 environment  
# It is defined by the  
# For example, here's  
  
import numpy as np # linear algebra  
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)  
  
# Input data files are available in the "../input/" directory.  
# For example, running this (by clicking run or pressing Shift+Enter) will list all files  
  
import os  
for dirname, _, filenames in os.walk('/kaggle/input'):  
    for filename in filenames:  
        print(os.path.join(dirname, filename))  
  
# Any results you write to the current directory are saved as output.
```

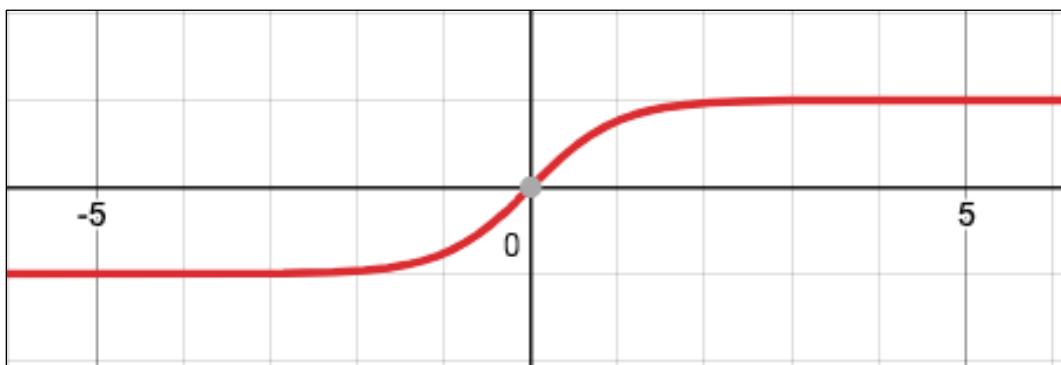
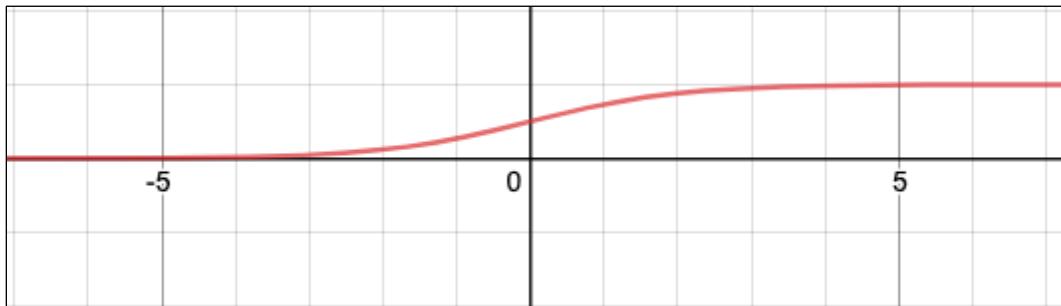
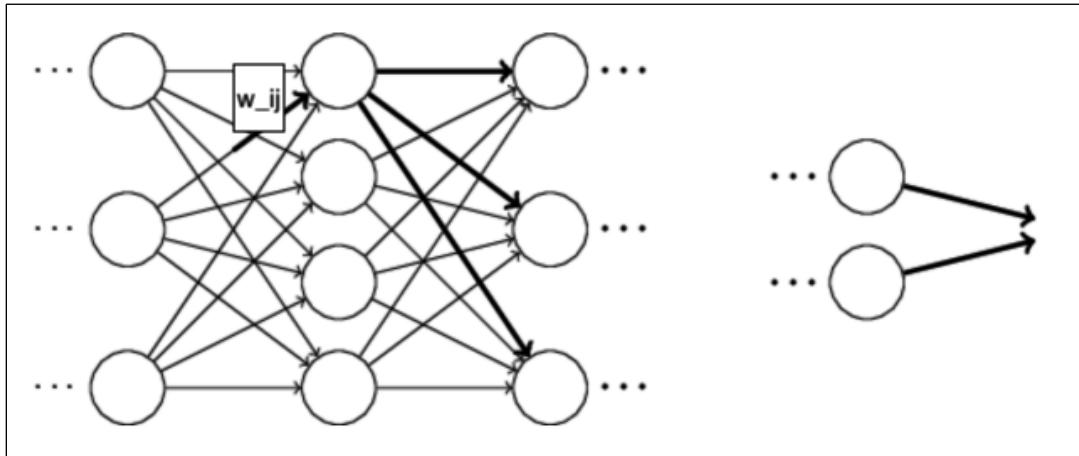


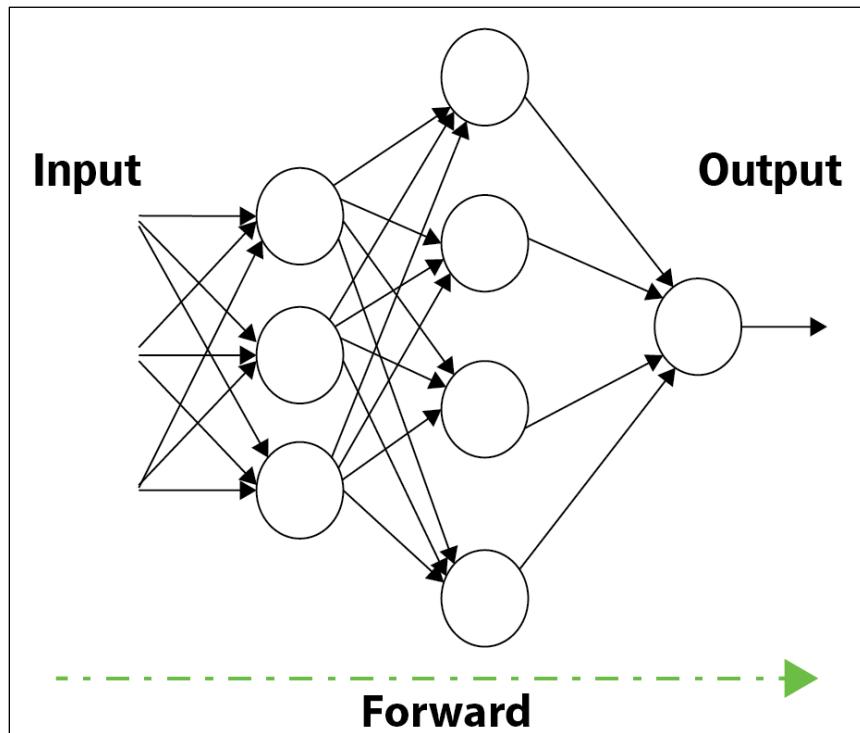
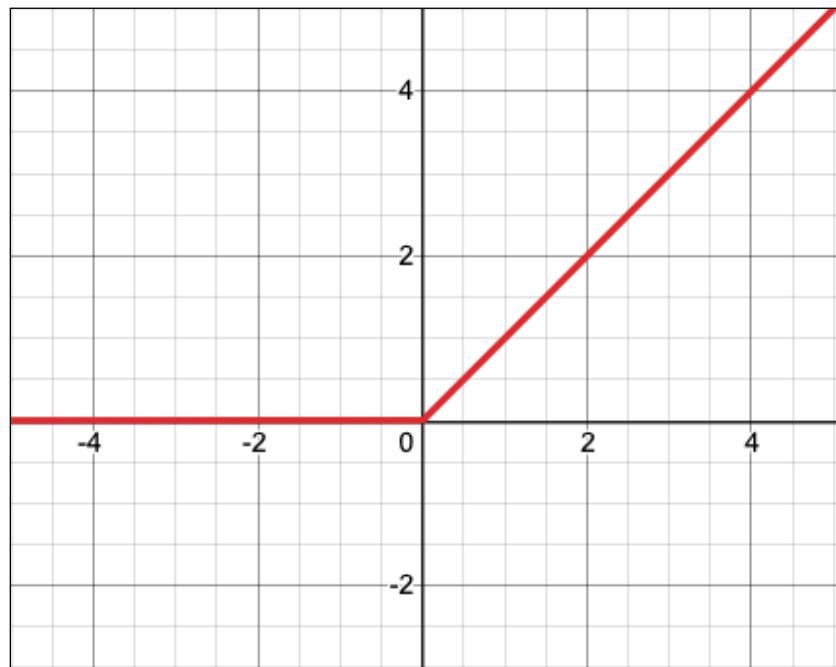
Chapter 15: The Math Behind Deep Learning

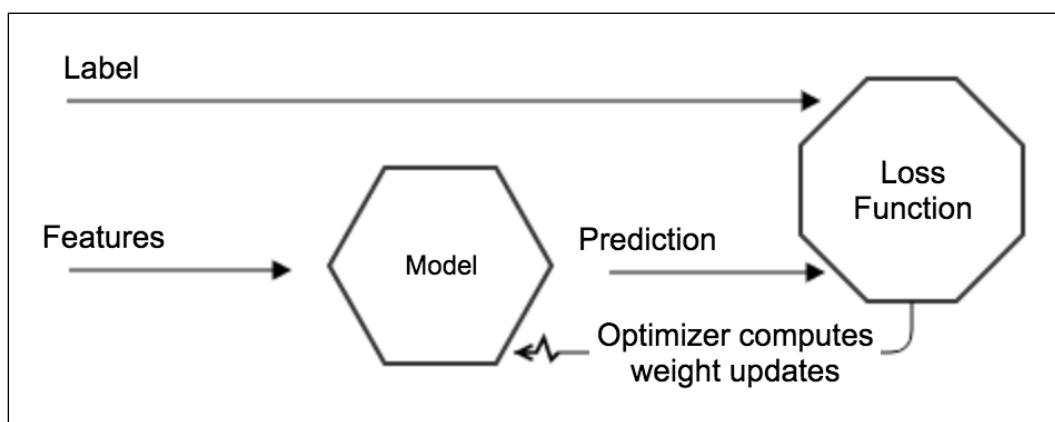
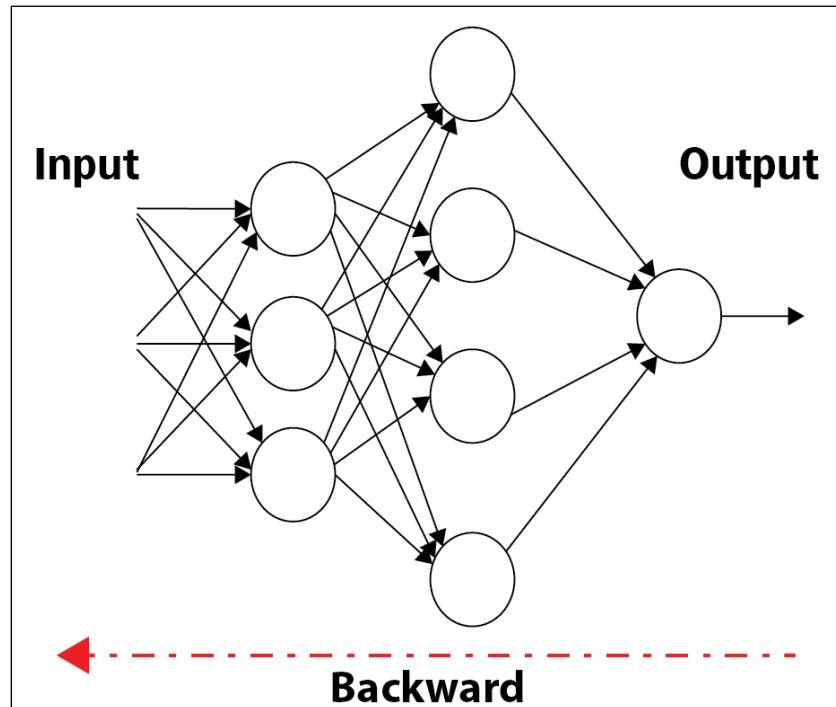


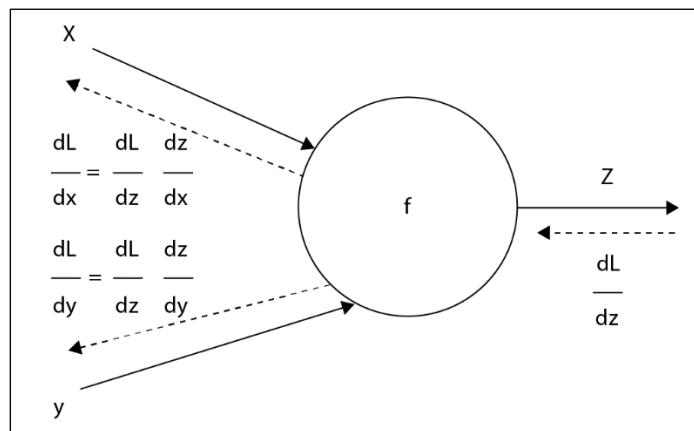
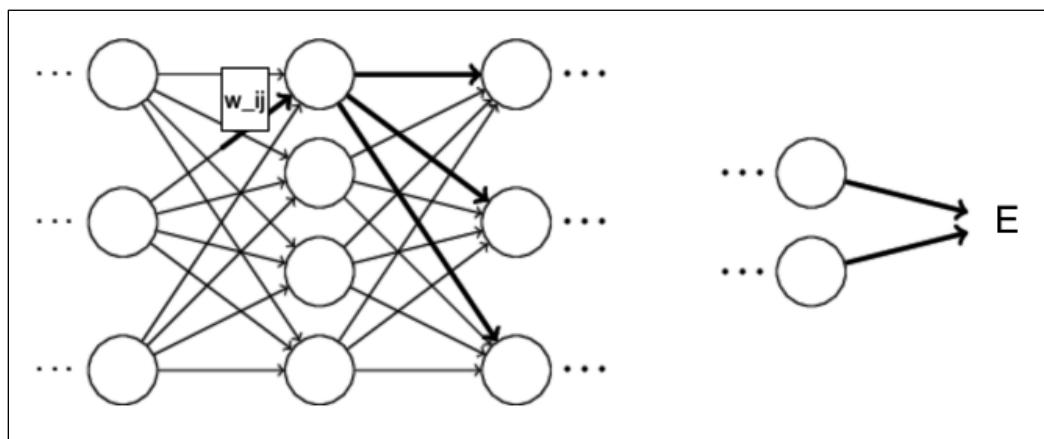
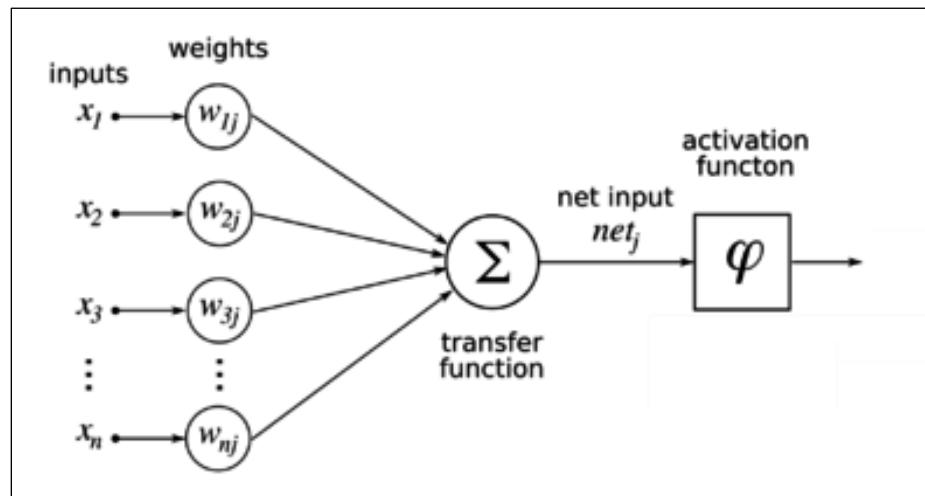








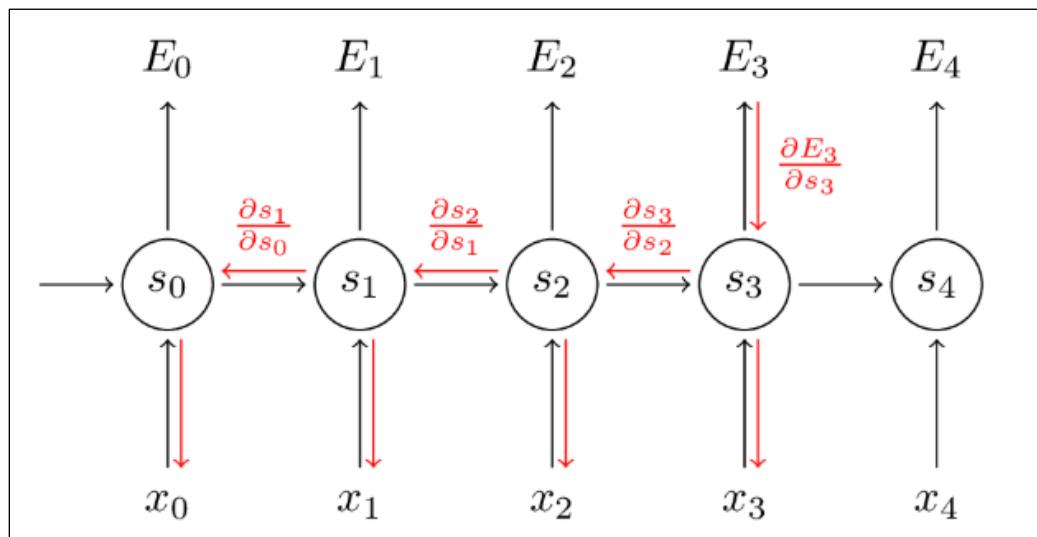
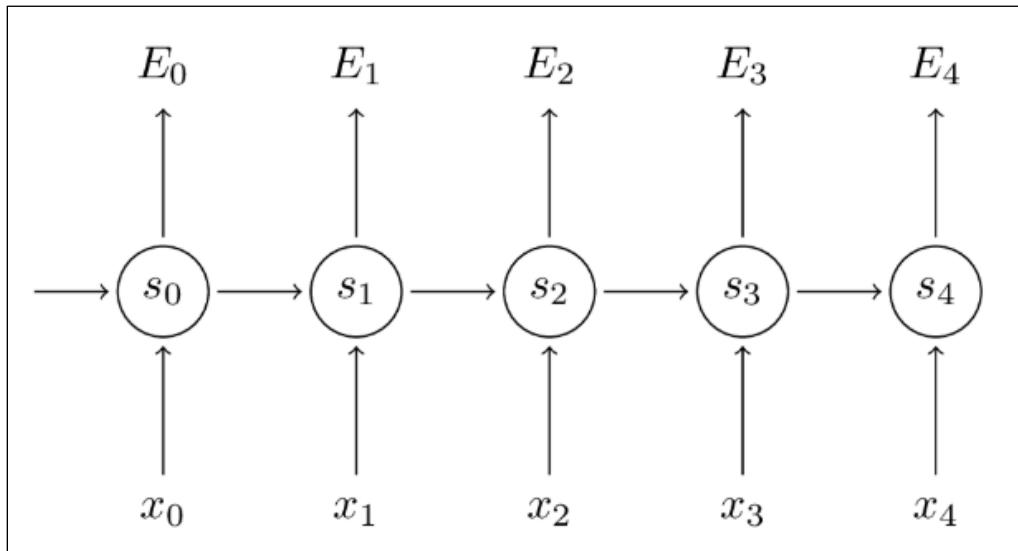




X_{11}	X_{12}	X_{13}
X_{21}	X_{22}	X_{23}
X_{31}	X_{32}	X_{33}

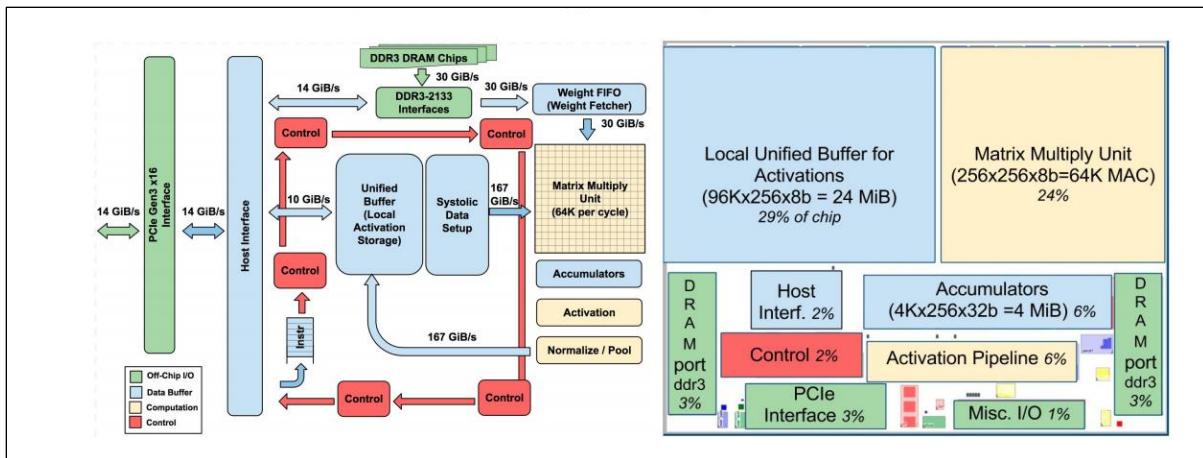
W_{11}	W_{12}
W_{21}	W_{22}

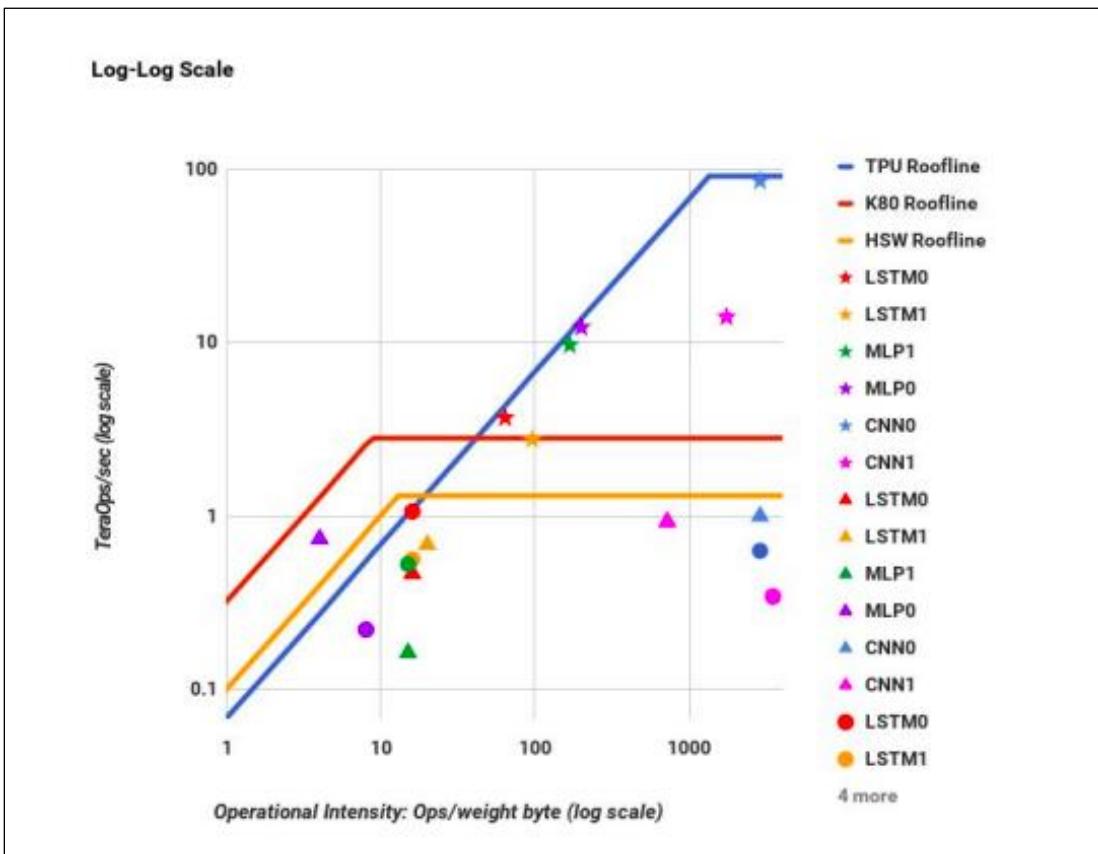
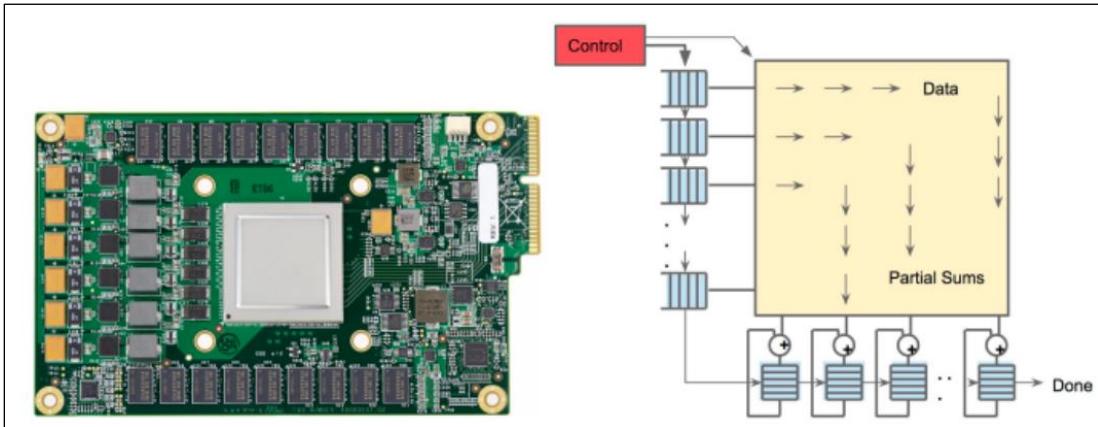
$W_{11}X_{11} + W_{12}X_{12} + W_{21}X_{21} + W_{22}X_{22}$	$W_{11}X_{12} + W_{12}X_{13} + W_{21}X_{21} + W_{22}X_{23}$
$W_{11}X_{21} + W_{12}X_{22} + W_{21}X_{31} + W_{22}X_{32}$	$W_{11}X_{22} + W_{12}X_{23} + W_{21}X_{32} + W_{22}X_{33}$

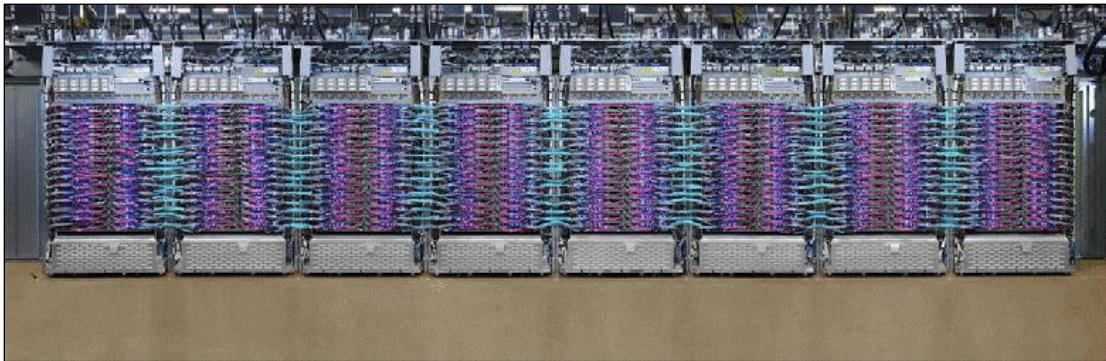
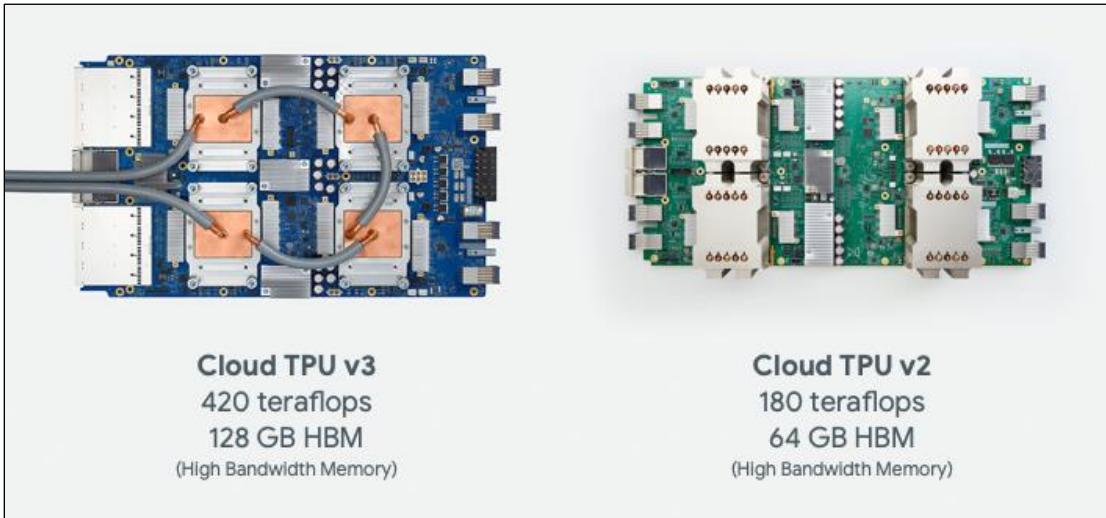


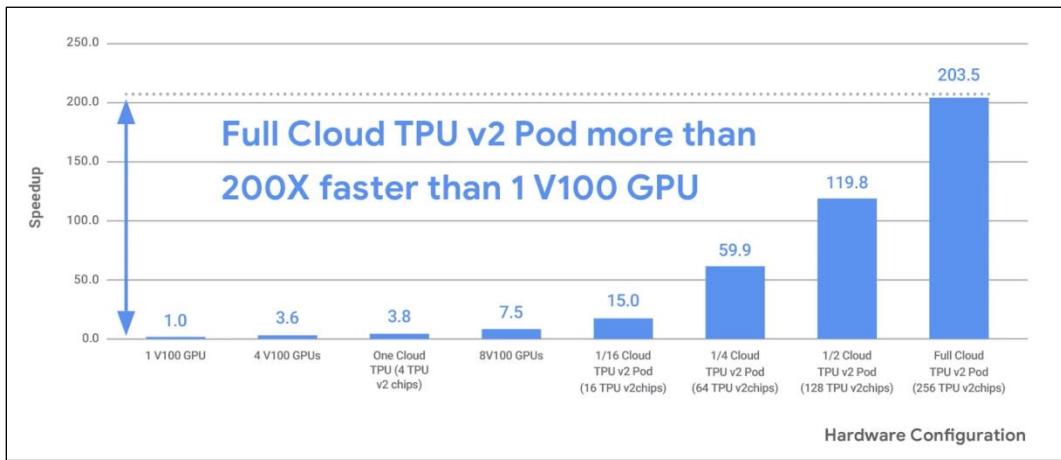
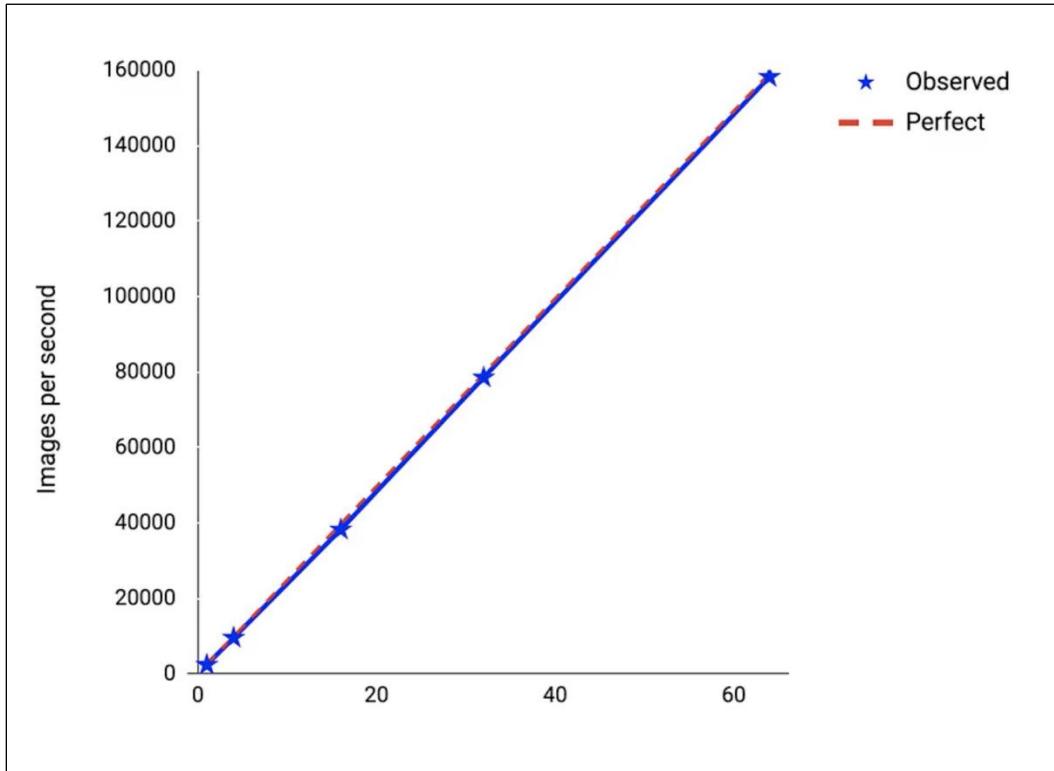
Chapter 16: Tensor Processing Unit

Name	LOC	Layers					Nonlinear function	Weights	TPUv1 Ops / Weight Byte	TPUv1 Batch Size	% Deployed
		FC	Conv	Vector	Pool	Total					
MLP0	0.1k	5				5	ReLU	20M	200	200	61%
MLP1	1k	4				4	ReLU	5M	168	168	
LSTM0	1k	24		34		58	sigmoid, tanh	52M	64	64	29%
LSTM1	1.5k	37		19		56	sigmoid, tanh	34M	96	96	
CNN0	1k		16			16	ReLU	8M	2888	8	5%
CNN1	1k	4	72		13	89	ReLU	100M	1750	32	









Notebook settings

Runtime type
Python 3

Hardware accelerator
TPU 

Omit code cell output when saving this notebook

CANCEL **SAVE**

Cloud TPUs

This repository is a collection of reference models and tools used with [Cloud TPUs](#).

The fastest way to get started training a model on a Cloud TPU is by following the tutorial. Click the button below to launch the tutorial using Google Cloud Shell.



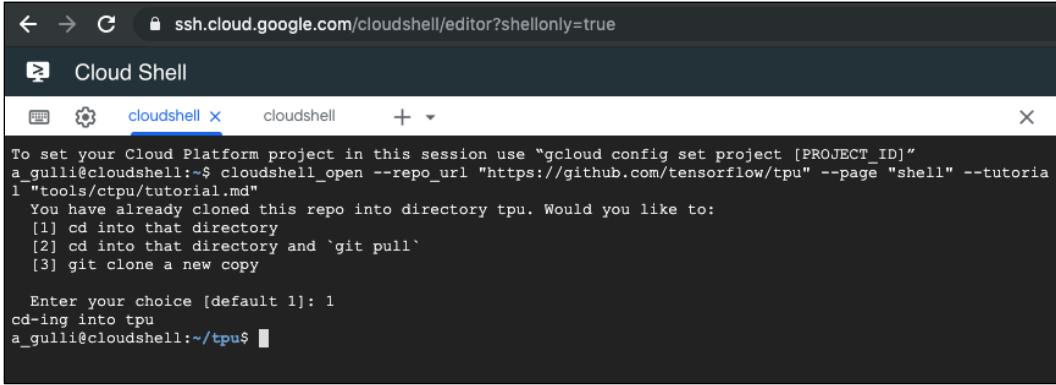
[OPEN IN GOOGLE CLOUD SHELL](#)

Note: This repository is a public mirror, pull requests will not be accepted. Please file an issue if you have a feature or bug request.

Running Models

To run models in the `models` subdirectory, you may need to add the top-level `/models` folder to the Python path with the command:

```
export PYTHONPATH="$PYTHONPATH:/path/to/models"
```



To set your Cloud Platform project in this session use "gcloud config set project [PROJECT_ID]"
a_gulli@cloudshell:~\$ cloudshell_open --repo_url "https://github.com/tensorflow/tpu" --page "shell" --tutorial
l "tools/ctpu/tutorial.md"
You have already cloned this repo into directory tpu. Would you like to:
[1] cd into that directory
[2] cd into that directory and `git pull`
[3] git clone a new copy
Enter your choice [default 1]: 1
cd-ing into tpu
a_gulli@cloudshell:~/tpu\$

ctpu quickstart

Introduction

This Google Cloud Shell tutorial walks through how to use the open source [ctpu](#) tool to train an image classification model on a Cloud TPU. In this tutorial, you will:

1. Confirm the configuration of `ctpu` through a few basic commands.
2. Launch a Cloud TPU "flock" (a Compute Engine VM and Cloud TPU pair).
3. Create a [Cloud Storage](#) bucket for your training data.
4. Download the [MNIST dataset](#) and prepare it for use with a Cloud TPU.
5. Train a simple convolutional neural network on the MNIST dataset to recognize handwritten digits.
6. Begin training a modern convolutional neural network ([ResNet-50](#)) on a simulated dataset.
7. View performance and other metrics using [TensorBoard](#).
8. Clean everything up!

Before you get started, be sure you have created a GCP Project with [billing enabled](#). When you have the [project ID](#) in hand (the "short name" found on the cloud console's main landing page), click "Continue" to get started!



Martin Görner
@martin_gorner

Full Keras / TPU support coming in Tensorflow 2.1. One line of code to get you model running on a TPU or TPU pod.

`model.fit()` or custom training loops.

I am presenting this at TF World now but you can already try with tf-nightly. Demo at bit.ly/keras-tpu-tf21